
Quantitative Risk Management

FIN-417

Homework 1

BY WILLIAM JALLOT (SCIPER 341540)
ANTOINE GARIN (SCIPER 327295)
MATTHIAS WYSS (SCIPER 329884)

QUANTITATIVE RISK MANAGEMENT, MA3

DR. URBAN ULRYCH

Contents

1 Exercise 1: Empirical Stylized Facts	1
1.A Log-returns	1
1.B ACF and CCF	1
1.C QQ-plots and Jarque–Bera test	1
2 Exercise 2: Value-at-Risk and Expected Shortfall	5
2.A Historical Simulation	5
2.B Gaussian Model	6
2.C Student-t Model	6
2.D Conditional parametric modeling: AR(p) + GARCH(1,1)	6
2.E Filtered Historical Simulation (FHS)	8
2.F Results	8
3 Exercise 3	10
3.A Backtesting VaR and ES backtesting procedure described	10
3.B Acerbi and Székely	12
4 Exercise 4	14
4.A Pseudo-observations	14
4.B Fit the copulas	15
4.C Simulate from the copulas	16
5 Exercise 5	19
5.A Univariate models	19
5.B Copulas	21
5.C Backtesting and Comparison	21
A Code Appendix	27

1 Exercise 1: Empirical Stylized Facts

1.A Log-returns

We collected daily adjusted closing prices for **AAPL**, **META**, and **JPM** between January 1, 2023 and June 30, 2025 using the `yfinance` Python package. The daily log-returns were computed as:

$$R_t^{(i)} = \log\left(\frac{P_t^{(i)}}{P_{t-1}^{(i)}}\right), \quad i \in \{\text{AAPL, META, JPM}\},$$

where $P_t^{(i)}$ denotes the adjusted closing price of asset i at day t . Missing data were removed after aligning the series by date.

Figure 1 displays the resulting log-return series. As expected, returns fluctuate around zero and show no persistent trend, consistent with weak-form market efficiency. However, volatility is clearly time-varying: META exhibits the largest fluctuations, AAPL intermediate, and JPM the lowest. Periods of high volatility tend to cluster together, illustrating the well-known phenomenon of *volatility clustering*.

1.B ACF and CCF

To study temporal and cross-asset dependence, we computed autocorrelation and cross-correlation functions for both raw and absolute returns. The empirical autocorrelation function (ACF) for asset i at lag h is defined as:

$$\rho_i(h) = \frac{\text{Cov}(R_t^{(i)}, R_{t-h}^{(i)})}{\sqrt{\text{Var}(R_t^{(i)}) \text{Var}(R_{t-h}^{(i)})}}, \quad h = 0, 1, \dots, 25.$$

The cross-correlation between two assets i and j is given by:

$$\rho_{i,j}(h) = \frac{\text{Cov}(R_t^{(i)}, R_{t-h}^{(j)})}{\sqrt{\text{Var}(R_t^{(i)}) \text{Var}(R_{t-h}^{(j)})}}.$$

Both ACFs and CCFs were also computed using absolute returns $|R_t^{(i)}|$ to analyze volatility spillovers. Figure 2 and Figure 3 shows the results.

The autocorrelation of log returns for AAPL, META, and JPM shows almost no correlation after lag 0. This indicates that past returns have little predictive power for future returns, confirming the weak linear dependence typical of stock prices. Absolute returns exhibit stronger but still weak autocorrelations, with lag 1 correlations around 0.2 for AAPL.

Cross-correlations between different stocks' log returns are moderate at lag 0 (roughly 0.27-0.41) and decay quickly for higher lags. This suggests some contemporaneous co-movement across assets, particularly between AAPL and META, but limited predictive influence across weeks.

Cross-correlations of absolute returns are weaker than the autocorrelations of absolute returns, typically ranging from 0.05 to 0.15 at lag 0. However, they remain positive for most lags, indicating modest volatility spillover across assets.

1.C QQ-plots and Jarque–Bera test

The normality of log-returns was examined through both graphical and statistical methods. Figure 4 presents QQ-plots comparing the empirical quantiles of returns with those from a standard normal distribution. For all three assets, the curves deviate from the 45° line: points lie below the line on the left tail and above it on the right, indicating left skewness and excess kurtosis (fat tails).

To confirm this, we applied the *Jarque–Bera* test, which jointly tests skewness and kurtosis deviations from the normal distribution:

$$\text{JB} = \frac{n}{6} \left(S^2 + \frac{(K - 3)^2}{4} \right),$$

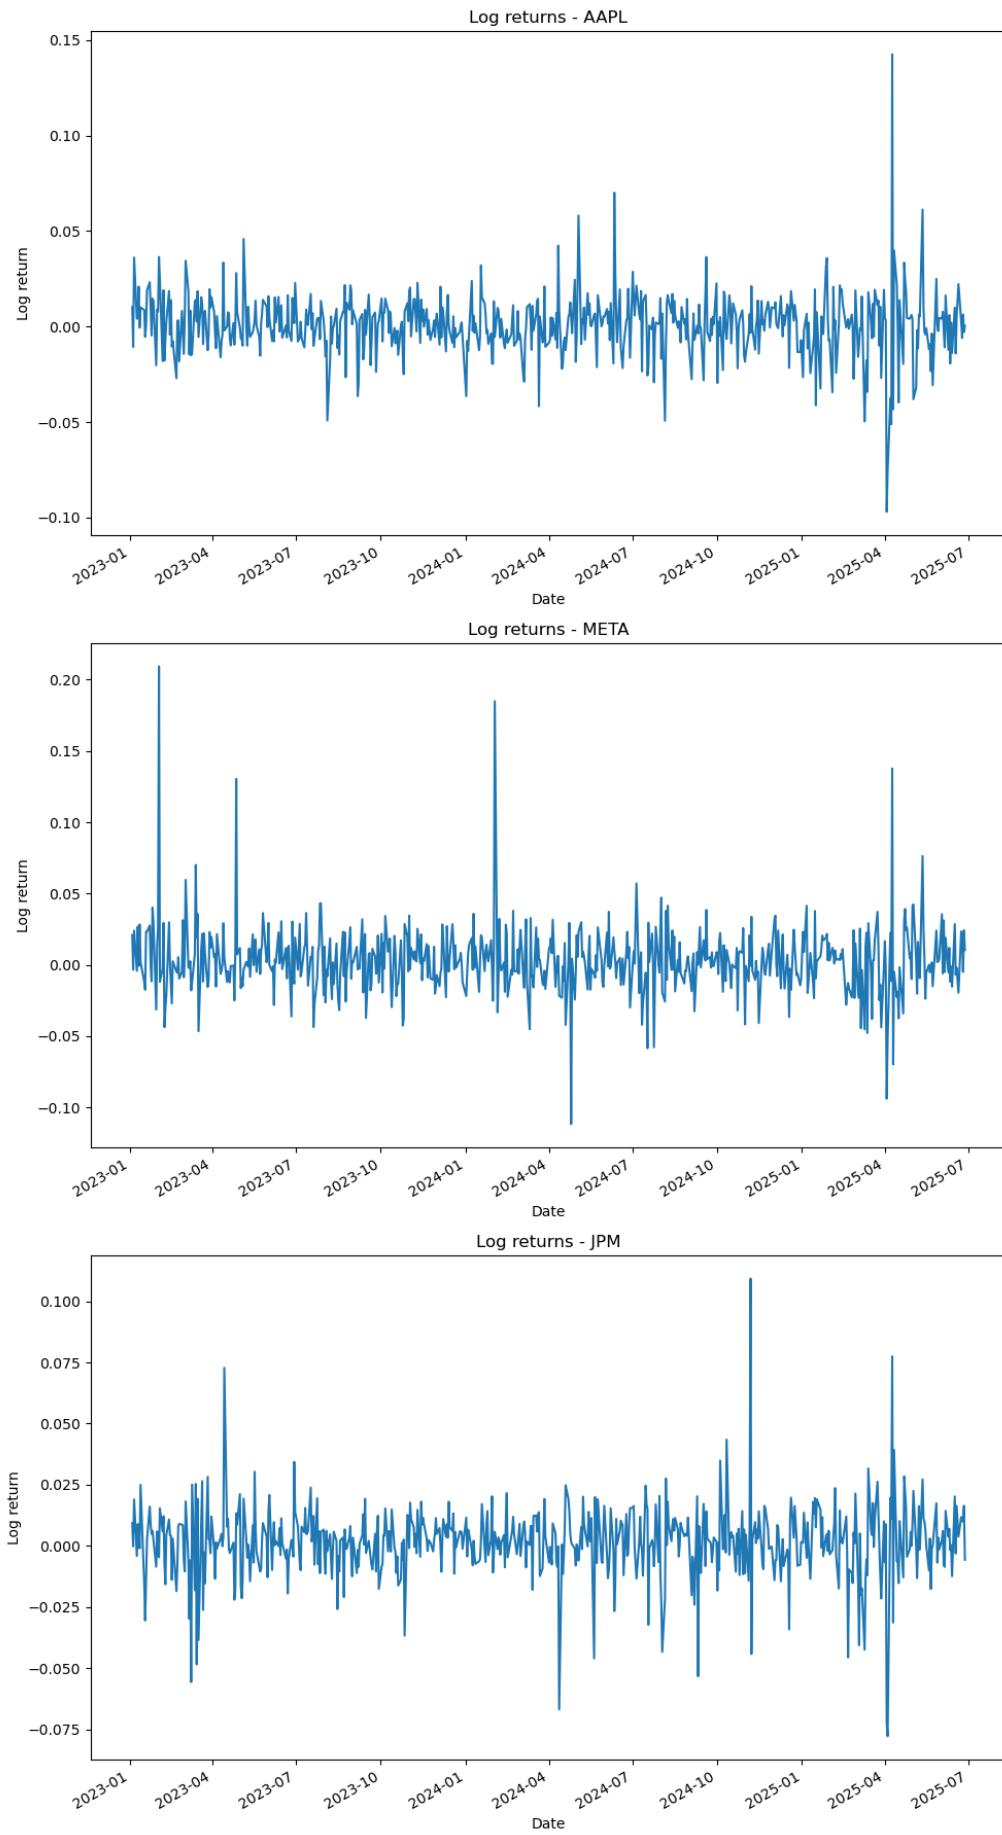


Figure 1: Daily log-returns for AAPL, META, and JPM over the period January 2023 – June 2025.

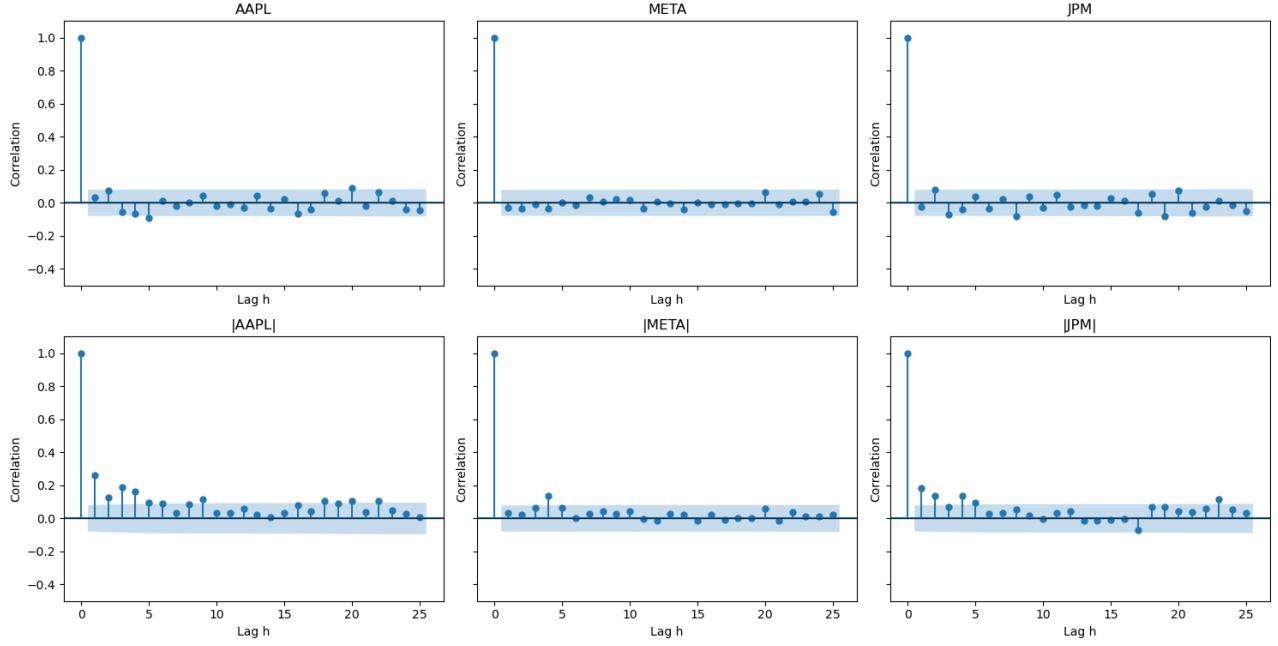


Figure 2: Autocorrelation function of raw and absolute log-returns up to lag 25.

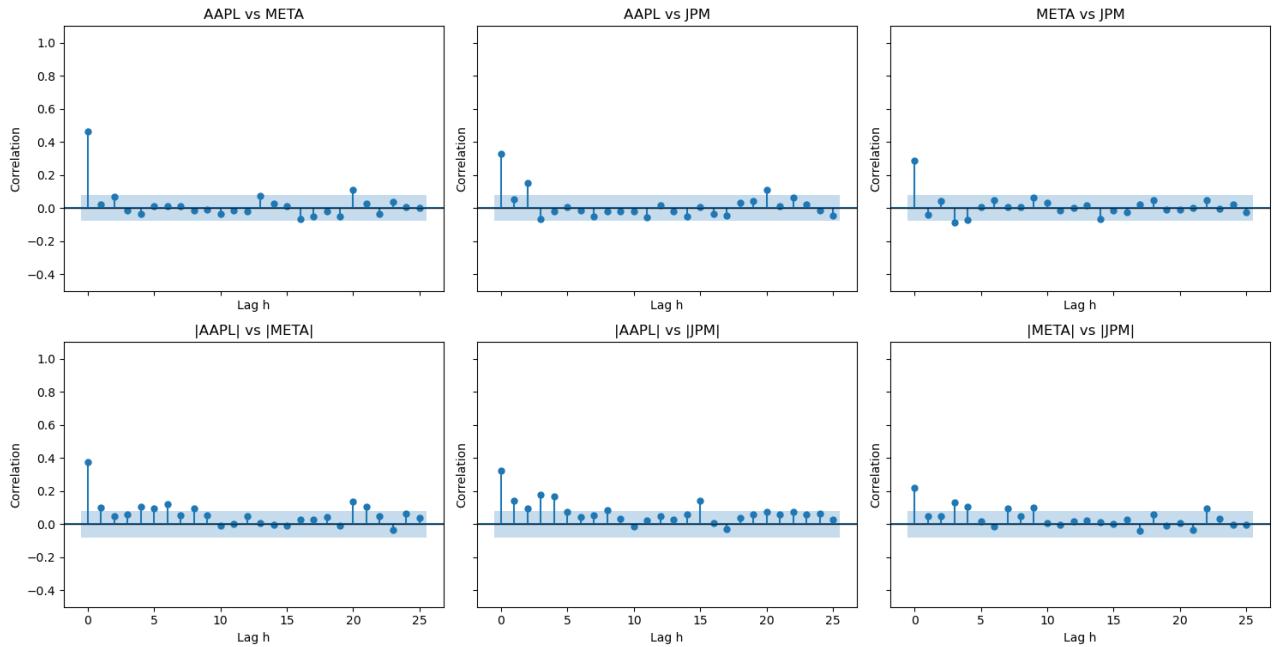


Figure 3: Cross-correlation function of raw and absolute log-returns up to lag 25.

Table 1: Jarque–Bera normality test results.

Asset	JB Statistic	p-value	Normality (5%)
AAPL	3619.30	0.000	Reject
META	7608.77	0.000	Reject
JPM	2203.51	0.000	Reject

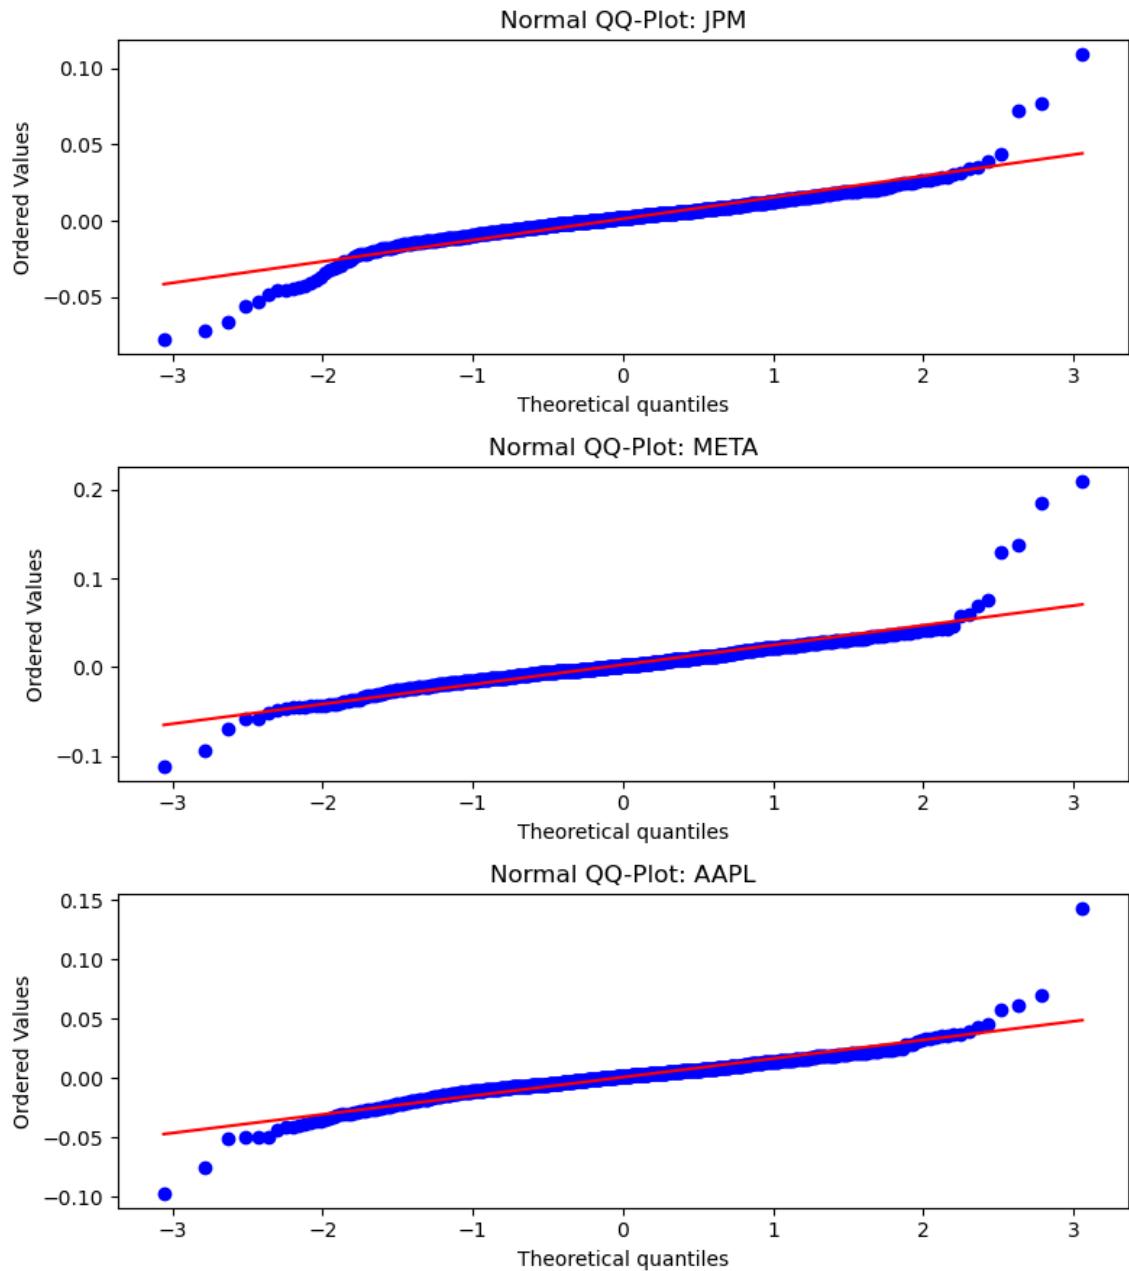


Figure 4: Normal QQ-plots of daily log-returns for AAPL, META, and JPM.

where S is the sample skewness and K the sample kurtosis. The null hypothesis H_0 : "the data follow a normal distribution" was rejected for all assets at the 5% significance level (Table 1). In summary, the log-returns exhibit the main empirical stylized facts of financial data: (i) absence of linear autocorrelation, (ii) volatility clustering, (iii) cross-asset volatility dependence, and (iv) non-normal, heavy-tailed return distributions.

2 Exercise 2: Value-at-Risk and Expected Shortfall

Single-window modeling: VaR, ES, and distributions. Analyze each asset separately on the first estimation window of length $W = 252$, i.e., approximately one trading year. Clearly describe your fitting procedure and any choices made. Use the loss convention $L_t = -R_t$, for $t = 1, \dots, W$ (i.e., right tail is risky). *Fit and compare* the following models and report the estimated 1-day ahead VaR and ES at levels $\alpha = 95\%$ and $\alpha = 99\%$. To compare different fitted loss distributions, plot their estimated probability density functions (pdfs) in a single plot. This allows you to visually assess the fit and tail behavior.

- a) *Historical simulation:* obtain the empirical cdf and use it to compute the VaR/ES of L .
- b) *Gaussian:* fit (μ, σ) and use the closed-form solutions for VaR/ES.
- c) *Student-t:* fit (ν, μ, σ) using maximum likelihood estimation (MLE) and use the closed-form solutions for VaR/ES; discuss the impact of degrees of freedom on tails.
- d) *Conditional parametric:* the previous models assume i.i.d. losses, now we make use of conditional models. The mean is defined through an autoregressive model AR(p) where you choose p based on the autocorrelation function (ACF) from 1.b) and the partial ACF. The volatility is modeled through a GARCH(1,1) model, with Gaussian innovations. Use the fitted model to obtain 1-step-ahead VaR/ES forecasts based on the conditional mean $\hat{\mu}_{W+1}$ and volatility $\hat{\sigma}_{W+1}$ estimates from the AR(p) and GARCH(1,1) parts, respectively.
- e) *Filtered Historical Simulation (FHS):* describe and implement the FHS procedure following Barone-Adesi et al. (1999) or a similar reference. Using the fitted AR(p) + GARCH(1,1) approach from 2.d), obtain the model's residuals $\hat{\epsilon}_t$ and standardized residuals $\tilde{\epsilon}_t = \hat{\epsilon}_t / \hat{\sigma}_t$, $t = 1, \dots, W$. Instead of assuming a parametric distribution for the innovation $\tilde{\epsilon}_{W+1}$ (such as in 2.d), employ the non-parametric bootstrap (i.e., resampling $M = 1000$ values with replacement) from the constructed sample of $\tilde{\epsilon}$ to obtain the 1-step-ahead VaR/ES forecasts.

In this exercise, we compute *Value-at-Risk (VaR)* and *Expected Shortfall (ES)* for the daily losses

$$L_t = -R_t,$$

where R_t are the log-returns from Exercise 1. We consider an estimation window of $W = 252$ days and compare five approaches: historical simulation, Gaussian, Student-t, conditional parametric (AR + GARCH), and filtered historical simulation (FHS).

2.A Historical Simulation

The *historical simulation* approach uses the empirical cumulative distribution function (CDF) of past losses. For a given confidence level α , the VaR is defined as the α -quantile of the empirical distribution:

$$\text{VaR}_\alpha^{\text{hist}} = \inf\{x \mid F_L(x) \geq \alpha\}.$$

The Expected Shortfall is the conditional expectation of losses exceeding the VaR:

$$\text{ES}_\alpha^{\text{hist}} = \mathbb{E}[L \mid L \geq \text{VaR}_\alpha^{\text{hist}}].$$

2.B Gaussian Model

Assuming losses follow a normal distribution $L_t \sim \mathcal{N}(\mu, \sigma^2)$, we fit the mean μ and standard deviation σ by maximum likelihood. The closed-form VaR and ES formulas are:

$$\text{VaR}_\alpha^{\text{Gauss}} = \mu + \sigma \Phi^{-1}(\alpha),$$

$$\text{ES}_\alpha^{\text{Gauss}} = \mu + \frac{\sigma \phi(\Phi^{-1}(\alpha))}{1 - \alpha},$$

where Φ^{-1} and ϕ denote the inverse CDF and PDF of the standard normal distribution, respectively.

2.C Student-t Model

To capture heavy tails, we fit a Student-t distribution with degrees of freedom ν , location μ_t , and scale σ_t using maximum likelihood. VaR and ES are:

$$\begin{aligned}\text{VaR}_\alpha^t &= \mu_t + \sigma_t t_\nu^{-1}(\alpha), \\ \text{ES}_\alpha^t &= \mu_t + \sigma_t \frac{\nu + (t_\nu^{-1}(\alpha))^2 t_\nu(t_\nu^{-1}(\alpha))}{\nu - 1} \frac{1}{1 - \alpha},\end{aligned}$$

where t_ν^{-1} and t_ν denote the inverse CDF and PDF of the Student-t distribution.

Lower degrees of freedom ν imply heavier tails, leading to higher VaR and ES estimates at high confidence levels. This highlights the importance of modeling fat tails in financial loss distributions.

2.D Conditional parametric modeling: AR(p) + GARCH(1,1)

To account for conditional dynamics in losses, we model the mean using an AR(p) process and the volatility using a GARCH(1,1) model with Gaussian innovations.

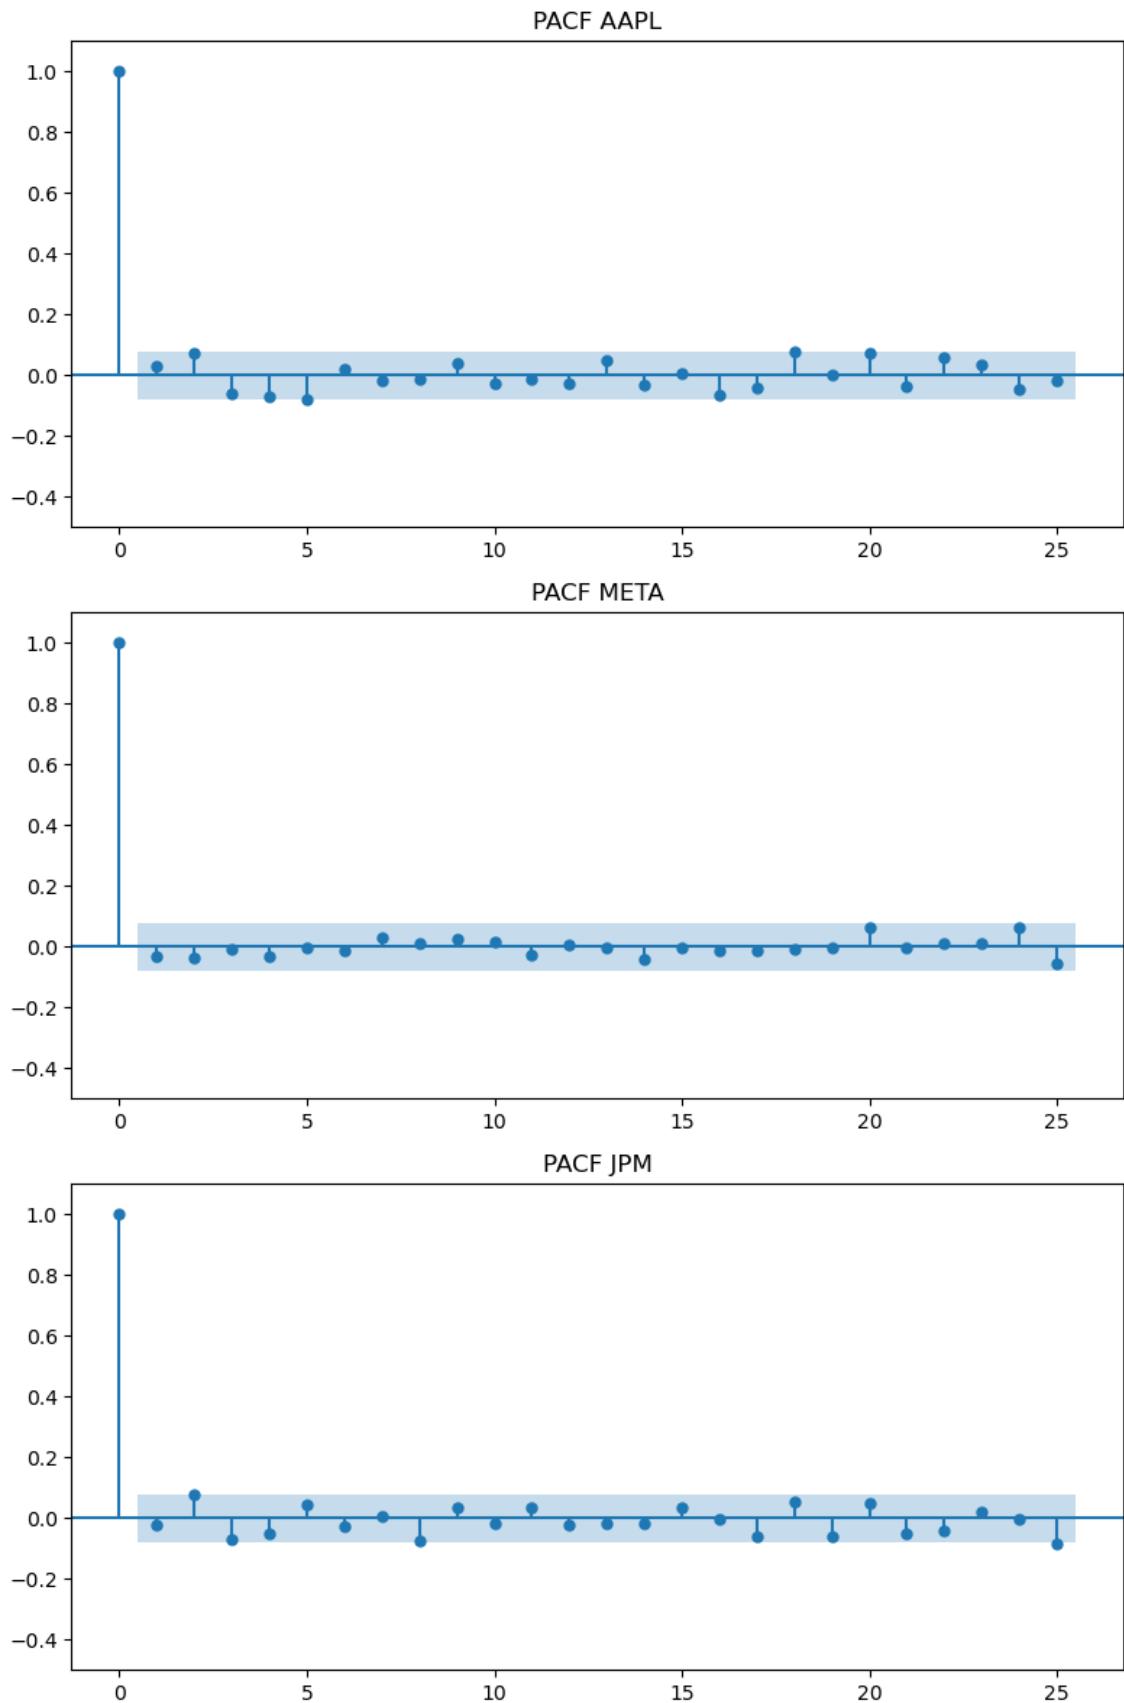


Figure 5: Partial autocorrelation function (PACF) of daily log-returns for AAPL, META, and JPM. The blue shaded area represents the 95% confidence interval under the null hypothesis of no autocorrelation.

The order p of the AR model is determined by inspecting the partial autocorrelation function (PACF)

of the log-returns in Figure 5. All PACF coefficients for lags greater than zero lie within the 95% confidence band and fluctuate around zero, indicating the absence of significant linear dependence at positive lags. Therefore, an AR(0) specification, corresponding to a pure white noise model, is sufficient to capture the linear structure in the mean of the losses.

Let L_t denote the loss at time t . Denoting by $\hat{\mu}_{W+1}$ and $\hat{\sigma}_{W+1}$ the 1-step-ahead forecasts of the conditional mean and volatility from the fitted AR(0) + GARCH(1,1) model, the conditional Value-at-Risk (VaR) and Expected Shortfall (ES) are computed using the standard closed-form VaR and ES formulas:

$$\text{VaR}_\alpha^{\text{cond}} = \hat{\mu}_{W+1} + \hat{\sigma}_{W+1} z_\alpha, \quad \text{ES}_\alpha^{\text{cond}} = \hat{\mu}_{W+1} + \frac{\hat{\sigma}_{W+1} \phi(\Phi^{-1}(\alpha))}{1 - \alpha},$$

where $\phi(\cdot)$ is the standard normal density.

This approach captures both the conditional mean and conditional heteroskedasticity in the loss series.

2.E Filtered Historical Simulation (FHS)

Filtered Historical Simulation combines the conditional AR(0) + GARCH(1,1) model with a non-parametric treatment of standardized residuals. First, we fit the AR(0) model to the losses and compute residuals $\hat{\varepsilon}_t = L_t - \hat{\mu}_t$. Then, a GARCH(1,1) model is fitted to the residuals to estimate conditional volatility $\hat{\sigma}_t$ and obtain standardized residuals:

$$\tilde{\varepsilon}_t = \frac{\hat{\varepsilon}_t}{\hat{\sigma}_t}, \quad t = 1, \dots, W.$$

Instead of assuming a parametric distribution for the innovation $\tilde{\varepsilon}_{W+1}$, we apply a non-parametric bootstrap: we resample $M = 1000$ values with replacement from the standardized residuals to generate 1-step-ahead scenarios. The simulated losses are then reconstructed as:

$$L_{W+1}^* = \hat{\mu}_{W+1} + \hat{\sigma}_{W+1} \tilde{\varepsilon}_{W+1}^*.$$

Finally, the 1-step-ahead VaR and ES forecasts are obtained from the empirical quantiles and averages of the simulated loss distribution:

$$\begin{aligned} \text{VaR}_\alpha^{\text{FHS}} &= \text{Quantile}_\alpha(L_{W+1}^*) = \inf\{x \mid F_{L_{W+1}^*}(x) \geq \alpha\}, \\ \text{ES}_\alpha^{\text{FHS}} &= \mathbb{E}[L_{W+1}^* \mid L_{W+1}^* \geq \text{VaR}_\alpha^{\text{FHS}}]. \end{aligned}$$

This method combines the benefits of conditional modeling (capturing time-varying volatility and mean) with the flexibility of a non-parametric distribution for innovations, providing a more realistic estimation of tail risk.

2.F Results

Table 2 and Table 3 report the estimated VaR and ES for each asset and confidence level. Figure 6 shows the estimated PDFs of losses.

Note that for the Conditional parametric model, the PDF used to compute VaR and ES is theoretical and conditional on the information available at time W . In contrast, the FHS approach estimates the next-day loss distribution non-parametrically via simulation of standardized residuals, rather than using an analytical formula. Historical simulation also relies on an empirical distribution and a kernel density estimate is used for visualization.

Across all three assets and both confidence levels, the five methods produce VaR estimates that lie within a comparable numerical range. This indicates that, despite relying on different distributional assumptions, the models agree on the order of magnitude of tail losses. But Gaussian VaR is systematically lower due to thin tails, while Student-t and Historical tend to produce higher values at the 99% level. FHS and the conditional parametric model generally fall between these extremes.

Table 2: Value-at-Risk (VaR) results for the first 252-day window.

Ticker	Alpha	Historical	Gaussian	Student-t	Parametric	FHS
AAPL	95%	0.017587	0.019421	0.018856	0.019611	0.017004
	99%	0.032997	0.028092	0.030776	0.028361	0.028708
META	95%	0.027310	0.035232	0.030027	0.031416	0.024515
	99%	0.043077	0.051512	0.053435	0.046116	0.036081
JPM	95%	0.018300	0.020292	0.017892	0.014451	0.011423
	99%	0.037581	0.029149	0.034002	0.020888	0.024574

Table 3: Expected Shortfall (ES) results for the first 252-day window.

Ticker	Alpha	Historical	Gaussian	Student-t	Parametric	FHS
AAPL	95%	0.026807	0.024738	0.026424	0.024976	0.025729
	99%	0.040701	0.032404	0.039146	0.032712	0.038851
META	95%	0.035920	0.045215	0.045386	0.040429	0.031771
	99%	0.044618	0.059608	0.073637	0.053425	0.039434
JPM	95%	0.030295	0.025723	0.028770	0.018398	0.018972
	99%	0.047523	0.033553	0.050030	0.024089	0.028465

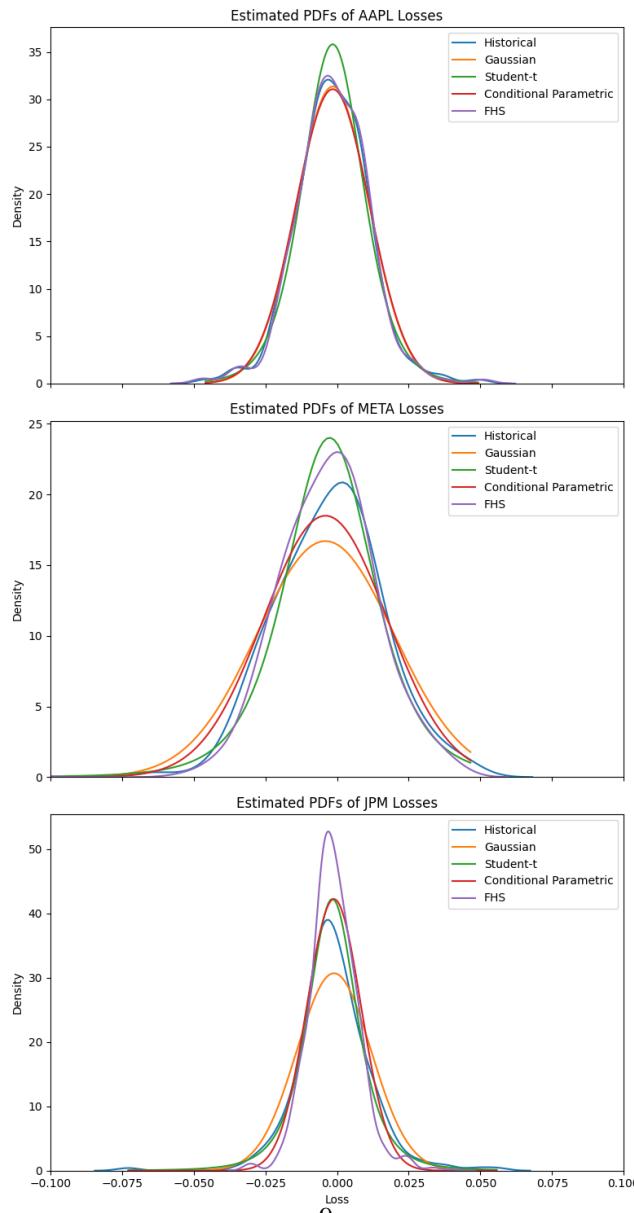


Figure 6: Estimated PDFs of losses for AAPL, META, and JPM using five approaches.

The PDF plots reveal deviations from Gaussianity across assets, with all three showing heavier tails and stronger kurtosis than the normal model can capture. Student-t and conditional parametric approaches provide better tail fitting, while FHS most closely reproduces the empirical distributional features. The differences across AAPL, META, and JPM further confirm that tail behaviour is asset-dependent.

3 Exercise 3

Backtesting VaR and ES. Use a rolling window of size $W = 252$ to produce 1-step-ahead forecasts of VaR/ES at the confidence levels $\alpha = 95\%$ and $\alpha = 99\%$ for each asset and each method from 2.a)–2.e) over the following *out-of-sample period*. Clearly describe your backtesting (i.e., rolling-window) procedure: at each time t , fit the models on $\{t - W + 1, \dots, t\}$, forecast at $t + 1$, advance by one day, and repeat until $t + 1 = T$, where T is the end of your sample. Plot the VaR forecasts along with the realized returns. Unless otherwise stated, use 95% as the primary confidence level.

- a) Implement and apply the Kupiec (1995) *Proportion-Of-Failures* (POF) test (unconditional coverage) and the Christoffersen (1998) conditional coverage test (joint test of correct exception rate and independence). Report test statistics, p -values, and discuss which models pass/fail. Interpret whether failures are due to biased coverage or clustered exceptions.
Refer to the referenced papers for further details, and go over the brief descriptions in the Appendix.
- b) *ES backtest:* Describe and apply the Acerbi and Székely (2014) Z_1 test for ES backtesting with $M = 1000$ simulations. State clearly the null and alternative hypotheses, emphasizing that the test is primarily one-sided (detecting underestimation of ES). Discuss which models pass/fail, report the corresponding p -values, and interpret the results.
Refer to the referenced paper for further details, and go over the brief descriptions in the Appendix.

Backtesting our VaR and ES estimation methods is fundamental for evaluating the adequacy of the risk model. It enables us to determine whether the predicted risk levels align with realized portfolio losses and to identify potential misspecification in the modeling of tail events. In this section, we perform the Kupiec Proportion-of-Failures test and the Christoffersen conditional coverage test for the Value-at-Risk, as well as the Acerbi and Székely Z_1 test for Expected Shortfall.

3.A Backtesting VaR and ES backtesting procedure described

For each confidence level (95%, 99%), ticker, and method, we computed one-step-ahead VaR and ES forecasts using a rolling window of 252 daily observations on the loss series. At each step, the model was re-estimated on the current window and used to generate the next-day forecast. This procedure produces a time series of VaR and ES estimates aligned with the out-of-sample period.

For 95% confidence we obtained the following VaR forecast for each ticker:

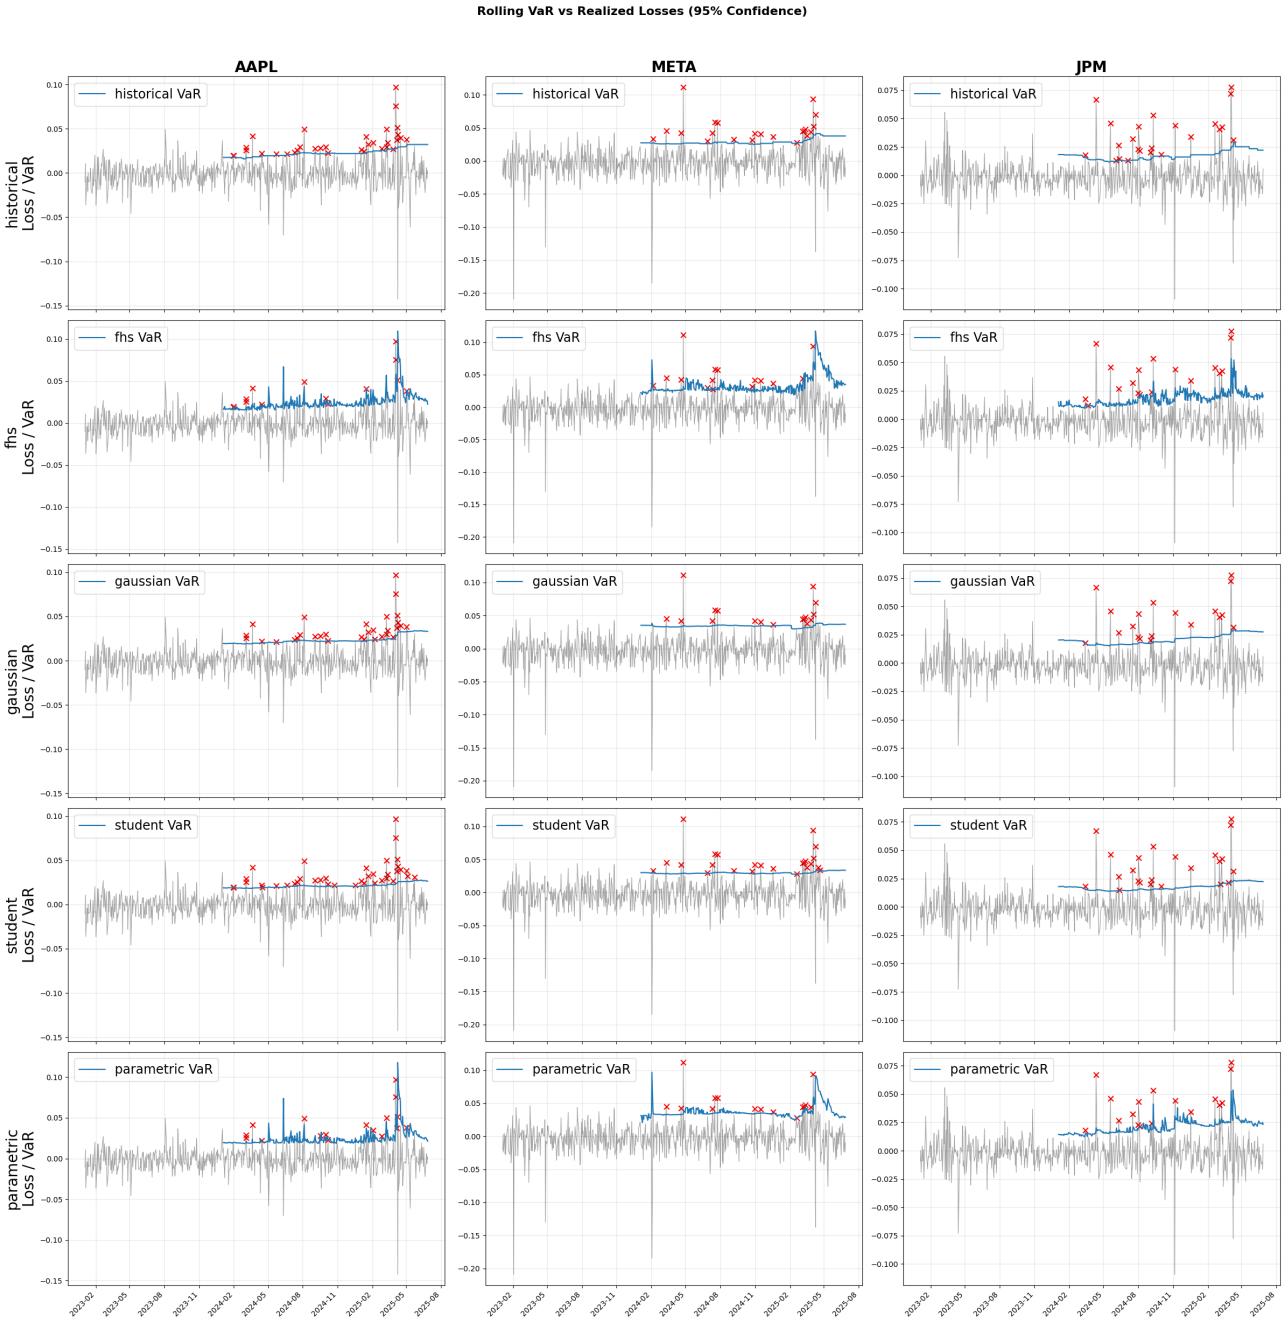


Figure 7: Time series of VaR forecasts versus realized losses.

Note that “Parametric” refers to the conditional parametric method for the remainder of this notebook.

From this plot, we observe that the non-parametric methods produce relatively flat VaR curves, since they do not explicitly react to changes in volatility as it is washed out by the estimation window. In contrast, the parametric approaches such as FHS and the conditional parametric model adjust the VaR threshold dynamically and capture local volatility spikes.

The Kupiec test examines whether the observed violation rate matches the model’s nominal level, while the Christoffersen test evaluates both correct coverage and independence of violations. Rejecting the null hypothesis in either test therefore signals that the VaR model is incorrect.

To perform both tests, we used each model’s one-step-ahead 95% VaR forecasts. For the Kupiec Proportion-of-Failures (POF) test, we evaluated the null hypothesis

$$\text{POF} \sim \chi_1^2,$$

and for the Christoffersen conditional coverage test we used

$$\text{POF} + \text{IND} \sim \chi^2_2.$$

Where IND is derived from the Christoffersen independence test. At the 5% significance level, we reject the null hypothesis when the p -value is below 0.05. We backtested our VaR across all assets using 5% significance test, and we obtained the following results:

Ticker	Method	Test statistic POF	Pvalue POF	Test statistic CD coverage	Pvalue CD coverage
AAPL	historical	11.077879	8.74e-04	14.674895	6.51e-04
	gaussian	7.455310	6.32e-03	9.212069	9.99e-03
	student	20.031440	7.62e-06	22.561494	1.26e-05
	parametric	5.389282	2.03e-02	5.425151	6.64e-02
	lhs	7.455310	6.32e-03	7.513060	2.34e-02
META	historical	1.073138	3.00e-01	1.273273	5.29e-01
	gaussian	0.131441	7.17e-01	0.186728	9.11e-01
	student	1.580101	2.09e-01	1.812080	4.04e-01
	parametric	0.743120	3.89e-01	1.962968	3.75e-01
	lhs	0.124872	7.24e-01	0.131576	9.36e-01
JPM	historical	1.580101	2.09e-01	4.263305	1.19e-01
	gaussian	0.014105	9.05e-01	2.720578	2.57e-01
	student	1.073138	3.00e-01	4.188907	1.23e-01
	parametric	0.014105	9.05e-01	2.720578	2.57e-01
	lhs	2.176202	1.40e-01	4.466326	1.07e-01

Across assets, all VaR models pass both the coverage and independence tests for META and JPM, indicating that their VaR and ES forecasts are statistically consistent with observed losses. In contrast, every model fails for AAPL, indicating an underestimation of the tail risk. We identify

whether failures come from biased coverage or clustered exceptions by checking which test rejects the null: Kupiec detects incorrect violation frequency, while Christoffersen's independence test reveals clustering, combining both indicates full model adequacy.

Since AAPL fails both the unconditional coverage test and the joint conditional coverage test for each methods, this confirms biased coverage, indicating that the models systematically underestimate its tail risk captured.

3.B Acerbi and Székely

Expected Shortfall is harder to backtest than VaR because ES measures the average severity of tail losses, not the frequency of exceptions. Since ES corresponds to the conditional mean of losses beyond the VaR, we cannot validate it simply by counting breaches, as it is done in VaR backtesting.

Acerbi and Székely use the following identity of ES for their backtesting procedure: $ES_{\alpha,t} = -\mathbb{E}[X_t | X_t + VaR_{\alpha,t} < 0]$. They define an indicator function that flags the days on which the VaR

is breached, and then analyse the distribution of realized losses on those breach days.

The core idea of the test is to compare the realized tail losses with the ES forecasts under a correctly specified model. To obtain the reference (null) distribution, they generate losses using the sample distribution of the null model. The realised Z-statistic is then compared with the simulated distribution. So the null hypothesis is that our model is correctly specified. The test is one-sided because

only underestimating Expected Shortfall is a real problem: if ES is too low, the model underestimates tail risk and leads to insufficient capital, whereas overestimating ES is simply conservative and not considered a failure. To conduct Acerbi and Székely test, for each method we created a sampler that

sample from the model distribution M samples for the Monte carlo simulation. By conducting Acerbi

and Székely test we obtained the following results :

Ticker	Method	Z1 stat	Pvalue	ES backtest
AAPL	historical	0.017973	3.64e-01	
	gaussian	0.166928	0.00e+00	
	student	-0.027741	5.52e-01	
	parametric	0.003810	4.54e-01	
	fhs	-0.114048	9.74e-01	
META	historical	0.057958	2.63e-01	
	gaussian	0.241137	0.00e+00	
	student	0.043681	2.70e-01	
	parametric	0.104055	2.00e-02	
	fhs	0.082961	2.40e-01	
JPM	historical	0.104375	1.61e-01	
	gaussian	0.534150	0.00e+00	
	student	0.299156	3.50e-02	
	parametric	0.343861	0.00e+00	
	fhs	0.090146	3.01e-01	

For AAPL, only the Gaussian ES model fails the Z_1 test, with a p-value of 0, indicating that the Gaussian assumption systematically underestimates tail severity. All other models exhibit p-values above 5%, indicating that their ES estimates are statistically compatible with the realized tail losses.

For META, the Gaussian model fails, with a very small p-value of 0, indicating that the Gaussian assumption systematically underestimates tail severity. The conditional parametric model also fails with a p-value of $2 \cdot 10^{-2}$. All other models exhibit p-values above 5% showing no statistical evidence of ES misspecification.

For JPM, the Gaussian model has p-value 0 and the parametric model 0 and the student model $3.50 \cdot 10^{-2}$ both fail the Z_1 test. These models underpredict the magnitude of extreme losses. The historical and FHS models pass, indicating that their ES forecasts align with realized tail behavior.

Across the three assets, the results show that ES model performance is strongly asset dependent, reflecting differences in tail behaviour across return distributions. A common pattern is that the Gaussian model systematically fails for all assets, confirming that its thin-tailed distributional assumption leads to persistent underestimation of extreme losses. However, the fact that different models fail for different assets highlights that the appropriateness of an ES model depends critically on asset-specific tail dynamics. Notably, the filtered historical simulation performs well across all cases, suggesting that its semi-parametric structure offers robust flexibility in practice.

4 Exercise 4

Copula fitting. On the single estimation window of length $W = 252$ (i.e., same as in question 2), model cross-asset *dependence* across AAPL, META, and JPM using copulas.

- a) For each asset i , let $U_{t,i} = \frac{\text{rank}(R_{t,i})}{W+1}$, $t = 1, \dots, W$ be the empirical quantiles, sometimes known as pseudo-observations. Plot dependence between asset pairs both in the raw returns space (i.e., via the scatter plot of $(R_{t,i}, R_{t,j})$) and the pseudo-observations space (i.e., via the scatter plot of $(U_{t,i}, U_{t,j})$). Comment on tail dependence, symmetry/asymmetry, and how these features motivate the choice of copula.
- b) Fit a Gaussian copula and a Student- t copula and explain the fitting procedure. Report the obtained parameter estimates.
- c) Explain how to simulate from a copula, and use it to generate synthetic data. Thereby, sample T points from the fitted copulas, with T being the total number of observations in the dataset. Using the inverse empirical cdf of each asset, transform the simulated uniform samples into simulated returns. Visually compare the original and simulated returns, and describe what you observe.

Copulas and Their Purpose

In multivariate modelling, it is often useful to separate the marginal behaviour of each variable from the dependence structure that links them. A copula provides exactly this decomposition. By Sklar's theorem, any multivariate distribution F with marginals F_1, \dots, F_d can be written as

$$F(x_1, \dots, x_d) = C(F_1(x_1), \dots, F_d(x_d)),$$

where C is a *copula*: a multivariate distribution on $[0, 1]^d$ with uniform marginals.

This representation allows us to model the *marginal distributions* F_i independently from *dependence structure* encoded by the copula C .

In financial applications, this flexibility is valuable because returns often exhibit heavy tails and skewness, while dependence may be nonlinear or exhibit tail co-movements. Gaussian and Student- t copulas are widely used examples: the Gaussian copula captures linear dependence through a correlation matrix, whereas the Student- t copula additionally allows for joint extreme events via its degrees of freedom.

In this exercise, we use copulas to isolate the dependence between AAPL, META, and JPM, fit parametric copula models to this structure, and generate simulated return series that preserve both the empirical marginals and the estimated dependence.

4.A Pseudo-observations

We begin by extracting the first estimation window ($W = 252$ observations) of daily log-returns for AAPL, META, and JPM. For each asset i and time t , we construct pseudo-observations using

$$U_{t,i} = \frac{\text{rank}(R_{t,i})}{W + 1},$$

where the ranking is applied column-wise. This transformation produces approximately uniform variables on $(0, 1)$ while preserving only the dependence structure.

Figure 8 shows the raw return pairs alongside their pseudo-observation counterparts. AAPL-META exhibits clear positive dependence in both spaces, while pairs involving JPM show weaker association.

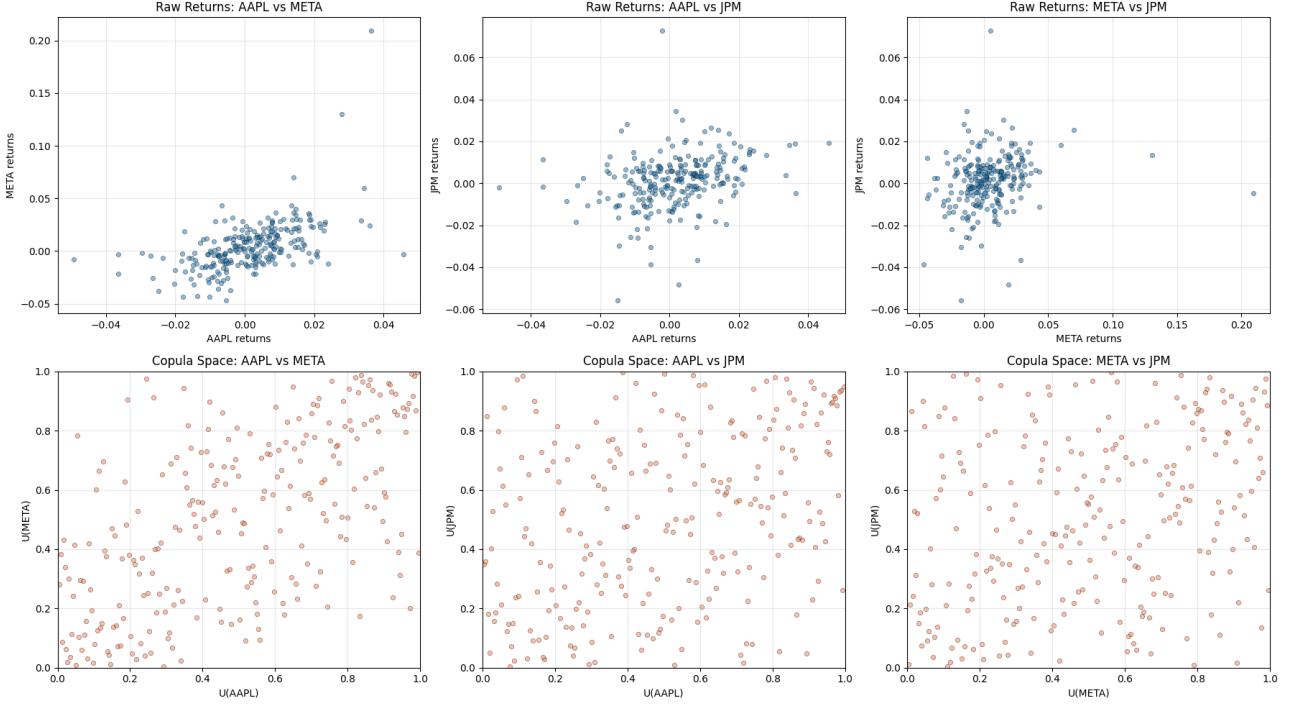


Figure 8: Raw returns (top row) and pseudo-observations (bottom row) for each asset pair.

Raw returns The **AAPL-META** pair exhibits a clear upward-sloping cloud, indicating a positive dependence and frequent joint large moves. In contrast, **AAPL-JPM** and **META-JPM** show more diffuse scatter patterns, reflecting weaker dependence between tech firms and a bank. Across all pairs, extreme returns are more dispersed than expected under normality, suggesting heavy-tailed behavior and occasional simultaneous large shocks.

Copula space In copula space, **AAPL-META** displays a marked diagonal structure, confirming a positive dependence. The clustering of observations near the corners $(0, 0)$ and $(1, 1)$ further reveals tail dependence, meaning that extreme moves tend to occur together. By comparison, **AAPL-JPM** and **META-JPM** resemble an almost uniform cloud over $[0, 1]^2$, confirming weaker dependence and only marginal tail co-movement.

Implications for copula choice Given the evident heavy tails and clear tail dependence in **AAPL-META**, a Gaussian copula, though adequate for symmetric dependence, cannot capture the probability mass in the joint extremes. The Student- t copula, however, explicitly models tail dependence and is therefore the appropriate copula for this portfolio. Its structure aligns with the empirical behavior of the assets, particularly the strong, tail-driven co-movements between the two tech stocks. Consequently, the Student- t copula provides a more realistic representation of joint downside risk than its Gaussian counterpart.

4.B Fit the copulas

We fit a Gaussian copula and a Student- t copula to the pseudo-observations via maximum likelihood. Given pseudo-observations $U = (U_{t,1}, U_{t,2}, U_{t,3})$, the log-likelihood of a copula C_θ with parameter vector θ is

$$\ell(\theta) = \sum_{t=1}^W \log c_\theta(U_{t,1}, U_{t,2}, U_{t,3}),$$

where c_θ is the copula density. The fitting procedure consists of maximizing $\ell(\theta)$ with respect to the dependence parameters (the correlation matrix, and additionally ν for the Student- t copula). Since the marginals are already transformed into uniforms, the estimation focuses purely on the dependence structure.

Gaussian copula. The Gaussian copula is parametrized solely by its 3×3 correlation matrix. The estimated correlation matrix and log-likelihood of ≈ 69 are obtained from the fitted model. With $d = 3$ assets, the Gaussian copula has 3 free parameters.

Student- t copula. The Student- t copula has the same correlation matrix as the Gaussian copula but includes an additional degrees-of-freedom parameter ν (making it a total of 4 parameters), which controls tail dependence. In our fit, we estimated $\nu \approx 13$, indicating moderately heavy joint tails. The Student- t copula achieves a slightly higher log-likelihood of ≈ 70 , consistent with the mild tail clustering observed in the data.

Correlation matrix.

$$\Sigma_{\text{Gaussian}} = \begin{matrix} & \text{AAPL} & \text{META} & \text{JPM} \\ \text{AAPL} & 1.0000 & 0.2870 & 0.3428 \\ \text{META} & 0.2870 & 1.0000 & 0.5979 \\ \text{JPM} & 0.3428 & 0.5979 & 1.0000 \end{matrix}$$

$$\Sigma_{\text{Student-}t} = \begin{matrix} & \text{AAPL} & \text{META} & \text{JPM} \\ \text{AAPL} & 1.0000 & 0.3021 & 0.3494 \\ \text{META} & 0.3021 & 1.0000 & 0.5995 \\ \text{JPM} & 0.3494 & 0.5995 & 1.0000 \end{matrix}$$

4.C Simulate from the copulas

Using the fitted copulas, we generate T synthetic observations for each asset. The simulation relies on the fundamental copula identity

$$F(x_1, \dots, x_d) = C_\theta(F_1(x_1), \dots, F_d(x_d)),$$

which implies that any joint sample can be obtained by first simulating the dependence (through the copula) and then restoring the marginal distributions through inverse CDFs.

Step 1: simulate uniforms. For each $t = 1, \dots, T$, we draw a vector of uniforms

$$U^{(t)} = (U_1^{(t)}, U_2^{(t)}, U_3^{(t)}) \sim C_\theta,$$

where C_θ is the fitted Gaussian or Student- t copula. These uniforms encode *only the dependence structure*: individually, $U_i^{(t)} \sim \mathcal{U}(0, 1)$, but jointly they reproduce the correlations or tail dependence implied by the fitted copula.

Step 2: inverse empirical CDFs. For each asset i , we construct its empirical inverse CDF \hat{F}_i^{-1} from the observed returns by sorting the data and performing monotone interpolation. Each simulated uniform is then transformed as

$$X_i^{(t)} = \hat{F}_i^{-1}(U_i^{(t)}).$$

This step ensures that the simulated values $X_i^{(t)}$ follow the same marginal distribution as the historical returns of asset i , regardless of the copula chosen. In particular, this approach reproduces skewness, kurtosis, and heavy tails present in the empirical data without assuming any parametric marginals.

Step 3: comparison. Figure 9 compares the simulated and original marginal distributions. Both copulas reproduce the empirical marginals very closely.

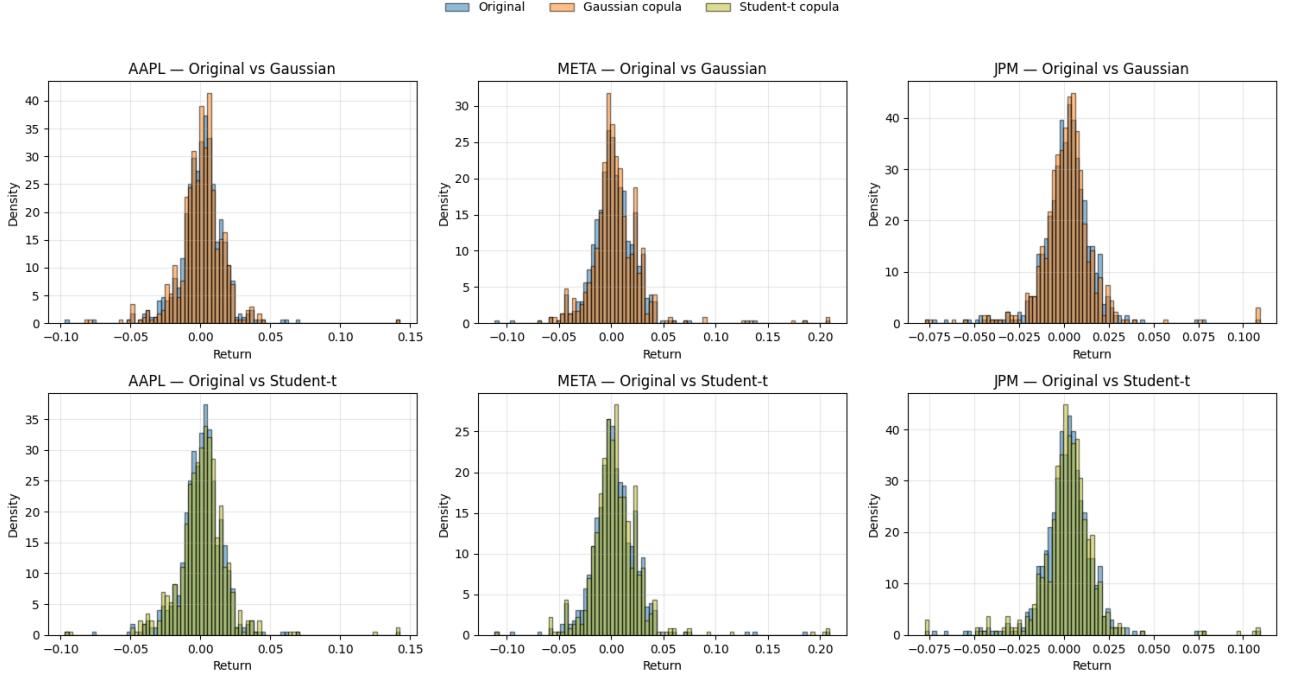


Figure 9: Comparison of original vs. copula-simulated marginal distributions (Gaussian and Student- t).

Figure 10 compares the pairwise dependence in the original and simulated datasets. The Gaussian copula captures the overall dependence well, while the Student- t copula produces slightly stronger joint extremes, though the effect is moderate due to the fitted degrees of freedom relatively moderate.

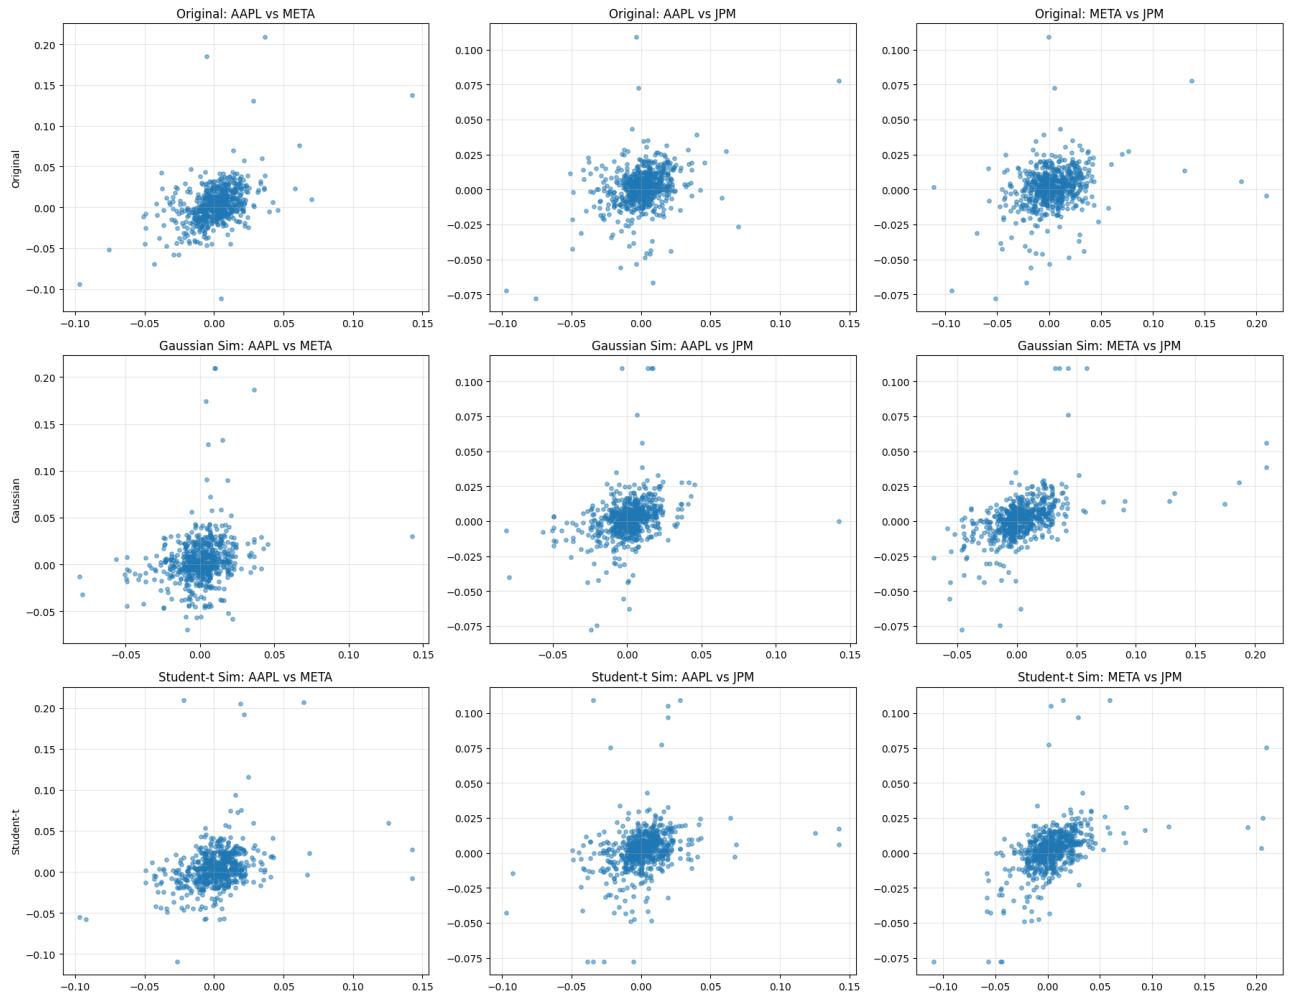


Figure 10: Dependence structure of original returns vs. Gaussian and Student- t copula simulations.

Figure 11 compares the evolution of the original return series with those generated by the Gaussian and Student- t copulas. Both copulas reproduce the short-term fluctuations around zero observed in the empirical returns. The Gaussian copula yields smoother trajectories with fewer pronounced peaks, reflecting its tendency to underestimate tail events. In contrast, the Student- t copula produces occasional large spikes and clusters of heightened volatility, more closely resembling the empirical behavior of financial returns and capturing heavy-tailed dynamics. This difference is particularly visible for META, where the Student- t simulations exhibit sequences of large movements indicative of volatility clustering.

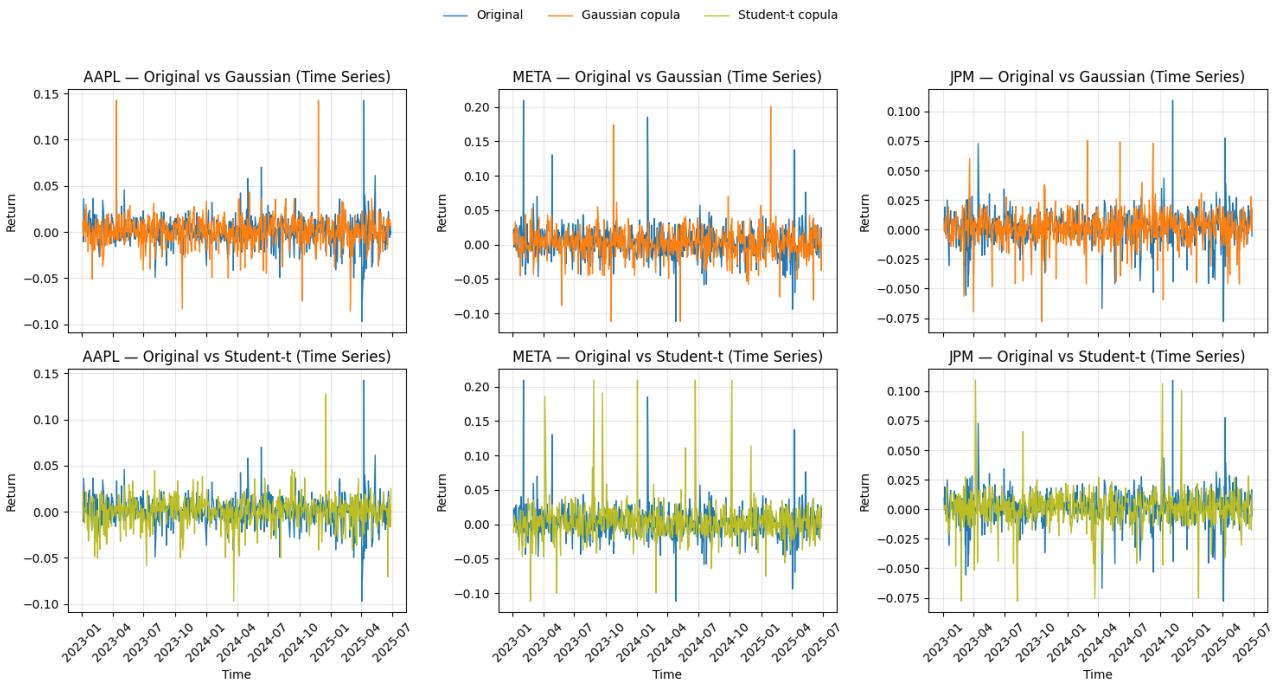


Figure 11: Time-series comparison of simulated returns vs. original returns

Overall, both copulas generate realistic synthetic returns that successfully replicate the empirical marginal distributions and preserve the dependence structure estimated in the first window. However, the Student-*t* copula offers a more faithful reproduction of extreme events and volatility clustering, making it better suited for scenarios involving risk analysis or stress testing, where accurate modeling of tail behavior is crucial.

5 Exercise 5

Backtesting portfolio VaR and ES. Construct an *equal-weighted* portfolio of AAPL, META, and JPM. Use the same univariate methods as in question 2 and include also the copula-based approach from question 4. Estimate 1-day ahead VaR/ES at $\alpha = 95\%$ and $\alpha = 99\%$, then *backtest and compare* them using rolling-window forecasts as in question 3. Compare portfolio VaR/ES backtests across univariate and copula-based approaches. Assess whether explicit dependence modeling improves accuracy, especially in the tails. Discuss potential reasons.

- Univariate models:* use the same backtesting procedure as in question 3, but now on the portfolio returns.
- Copulas:* at each time t , fit the copulas as in 4.b) on the last W days of portfolio component returns. Simulate $N = 1000$ returns from each fitted copula, then estimate VaR/ES from the simulated portfolio returns.
- Backtesting and comparison:* present a compact comparison of backtests. Does dependence modeling improve results? Why/why not?

5.A Univariate models

We construct an equal-weighted portfolio of the three assets,

$$R_t^{\text{port}} = \frac{1}{3}R_t^{\text{AAPL}} + \frac{1}{3}R_t^{\text{META}} + \frac{1}{3}R_t^{\text{JPM}},$$

and compute its daily log-returns over the full sample. The resulting series has mean 0.00165 and standard deviation 0.0142, with extremes ranging from -8.8% to $+11.9\%$.

We evaluate several univariate risk models on the portfolio loss series $L_t^{\text{port}} = -R_t^{\text{port}}$, using the same rolling-window setup as in Exercise 3 ($W = 252$ days). For each method and each confidence level $1 - \alpha \in \{0.95, 0.99\}$, we compute one-step-ahead forecasts of VaR_α and ES_α :

We evaluate five univariate risk models: Historical Simulation (HS), a Gaussian parametric model, a Student- t parametric model, a conditional parametric model (GARCH with Gaussian innovations), and Filtered Historical Simulation (FHS).

The resulting VaR and ES sequences are backtested using the Kupiec Proportion-of-Failures (POF) test, the Christoffersen independence test (serial independence of violations) which we use to implement the Conditional Coverage tests. Finally, we use the Z1 ES backtest based on simulated p-values.

Method	95%				99%							
	VaR	POF	VaR	CC	ES	Z1	VaR	POF	VaR	CC	ES	Z1
Historical	✓		✓		✗		✗		✗		✓	
Gaussian	✓		✓		✗		✗		✗		✗	
Student- t	✓		✓		✓		✗		✗		✓	
Parametric (GARCH)	✓		✓		✗		✓		✓		✓	
FHS	✓		✓		✓		✓		✓		✓	

Table 4: Combined backtesting results for portfolio VaR and ES (univariate models). A checkmark denotes failure to reject the null hypothesis at the 5% significance level.

Confidence Level	Method	POF p-value	CC p-value	ES Z1 stat	ES p-value
95%	historical	0.1436	0.2264	0.1641	0.0190
	gaussian	0.1436	0.2264	0.3136	0.0000
	student	0.0587	0.1326	0.1785	0.0630
	parametric	0.3061	0.3064	0.1465	0.0020
	fhs	0.5673	0.6757	0.0598	0.2170
99%	historical	0.0067	0.0141	0.0576	0.2300
	gaussian	0.0002	0.0006	0.2675	0.0050
	student	0.0021	0.0056	-0.0189	0.3870
	parametric	0.0523	0.1280	0.1151	0.0760
	fhs	0.1263	0.2722	-0.0336	0.8100

Table 5: Backtesting results for univariate portfolio models. POF = Kupiec proportion of failures test; CC = Christoffersen conditional coverage test; ES = Expected Shortfall Z1 backtest.

At the 95% confidence level, all models successfully pass the VaR backtests. However, at the more stringent 99% level, the pure historical and simple parametric specifications (Gaussian and Student- t) fail both Kupiec and Christoffersen tests, confirming their inability to accurately capture tail risk. In contrast, dynamic and volatility sensitive approaches, namely the conditional GARCH model and Filtered Historical Simulation (FHS), display markedly superior performance, passing both VaR diagnostics across confidence levels.

Expected Shortfall forecasts, evaluated using the Acerbi-Szekely Z1 backtest, exhibit a similar pattern. While Gaussian-based ES forecasts systematically fail, heavy-tailed or conditional models (Student- t , GARCH, and especially FHS) generally produce statistically consistent ES estimates, indicating a more reliable characterization of downside risk.

The backtesting results demonstrate that static or distributionally simplistic models are inadequate for high-confidence risk estimation on this portfolio. Models that account for time-varying volatility or tail dependence, in particular GARCH and FHS, provide materially more robust risk forecasts, validating their use in environments where accurate extreme-loss modeling is essential.

5.B Copulas

We next estimate portfolio VaR and ES using multivariate dependence models. At each time $t \geq W$ (with $W = 252$), we fit a Gaussian and a Student- t copula to the pseudo-observations of AAPL, META, and JPM over the most recent window. The fitted copulas capture the cross-asset dependence structure for that period.

Intuition. The copula approach generates many plausible future market scenarios by combining the dependence structure between assets with their empirical marginal distributions. The fitted copula produces correlated uniform random variables that encode how the three assets tend to co-move, including the possibility of joint extremes (in the Student- t case). Each uniform component is then mapped through the empirical inverse CDF of the corresponding asset, ensuring that the simulated returns follow the same marginal behavior as in the historical window. This yields realistic joint return scenarios from which portfolio losses can be computed.

Simulation and risk estimation. For each copula, we simulate $N_{\text{sim}} = 1000$ joint scenarios

$$U^{(m)} \sim C_\theta, \quad m = 1, \dots, N_{\text{sim}},$$

and transform the components via

$$X_i^{(m)} = \hat{F}_i^{-1}(U_i^{(m)}).$$

This produces simulated returns with the correct marginals and copula-based dependence. Portfolio returns are then computed using the equal-weight vector $w = (\frac{1}{3}, \frac{1}{3}, \frac{1}{3})$, and simulated losses

$$L^{(m)} = -w^\top X^{(m)}$$

yield Monte Carlo estimates of VaR and ES at confidence levels $1 - \alpha \in \{0.95, 0.99\}$. Results are discussed in the next subsection 5.C.

5.C Backtesting and Comparison

We now compare the out-of-sample performance of all portfolio risk models: both univariate and copula-based using the VaR and ES forecasts obtained in the previous subsections. The evaluation covers the final 370 days of the sample, corresponding to the rolling-window setup with $W = 252$. The main results are shown in Figure 12.

Method	VaR at 95%		VaR at 99%		ES	
	POF	CC	POF	CC	95%	99%
Gaussian Copula	✓	✓	✗	✗	✗	✓
Student- t Copula	✓	✓	✓	✓	✓	✓

Table 6: Backtesting results for copula-based portfolio models. A checkmark denotes failure to reject the null hypothesis at 5%.

Confidence	Approach	Method	Violations	Total	Violation Rate	Expected Rate	POF Stat	POF p-value	Christoffersen Stat	Christoffersen p-value	Z1 Stat	Z1 p-value
0	95% Univariate	historical	25	371	0.0674	0.0500	2.1391	0.1436	2.9707	0.2264	0.1641	0.0130
1	95% Univariate	gaussian	25	371	0.0674	0.0500	2.1391	0.1436	2.9707	0.2264	0.3136	0.0000
2	95% Univariate	student	27	371	0.0728	0.0500	3.5740	0.0587	4.0411	0.1326	0.1785	0.0600
3	95% Univariate	parametric	23	371	0.0620	0.0500	1.0475	0.3061	2.3654	0.3064	0.1513	0.0010
4	95% Univariate	fhs	19	371	0.0512	0.0500	0.0114	0.9150	0.8387	0.6575	0.1120	0.0650
5	95% Copula	Gaussian Copula	26	370	0.0703	0.0500	2.8581	0.0909	3.4867	0.1749	0.1561	0.0310
6	95% Copula	Student-t Copula	25	370	0.0676	0.0500	2.1762	0.1402	3.0007	0.2230	0.1597	0.0530
7	99% Univariate	historical	10	371	0.0270	0.0100	7.3594	0.0067	8.5207	0.0141	0.0576	0.2040
8	99% Univariate	gaussian	13	371	0.0350	0.0100	14.2588	0.0002	14.7410	0.0006	0.2675	0.0040
9	99% Univariate	student	11	371	0.0296	0.0100	9.4767	0.0021	10.3690	0.0056	-0.0189	0.4360
10	99% Univariate	parametric	8	371	0.0216	0.0100	3.7649	0.0523	4.1108	0.1280	0.1215	0.0590
11	99% Univariate	fhs	6	371	0.0162	0.0100	1.2030	0.2727	1.3976	0.4972	0.0714	0.3270
12	99% Copula	Gaussian Copula	9	370	0.0243	0.0100	5.4771	0.0193	6.9544	0.0309	0.1684	0.0740
13	99% Copula	Student-t Copula	8	370	0.0216	0.0100	3.7884	0.0516	5.6469	0.0594	0.0556	0.3080

Figure 12: Summary table of VaR and ES backtesting statistics for all univariate and copula-based models at the 95% and 99% confidence levels.

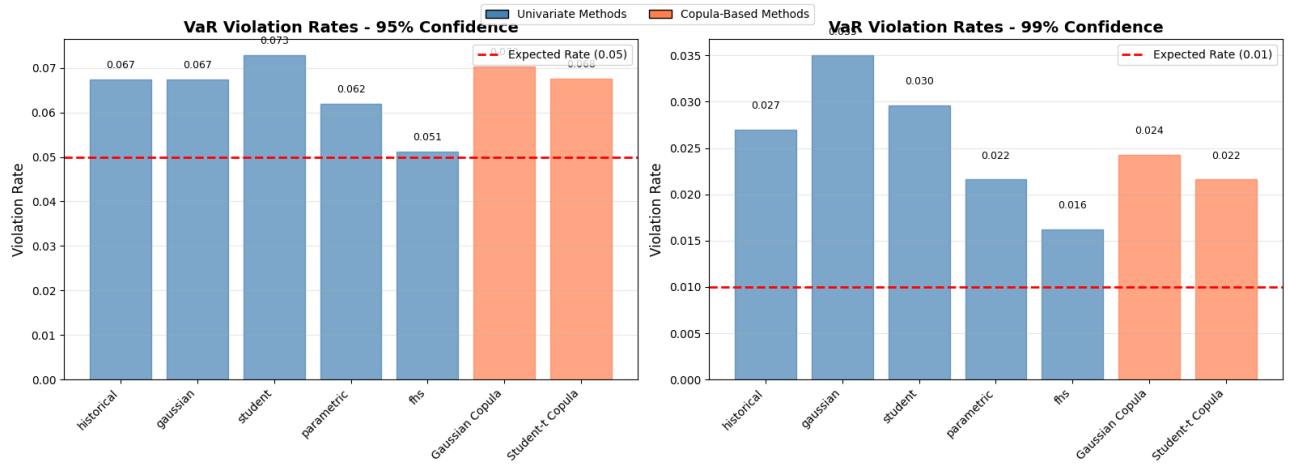


Figure 13: Violation rates for univariate and copula-based VaR forecasts at 95% and 99% confidence.

Interpretation of Copula-Based Backtesting Results. At the 95% confidence level, the Figure 13 shows that both Gaussian and Student- t copulas exhibit violation rates that exceed the theoretical benchmark, performing similarly to or slightly worse than several univariate methods. At the 99% level, copula models show a modest improvement relative to some univariate approaches; however, their violation rates remain above the expected threshold, indicating persistent underestimation of tail risk. In contrast, univariate models that account for time-varying volatility, in particular the Filtered Historical Simulation (FHS) and parametric GARCH, achieve violation rates closest to the theoretical expectations across both confidence levels. This suggests that, for our particular portfolio and the period under study, modeling temporal dependence and volatility clustering is more impactful for accurate risk measurement than modeling cross-sectional dependence alone.

Portfolio Rolling VaR vs Realized Losses Methods x Confidence Levels

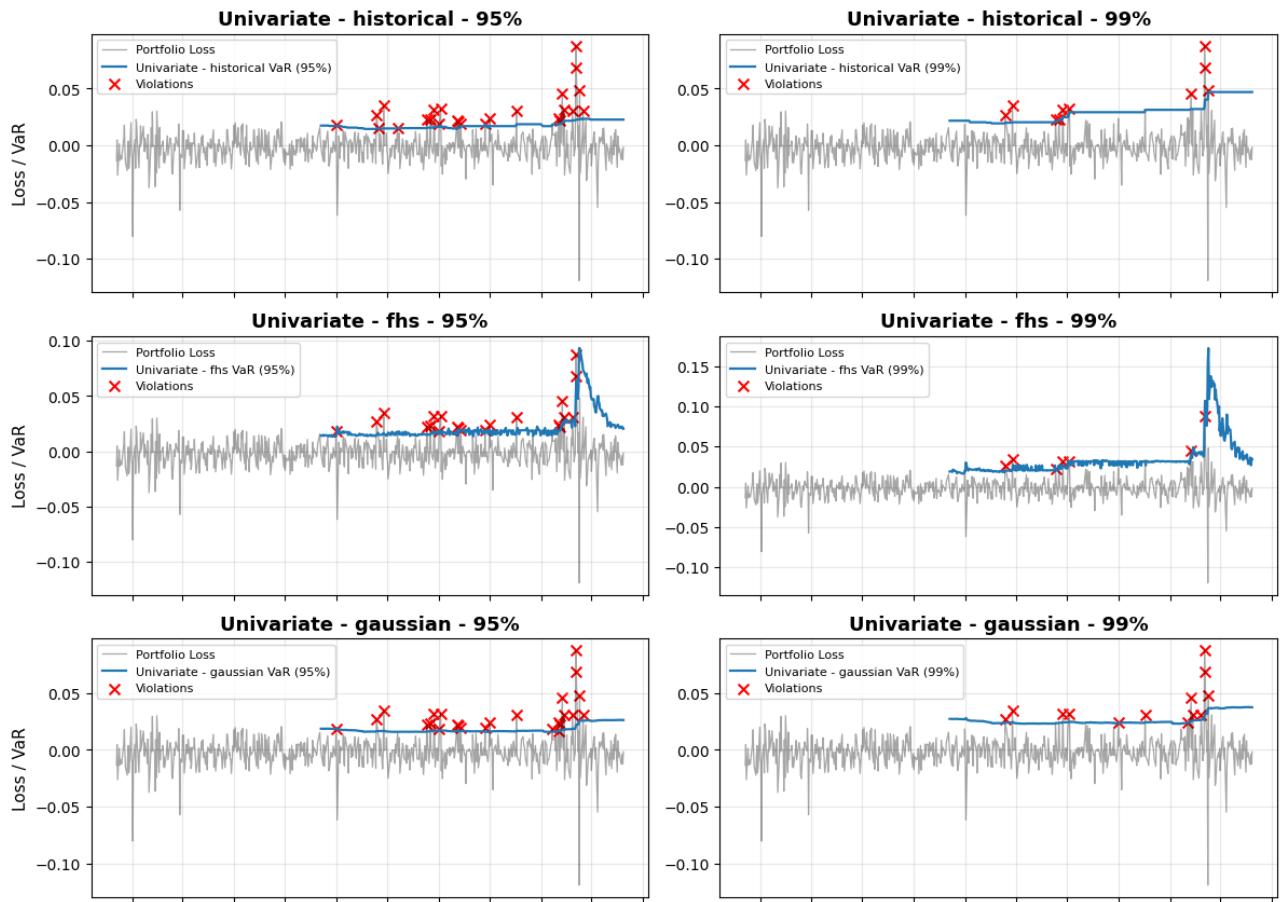


Figure 14: Time series of portfolio VaR forecasts (Part 1).

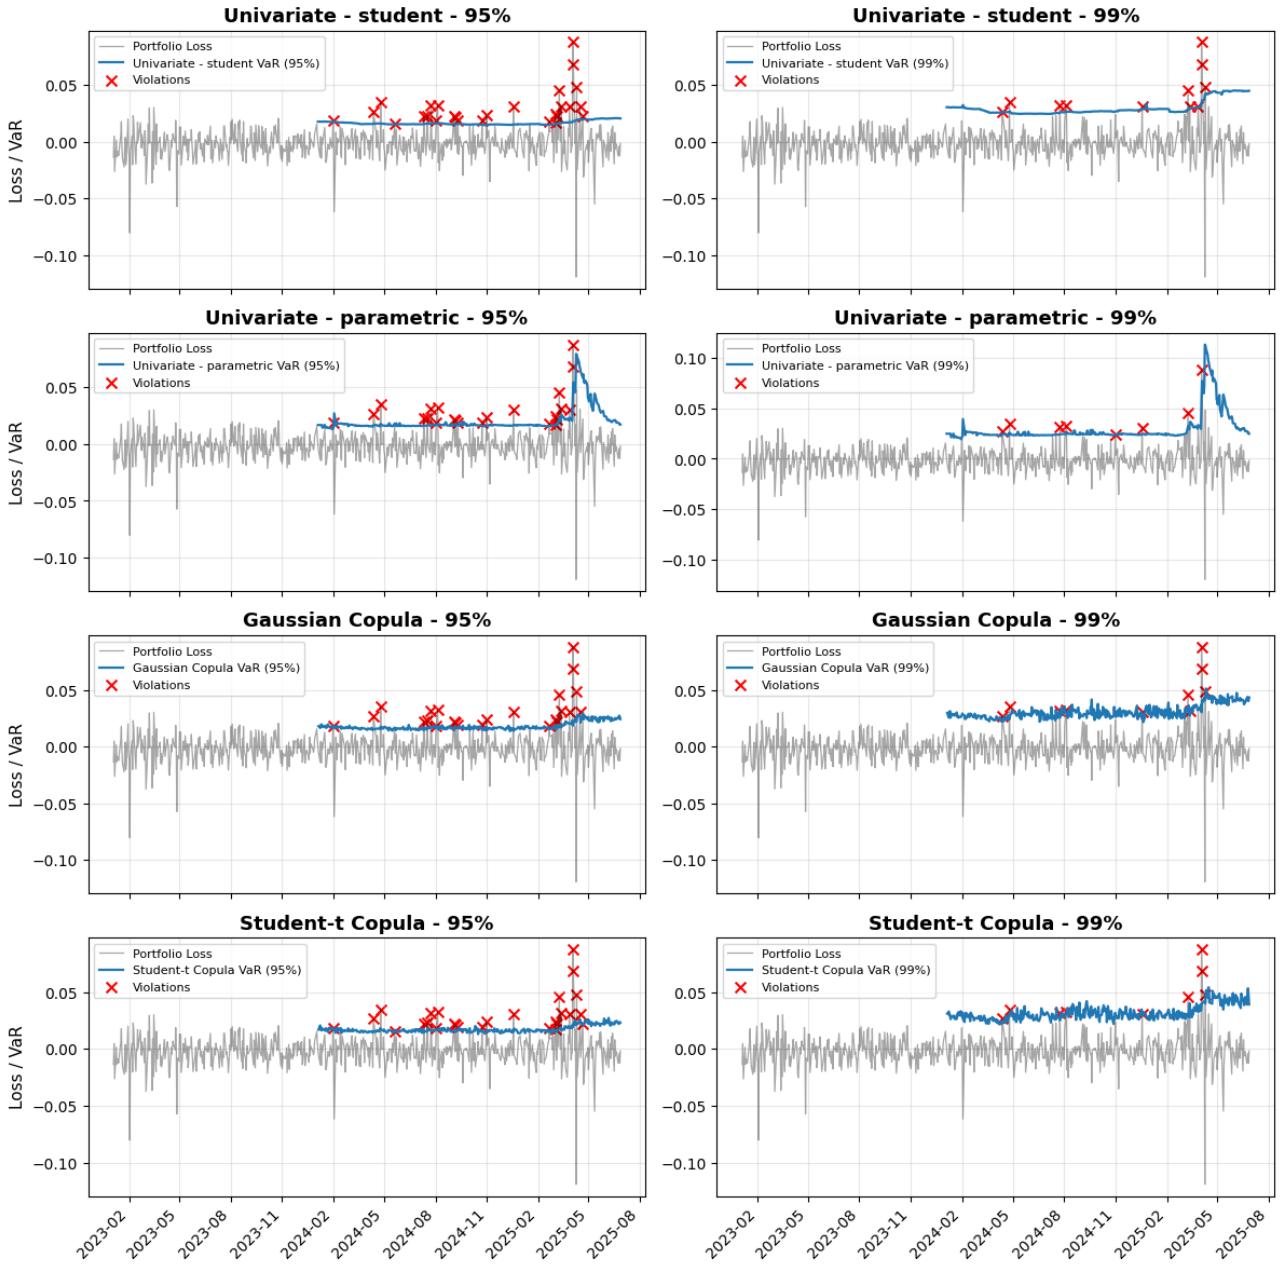


Figure 15: Time series of portfolio VaR forecasts (Part 2).

Figures 14 and 15 present the rolling VaR estimates alongside realized portfolio losses for all considered models. A visual inspection reveals that copula-based approaches do not systematically outperform univariate models such as FHS or GARCH-type specifications. Although copulas provide a more flexible framework to model cross-sectional dependence, their VaR curves often remain too smooth and fail to adjust rapidly during periods of heightened volatility. As a result, copula-based VaR estimates are frequently exceeded during market stress, leading to a substantial number of violations.

In contrast, univariate methods that explicitly incorporate time-varying volatility, particularly FHS, tend to align more closely with the observed loss dynamics. These models react more promptly to volatility bursts and adjust their risk estimates accordingly, which explains their lower violation frequencies. This behavior is especially visible during the sharp increase in losses towards the end of the sample: while FHS-based VaR rapidly escalates to reflect the changing risk environment, copula-based VaR remains comparatively muted and is repeatedly breached.

Portfolio Rolling ES vs Realized Losses Methods x Confidence Levels

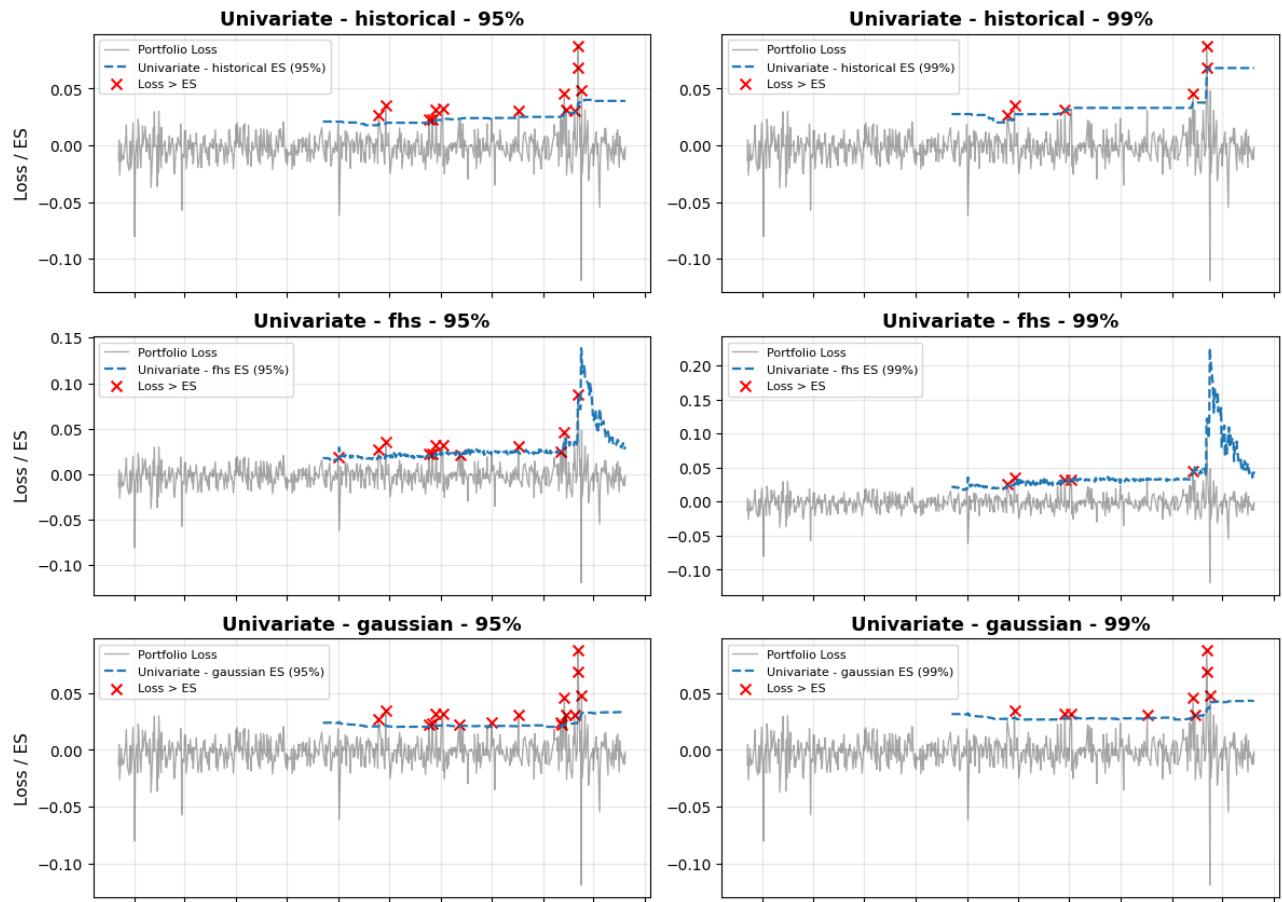


Figure 16: Time series of portfolio ES forecasts (Part 1).

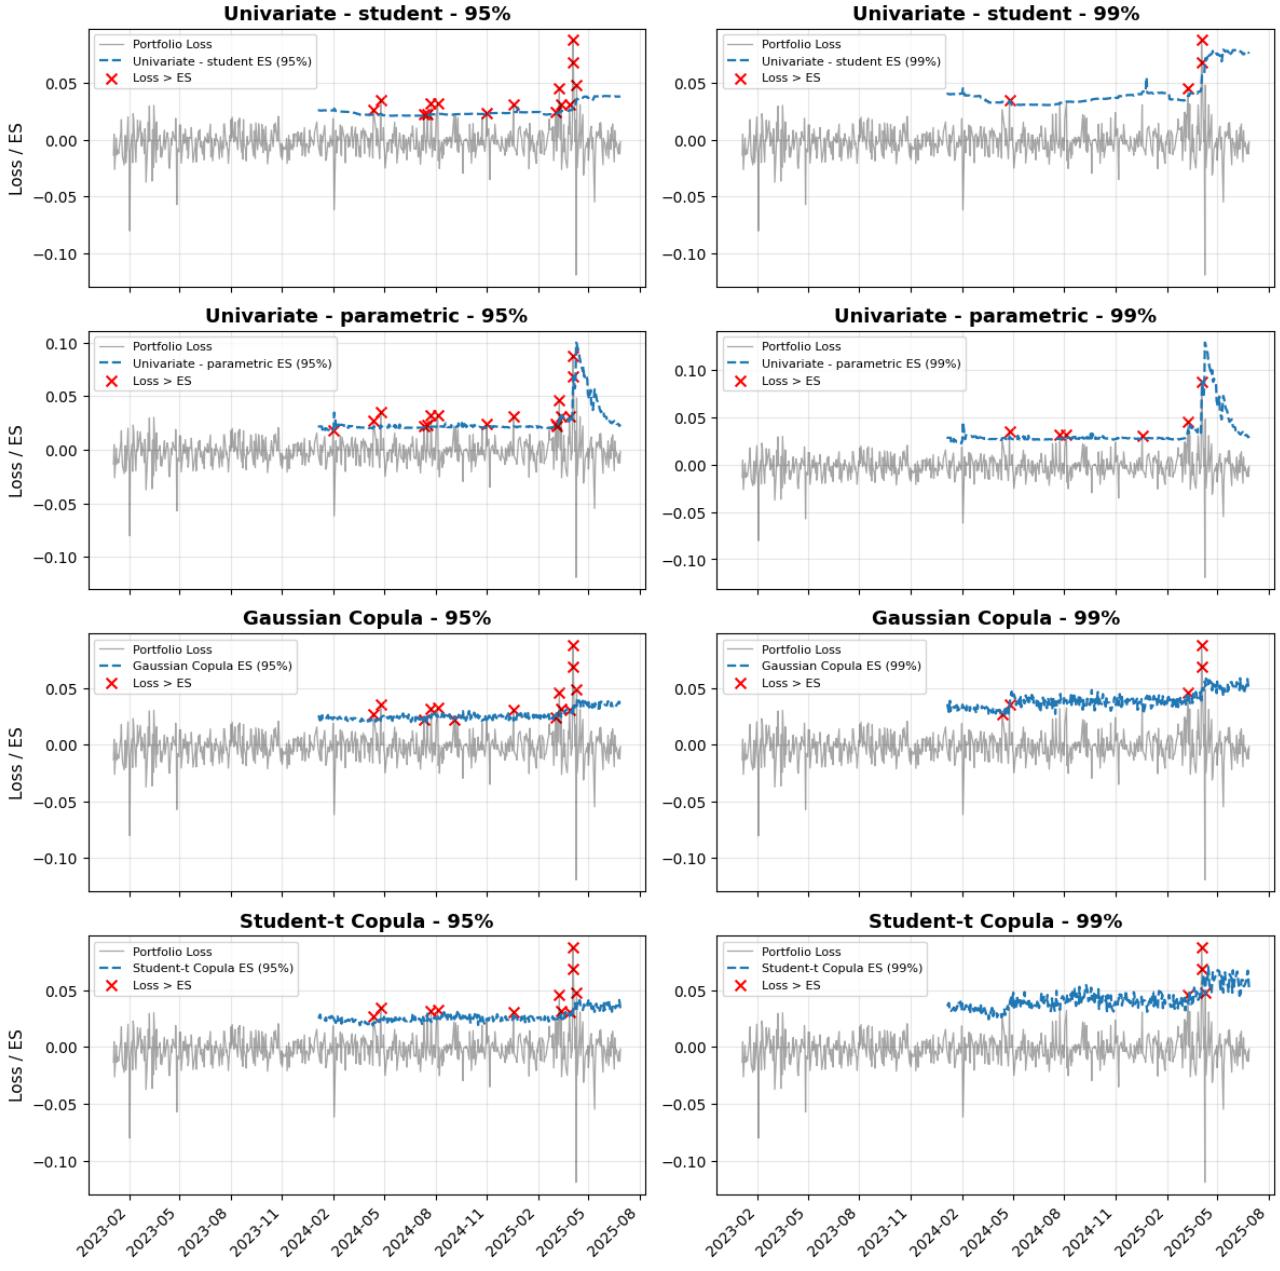


Figure 17: Time series of portfolio ES forecasts (Part 2).

Figures 16 and 17 display the rolling Expected Shortfall (ES) forecasts together with realized portfolio losses for all risk models. Similar to the VaR analysis, the copula-based approaches do not demonstrate a clear performance advantage over univariate methods. Although ES is a coherent risk measure and theoretically more sensitive to tail events, the copula-based ES estimates remain relatively stable throughout the sample and fail to escalate during periods of heightened market stress. Consequently, several realized losses exceed the copula-based ES thresholds, indicating that these models tend to underestimate tail risk.

By contrast, univariate models that incorporate time-varying volatility, particularly the Filtered Historical Simulation (FHS) and, to a lesser extent, parametric GARCH-type, provide more reactive ES estimates. These methods adjust rapidly to volatility bursts, producing higher ES values when market conditions deteriorate. This dynamic behavior results in fewer ES breaches, especially in the latter part of the sample, where the portfolio experiences the most extreme losses. The responsiveness of FHS to evolving volatility patterns highlights the importance of modeling temporal dependence and volatility clustering when forecasting downside risk.

Conclusion. The objective of this analysis was to assess whether dependence modeling improves portfolio risk estimation. While copula-based approaches explicitly capture the cross-sectional dependence structure between assets, the empirical results do not provide clear evidence of superior performance compared to univariate models.

An explanation for these findings lies in the characteristics of the underlying portfolio. The three assets considered do not exhibit pronounced dependence, which limits the potential gains from employing copula-based models. In such a context, explicitly modeling cross-sectional dependence offers little improvement. Moreover, copulas do not account for volatility clustering or sudden volatility bursts if we give them simple Gaussian or Student- t marginals, unlike FHS and GARCH. Since extreme losses during the sample period appear to be primarily driven by volatility dynamics rather than dependence structures, correctly modeling time-varying volatility proves more important than capturing marginal dependence. Consequently, the relative underperformance of copula-based methods is consistent with the underlying data-generating process and the nature of the portfolio risk.

Therefore, we cannot conclude that dependence modeling via copulas systematically improves the results. Copula-based methods do not outperform well-specified univariate models such as FHS, and their benefits become visible only at extreme confidence levels without consistently translating into better backtesting outcomes. While copulas generally perform better than some simpler univariate approaches, this incremental improvement remains insufficient to claim superiority over methods that explicitly model volatility dynamics. As a result, the evidence does not support the claim that copulas provide superior VaR or ES estimates in this setting.

A Code Appendix

```
# %% [markdown]
# # Project 1: Market Risk - VaR, ES, and Copulas

# %% [markdown]
# **Names of all group members:**
# - William Jallot (william.jallot@epfl.ch)
# - Matthias Wyss (matthias.wyss@epfl.ch)
# - Antoine Garin (antoine.garin@epfl.ch)
#
# ---

# %% [markdown]
# ## 0 - Setup
#
# Creates necessary folders and sets hyperparameters for the project.

# %%
# If a package import fails, install it in your environment, e.g.:
# %pip install yfinance arch copulae statsmodels seaborn

import os
import numpy as np
import copulae
import pandas as pd
import scipy.stats as stats
from pathlib import Path
from itertools import combinations
from statsmodels.graphics.tsaplots import plot_acf, plot_ccf, plot_pacf
from IPython.display import clear_output
from scipy.stats import chi2
from statsmodels.tsa.ar_model import AutoReg
from arch import arch_model
from scipy.stats import norm
import matplotlib.patches as mpatches
```

```

from matplotlib.patches import Patch
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')

# your own script with helper functions, if any
# import utils as U

# Locate the Project 1 directory to this notebook's working directory
PROJECT_DIR = Path.cwd()
DATA_DIR = PROJECT_DIR / 'data'
OUT_DIR = PROJECT_DIR / 'output'
for d in [DATA_DIR, OUT_DIR]:
    d.mkdir(parents=True, exist_ok=True)

# Parameters
TICKERS = ['AAPL', 'META', 'JPM']
START = '2023-01-01'
END = '2025-06-30'
WINDOW = 252
CONFIDENCE = [0.95, 0.99]
np.random.seed(0)

print('Project_dir:', PROJECT_DIR)
print('Output->', OUT_DIR)

# %% [markdown]
# Download and save Adjusted Close for the tickers over the given range into 'data/' (CSV per ticker).

# %%
import yfinance as yf

print('Downloading_data_to', DATA_DIR)
for t in TICKERS:
    print(f'> {t}')
    df = yf.download(t, start=START, end=END, progress=False, auto_adjust=False)
    if df.empty:
        print(f'Warning: no data for {t}')
        continue
    out = df.reset_index()
    out = out[['Date', 'AdjClose']]
    out.to_csv(DATA_DIR / f'{t}.csv', index=False)
print('Done.')

# %% [markdown]
# Now, load the data back from CSVs

# %%
files = [f for f in os.listdir(DATA_DIR) if f.lower().endswith('.csv')]
frames = []
for f in files:
    p = os.path.join(DATA_DIR, f)
    df = pd.read_csv(p, parse_dates=['Date'])
    df = df[['Date', 'AdjClose']]
    # Coerce to numeric and drop malformed rows
    df['AdjClose'] = pd.to_numeric(df['AdjClose'], errors='coerce')

```

```

df = df.dropna(subset=['Date', 'AdjClose'])
df = df.rename(columns={'AdjClose': f.split('.')[0]})  

df = df.set_index('Date').sort_index()  

frames.append(df)  

prices = pd.concat(frames, axis=1).dropna(how='all')  
  

print(prices.head(15))  
  

# %% [markdown]  

# ## 1 - Empirical stylized facts  
  

# %% [markdown]  

# - a. Construct log-returns for AAPL, MSFT, JPM; plot series and comment on trends/  

# volatility.  

# - b. Estimate correlation functions of returns and absolute returns across assets and  

# lags 025; comment.  

# - c. QQ plots vs Normal; perform JarqueBera test and discuss normality.  
  

# %% [markdown]  

# ### a. Log-returns and plots  
  

# %%  

log_returns = np.log(prices).diff().dropna()  
  

# Non-overlapping weekly returns: resample to weekly (Friday close) then compute log returns  

# weekly_prices = prices.resample('W-FRI').last()  

# weekly_log_returns = np.log(weekly_prices / weekly_prices.shift(1)).dropna()  
  

fig, axs = plt.subplots(len(TICKERS), 1, figsize=(10, 3 * len(TICKERS)))  

for i, ticker in enumerate(TICKERS):  

    log_returns[ticker].plot(ax=axs[i], title=f"Log_returns_{ticker}")  

    axs[i].set_ylabel("Log_return")  

plt.tight_layout()  

plt.savefig(os.path.join(OUT_DIR, f"log_returns_all.png"))  

plt.show()  
  

# %% [markdown]  

# Log-returns show no persistent trend (mean around zero).  

# Volatility clusters show us that META has the highest volatility, AAPL is moderate, and  

# JPM is the lowest.  
  

# %% [markdown]  

# ### b. Cross-correlation and autocorrelation functions  
  

# %%  

lags = 25  
  

# ACF align per asset only  

fig_acf, axs_acf = plt.subplots(2, len(TICKERS), figsize=(5*len(TICKERS), 8), sharex=True,  

    sharey=True)  

for j, asset in enumerate(TICKERS):  

    s = log_returns[asset].dropna()  
  

    plot_acf(s, lags=lags, ax=axs_acf[0, j])  

    axs_acf[0, j].set_title(f"{asset}")  

    axs_acf[0, j].axhline(0, lw=1, color="k", alpha=0.7)  
  

    plot_acf(s.abs(), lags=lags, ax=axs_acf[1, j])  

    axs_acf[1, j].set_title(f"|\{asset}|")  

    axs_acf[1, j].axhline(0, lw=1, color="k", alpha=0.7)

```

```

for ax in axs_acf.flat:
    ax.set_xlabel("Lagh")
    ax.set_ylabel("Correlation")
    ax.set_ylim(-0.5, 1.1)

fig_acf.suptitle("Autocorrelation_(returns_and_absolute_returns)", y=1.02)
fig_acf.tight_layout()
fig_acf.savefig(os.path.join(OUT_DIR, "acf_returns_abs_returns.png"))

# CCF - align per pair
pairs = list(combinations(TICKERS, 2)) # We use combinations to avoid duplicate pairs (A,B)
                                         and (B,A) for cross-correlation
fig_ccf, axs_ccf = plt.subplots(2, len(pairs), figsize=(5*len(pairs), 8), sharex=True,
                                sharey=True)

for j, (n, m) in enumerate(pairs):
    # raw returns for direction Corr(n_t, m_{t-h})
    xy = log_returns[[n, m]].dropna()
    x, y = xy[n], xy[m]

    plot_ccf(x, y, lags=lags, ax=axs_ccf[0, j])
    axs_ccf[0, j].set_title(f" $\{n\}$  vs  $\{m\}$ ")
    axs_ccf[0, j].axhline(0, lw=1, color="k", alpha=0.7)

    # absolute returns: take abs on raw series
    xa, ya = x.abs(), y.abs()
    plot_ccf(xa, ya, lags=lags, ax=axs_ccf[1, j])
    axs_ccf[1, j].set_title(f" $| \{n\} |$  vs  $| \{m\} |$ ")
    axs_ccf[1, j].axhline(0, lw=1, color="k", alpha=0.7)

for ax in axs_ccf.flat:
    #ax.set_xlim(0, lags)
    ax.set_xlabel("Lagh")
    ax.set_ylabel("Correlation")
    ax.set_ylim(-0.5, 1.1)

fig_ccf.suptitle("Cross-correlation_(returns_and_absolute_returns)", y=1.02)
fig_ccf.tight_layout()
fig_ccf.savefig(os.path.join(OUT_DIR, "ccf_returns_abs_returns.png"))

plt.show()

# %% [markdown]
# The autocorrelation of log returns for AAPL, META, and JPM shows almost no correlation
# after lag 0. This indicates that past returns have little predictive power for future
# returns, confirming the weak linear dependence typical of stock prices.
#
# Absolute returns exhibit stronger but still weak autocorrelations, with lag 1 correlations
# around 0.2 for AAPL.
#
# Cross-correlations between different stocks log returns are moderate at lag 0 (roughly
# 0.27-0.41) and decay quickly for higher lags. This suggests some contemporaneous co-
# movement across assets, particularly between AAPL and META, but limited predictive
# influence across weeks.
#
# Cross-correlations of absolute returns are weaker than the autocorrelations of absolute
# returns, typically ranging from 0.05 to 0.15 at lag 0. However, they remain positive for
# most lags, indicating modest volatility spillover across assets.

```

```

# %% [markdown]
# ### c. QQ plots and normality tests

# %%
n_assets = len(log_returns.columns)
jarque_bera_results = {}

fig, axes = plt.subplots(n_assets, 1, figsize=(8, 3*n_assets))
for i, ticker in enumerate(log_returns.columns):
    stats.probplot(log_returns[ticker], dist="norm", plot=axes[i])
    axes[i].set_title(f"Normal QQ-Plot: {ticker}")

    jarque_bera_results[ticker] = stats.jarque_bera(log_returns[ticker])

fig.tight_layout()
fig.savefig(os.path.join(OUT_DIR, "qq_plots.png"))
fig.show()

# %% [markdown]
# We can see that across the three assets, the result of the QQ plot exhibit the same 'shape'. At the two extremes of the 45 degree line, the theoretical quantiles are below on the left and above on the right. This means our log returns are left skewed (long negative tail) and are fat tailed.

# %% [markdown]
# The Jarque-Bera test tests whether the sample data has the skewness and kurtosis matching a normal distribution. It tests, if we can reject the null hypothesis that our returns are normally distributed. If we want to test at 5% level of significance, in case our  $p_{\text{value}}$ , the lowest significance at which you can reject, if it is lower than 5% we would reject normality.

#
#
# %%
for ticker, result in jarque_bera_results.items():
    jb_stat, jb_pvalue = result[0], result[1]
    print(f"Jarque-Bera test for {ticker}:")
    print(f"JB Statistic: {jb_stat:.4f}")
    print(f"p-value: {jb_pvalue:.4f}")
    if jb_pvalue < 0.05:
        print("=> Reject the null hypothesis of normality at the 5% significance level.\n")
    else:
        print("=> Fail to reject the null hypothesis of normality at the 5% significance level.\n")

# %% [markdown]
# ## 2 - First-window modeling: VaR, ES, and distributions

# %% [markdown]
# Use the first estimation window  $W=252$  days on each asset separately with losses  $L_t = R_t$ .
# Compare:
#
# - Historical,
# - Gaussian,
# - Student-t,
# - AR(p)+GARCH(1,1) with Normal/Student-t,
# - Filtered Historical Simulation (FHS).

```

```

# %% [markdown]
# We use the first estimation window of length  $W = 252$  days ( $\sim 1$  trading year) for each asset
# separately.
# Losses are defined as  $L_t = -R_t$ , so the right tail corresponds to risk.
#
# ### a. Historical Simulation (Empirical CDF)
#
# The **empirical cumulative distribution function (CDF)** of losses is:
#
# 
$$\hat{F}_L(x) = \frac{1}{W} \sum_{t=1}^W \mathbf{1}_{\{L_t \leq x\}}$$

#
# - **Value-at-Risk (VaR)** at confidence level  $\alpha$ :
#
# 
$$VaR_\alpha = \inf \{x : \hat{F}_L(x) \geq \alpha\}$$

#
# - **Expected Shortfall (ES)**:
#
# 
$$ES_\alpha = \mathbb{E}[L_t | L_t \geq VaR_\alpha] \approx \text{mean of losses}_{VaR_\alpha}$$

#
# **Notes:**
# - Non-parametric approach: no assumption on the loss distribution.
# - Directly based on historical data.
#
# ---

# %%
def historical_simulation(L, confidence_levels) :
    VaR_hist = {a: np.quantile(L, a) for a in confidence_levels}
    ES_hist = {a: L[L >= VaR_hist[a]].mean() for a in confidence_levels}

    out = {
        "VaR": {},
        "ES": {}
    }
    for a in confidence_levels:
        out["VaR"][a] = VaR_hist[a]
        out["ES"][a] = ES_hist[a]
    return out

# We will need this later for Monte Carlo simulations
def historical_sampler(L, M) :
    return np.random.choice(L, replace=True, size=M)

# %% [markdown]
# ### b. Gaussian (Normal) Distribution
#
# Assume losses are **normally distributed**:
#
# 
$$L_t \sim \mathcal{N}(\mu, \sigma^2)$$

#
# where:
#

```

```

# $$ 
# \mu = \text{mean}(L), \quad \sigma = \text{std}(L)
# $$ 
#
# - **Value-at-Risk (VaR)**:
#
# $$ 
# \text{VaR}_\alpha = \mu + \sigma \cdot z_\alpha
# $$ 
#
# where  $z_\alpha = \Phi^{-1}(\alpha)$  is the  $\alpha$ -quantile of the standard normal distribution.
#
# - **Expected Shortfall (ES)**:
#
# $$ 
# \text{ES}_\alpha = \mu + \sigma \frac{\phi(z_\alpha)}{1-\alpha}
# $$ 
#
# with  $\phi(z) = \frac{1}{\sqrt{2\pi}} e^{-z^2/2}$ , the standard normal density.
#
# **Notes:**
# - Closed-form formulas make Gaussian VaR/ES fast to compute.
# - Limitation: does not capture fat tails in the loss distribution.
#
# ---
#
# %%

def gaussian_normal_distribution(L, confidence_levels) :
    mu, sigma = norm.fit(L)
    VaR_gauss = {a: mu + sigma * norm.ppf(a) for a in confidence_levels}
    ES_gauss = {a: mu + sigma * norm.pdf(norm.ppf(a)) / (1 - a) for a in confidence_levels}
    out = {
        "mu": mu,
        "sigma": sigma,
        "VaR": {},
        "ES": {}
    }
    for a in confidence_levels:
        out["VaR"][a] = VaR_gauss[a]
        out["ES"][a] = ES_gauss[a]
    return out

# We will need this later for Monte Carlo simulations
def gaussian_sampler(L, M) :

    mu, sigma = norm.fit(L)
    return np.random.normal(loc=mu, scale=sigma, size=M)

# %% [markdown]
# ### c. Student-t Distribution
#
# Assume losses follow a **Student-t distribution**:
#
# $$ 
# L_t \sim t_{nu}(\mu, \sigma)
# $$ 
#
# where:
# -  $nu$  = degrees of freedom (controls tail thickness)

```

```

# -  $\mu$  = location parameter
# -  $\sigma$  = scale parameter
#
# - **Value-at-Risk (VaR)** at confidence level  $\alpha$ :
#
# $$
# \text{VaR}_{\alpha} = \mu + \sigma \cdot t_{\nu^{-1}}(\alpha)
# $$
#
# where  $t_{\nu^{-1}}(\alpha)$  is the  $\alpha$ -quantile of the standard Student-t with  $\nu$  degrees of freedom.
#
# - **Expected Shortfall (ES)** for  $\nu > 1$ :
#
# $$
# \text{ES}_{\alpha} = \mu + \sigma \frac{\nu + \big(t_{\nu^{-1}}(\alpha)\big)^2}{\nu - 1} \cdot \frac{f_{\nu}(t_{\nu^{-1}}(\alpha))}{1 - \alpha}
# $$
#
# where  $f_{\nu}(x)$  is the standard Student-t density.
#
# **Impact of  $\nu$ :**
# - Small  $\nu$  fatter tails higher VaR and ES (more conservative estimates)
# - Large  $\nu$  approaches normal distribution VaR/ES similar to Gaussian
#
# **Notes:**
# - Student-t captures extreme losses better than Gaussian due to fat tails.
# - VaR and ES depend strongly on the estimated degrees of freedom  $\nu$ .
#
# %%

from scipy.stats import t

def student_var_es(L, confidence_levels) :
    nu, mu_t, sigma_t = t.fit(L)
    VaR_t = {a: mu_t + sigma_t * t.ppf(a, nu) for a in confidence_levels}
    ES_t = {a: mu_t + sigma_t * (nu + (t.ppf(a, nu))**2)/(nu - 1) * t.pdf(t.ppf(a, nu), nu) / (1 - a) for a in confidence_levels}

    out = {
        "VaR": [],
        "ES": []
    }
    for a in confidence_levels:
        out['nu'] = nu
        out['mu_t'] = mu_t
        out['sigma_t'] = sigma_t
        out["VaR"][a] = VaR_t[a]
        out["ES"][a] = ES_t[a]
    return out

# We will need this later for Monte Carlo simulations
def student_sampler(L, M) :
    nu, mu_t, sigma_t = t.fit(L)
    return t.rvs(df=nu, loc=mu_t, scale=sigma_t, size=M)

# %% [markdown]
# ### d. Conditional parametric

# %%
n_assets = len(TICKERS)

```

```

fig, axes = plt.subplots(n_assets, 1, figsize=(8, 4 * n_assets))
for i, asset in enumerate(TICKERS):
    plot_pacf(
        log_returns[asset],
        lags=lags,
        ax=axes[i],
        title=f'PACF_{asset}'
    )
    axes[i].set_ylim(-0.5, 1.1)

plt.tight_layout()
plt.savefig(os.path.join(OUT_DIR, "pacf_plots.png"))

# %% [markdown]
# After lag 0, most values fluctuate around zero and are within the blue confidence band so
# we should use an AR(0) model, a pure noise model.

# %%
def conditional_parametric_var_es(L, confidence_levels) :
    res_mean = AutoReg(L, lags=0).fit()
    eps = res_mean.resid

    garch = arch_model(eps, mean='Zero', vol='GARCH', p=1, q=1)
    res_vol = garch.fit(show_warning=False, disp='off')
    vol_forecast = res_vol.forecast(horizon=1)

    sigma_forecast = np.sqrt(vol_forecast.variance.values[-1, 0])
    mean_forecast = res_mean.predict(start=len(L), end=len(L))

    VaR_parametric = {a: mean_forecast.values[0] + sigma_forecast * norm.ppf(a) for a in
                      confidence_levels}
    ES_parametric = {a: mean_forecast.values[0] + sigma_forecast * norm.pdf(norm.ppf(a)) /
                     (1 - a) for a in confidence_levels}

    out = {
        "VaR": {},
        "ES": {},
        "mu_forecast": mean_forecast,
        "sigma_forecast": sigma_forecast
    }
    for a in confidence_levels:
        out["VaR"][a] = VaR_parametric[a]
        out["ES"][a] = ES_parametric[a]

    return out

# We will need this later for Monte Carlo simulations
def conditional_parametric_sampler(L, M) :

    res_mean = AutoReg(L, lags=0).fit()
    eps = res_mean.resid

    garch = arch_model(eps, mean='Zero', vol='GARCH', p=1, q=1)
    res_vol = garch.fit(show_warning=False, disp='off')
    vol_forecast = res_vol.forecast(horizon=1)

    sigma_forecast = np.sqrt(vol_forecast.variance.values[-1, 0])
    mean_forecast = res_mean.predict(start=len(L), end=len(L))

    return np.random.normal(loc=mean_forecast, scale=sigma_forecast, size=M)

```

```

# %% [markdown]
# ### e. Filtered Historical Simulation

# %%
def filtered_historical_simulation(L, alpha_levels, M=1000, random_state=0):
    """
    Filtered_Historical_Simulation(FHS) on LOSS series L.

    Parameters
    -----
    L: array-like or pd.Series
        Losses (positive on bad days, small/possibly negative on good days).
    alpha_levels: list of float
        Confidence levels, e.g. [0.95, 0.99].
    M: int
        Number of Monte Carlo simulations.
    """
L = pd.Series(L).dropna()

# AR(0)
ar = AutoReg(L, lags=0).fit()
eps = ar.resid # residuals of losses

# GARCH(1,1) on residuals
garch = arch_model(eps, mean='Zero', vol='GARCH', p=1, q=1)
garch_res = garch.fit(show_warning=False, disp='off')

# Conditional volatility and 1-step-ahead forecast
z_t = garch_res.std_resid

h_t1 = garch_res.forecast(horizon=1).variance.iloc[-1, 0]
sigma_t1 = float(np.sqrt(h_t1))

# One-step-ahead conditional mean of losses
mu_t1 = ar.forecast(steps=1).values[0]

z_star = np.random.choice(z_t, size=M, replace=True)
eps_t1_star = sigma_t1 * z_star

# Simulated next-period losses
L_simulated = mu_t1 + eps_t1_star

# VaR and ES
out = {
    "mu_forecast": mu_t1,
    "h_forecast": float(h_t1),
    "r_sims": L_simulated,
    "VaR": {},
    "ES": {}
}

for a in alpha_levels:
    # a is the confidence level
    var_a = np.quantile(L_simulated, a)
    tail = L_simulated[L_simulated >= var_a]
    es_a = float(tail.mean())

    out["VaR"][a] = float(var_a)
    out["ES"][a] = es_a

```

```

    return out

# We will need this later for Monte Carlo simulations
def fhs_sampler(L, M):
    # AR(0) on LOSSES
    ar = AutoReg(L, lags=0).fit()
    eps = ar.resid

    # GARCH(1,1) on residuals
    garch = arch_model(eps, mean='Zero', vol='GARCH', p=1, q=1)
    garch_fitted = garch.fit(show_warning=False, disp='off')

    sigma_t = garch_fitted.conditional_volatility
    h_t1 = garch_fitted.forecast(horizon=1).variance.values[-1, 0] # variance
    sigma_t1 = np.sqrt(h_t1)

    # 1-step ahead mean
    mu_t1 = ar.forecast(steps=1) # shape (1,)

    # Standardized residuals
    z_t = eps / sigma_t

    # Resampling
    e_star = np.random.choice(z_t, size=M, replace=True)
    z_t1_star = e_star * sigma_t1

    # Simulated next period
    r_simulated = mu_t1.ravel() + z_t1_star

    return r_simulated

# %%
# Plot PDFs for all tickers in one figure
fig, axes = plt.subplots(len(TICKERS), 1, figsize=(8, 5*len(TICKERS)), sharex=True)

# Dictionaries to store results
VaR_results = []
ES_results = []

for i, ticker in enumerate(TICKERS):
    ax = axes[i] # Select subplot
    ax.set_xlim(-0.1, 0.1)

    # First W observations
    L = -log_returns[ticker].iloc[:WINDOW]

    # === Historical ===
    out_hist = historical_simulation(L, CONFIDENCE)
    sns.kdeplot(L, label='Historical', bw_method=0.3, ax=ax)

    # === Gaussian ===
    out_gauss = gaussian_normal_distribution(L, CONFIDENCE)
    x = np.linspace(L.min(), L.max(), 200)
    ax.plot(x, norm.pdf(x, out_gauss['mu'], out_gauss['sigma']), label='Gaussian')

    # === Student-t ===
    out_t = student_var_es(L, CONFIDENCE)
    ax.plot(x, t.pdf((x - out_t['mu_t'])/out_t['sigma_t'], out_t['nu'])/out_t['sigma_t'],
             label='Student-t')

```

```

# === Conditional parametric ===
out_parametric = conditional_parametric_var_es(L, CONFIDENCE)
ax.plot(x, norm.pdf(x, out_parametric['mu_forecast'], out_parametric['sigma_forecast']),
        label='Conditional Parametric')

# === Filtered Historical Simulation (FHS) ===
out_fhs = filtered_historical_simulation(L, CONFIDENCE)
sns.kdeplot(out_fhs['r_sims'], label='FHS', bw_method=0.3, ax=ax)

# === Store results ===
for a in CONFIDENCE:
    VaR_results.append({
        'Ticker': ticker,
        'Alpha': a,
        'Historical': out_hist['VaR'][a],
        'Gaussian': out_gauss['VaR'][a],
        'Student-t': out_t['VaR'][a],
        'Parametric' : out_parametric['VaR'][a],
        'FHS': out_fhs['VaR'][a]
    })
    ES_results.append({
        'Ticker': ticker,
        'Alpha': a,
        'Historical': out_hist['ES'][a],
        'Gaussian': out_gauss['ES'][a],
        'Student-t': out_t['ES'][a],
        'Parametric' : out_parametric['ES'][a],
        'FHS': out_fhs['ES'][a]
    })
}

ax.set_title(f"Estimated PDFs of {ticker} Losses")
ax.set_ylabel("Density")
ax.legend()

# Convert results to DataFrames
VaR_df = pd.DataFrame(VaR_results).set_index(['Ticker', 'Alpha'])
ES_df = pd.DataFrame(ES_results).set_index(['Ticker', 'Alpha'])

# Format the Alpha index as percentages
VaR_df.index = VaR_df.index.set_levels([VaR_df.index.levels[0],
                                         [f"{int(a*100)}%" for a in VaR_df.index.levels[1]]])
ES_df.index = ES_df.index.set_levels([ES_df.index.levels[0],
                                         [f"{int(a*100)}%" for a in ES_df.index.levels[1]]])

# Display VaR DataFrames
print("== Value-at-Risk (VaR) Results ==")
display(VaR_df)

# Display ES DataFrames
print("\n== Expected Shortfall (ES) Results ==")
display(ES_df)

# Save and show plots
axes[-1].set_xlabel("Loss")
plt.tight_layout()
plt.savefig(os.path.join(OUT_DIR, "loss_pdfs.png"))
plt.show()

# %% [markdown]
# The PDF plots reveal deviations from Gaussianity across assets, with all three showing
# heavier tails and stronger kurtosis than the normal model can capture. Student-t and

```

conditional parametric approaches provide better tail fitting, while FHS most closely reproduces the empirical distributional features. The differences across AAPL, META, and JPM further confirm that tail behaviour is asset-dependent.

```
# %% [markdown]
# ## 3) Rolling-window backtesting of VaR and ES

# %% [markdown]
# Use a rolling window of size  $W$  to produce 1-step-ahead VaR/ES at 95% and 99% for each
# method in Exercise 2. Then, implement the following statistical tests:
#
# - VaR backtests: Kupiec POF and Christoffersen independence tests.
# - ES backtest: AcerbiSzékely Z1 test.

# %%
METHODS = {
    "historical": historical_simulation,
    "gaussian": gaussian_normal_distribution,
    "student": student_var_es,
    "parametric": conditional_parametric_var_es,
    "fhs": filtered_historical_simulation
}

# Put results into a nice rolling forecast dataframe
def rolling_forecast(L, W, methods, confidence):
    results = []
    for name, func in methods.items():
        for metric in ["VaR", "ES"]:
            for a in confidence:
                series = L.rolling(W).apply(
                    lambda loss: func(loss, confidence)[metric][a]
                )
                series.name = (name, metric, a)
                results.append(series)

    out = pd.concat(results, axis=1)
    # Make it into a nice dataframe to facilitate computations
    out.columns = pd.MultiIndex.from_tuples(out.columns, names=["method", "metric", "alpha"])
]

    return out.dropna(how="all")

rolling_results = {}
for ticker in TICKERS:
    print(f'Processing {ticker}... ')
    L = -log_returns[ticker]
    rolling_results[ticker] = rolling_forecast(L, WINDOW, METHODS, CONFIDENCE)

clear_output()
display(rolling_results)

# %%
import matplotlib.dates as mdates

def plot_var_subplots(rolling_results, log_returns, tickers, methods_to_plot, confidence
                      =0.95):
    n_methods = len(methods_to_plot)
    n_tickers = len(tickers)
```

```

fig, axes = plt.subplots(n_methods, n_tickers, figsize=(6*n_tickers, 3*n_methods),
                        sharex=True)

for i, method in enumerate(methods_to_plot):
    for j, ticker in enumerate(tickers):

        ax = axes[i, j]

        L = -log_returns[ticker].dropna()
        df = rolling_results[ticker]

        VaR_series = df[(method, "VaR", confidence)].dropna()

        # Align loss and VaR indices
        common_idx = L.index.intersection(VaR_series.index)
        L_aligned = L.loc[common_idx]
        V_aligned = VaR_series.loc[common_idx]

        ax.plot(L.index, L.values, color="black", alpha=0.35, linewidth=0.9)
        ax.plot(V_aligned.index, V_aligned.values, linewidth=1.6, label=f"{method}\u222a VaR")

        violations = L_aligned > V_aligned
        ax.scatter(
            L_aligned[violations].index, L_aligned[violations].values,
            color="red", marker="x", s=50
        )

        # Titles
        if i == 0:
            ax.set_title(f"{ticker}", fontsize=14, fontweight="bold")
        if j == 0:
            ax.set_ylabel(f"{method}\nLoss\u222a VaR", fontsize=11)

        # Grid
        ax.grid(alpha=0.3)
        ax.legend(fontsize=8, loc="upper_left")

        # Prevent overlapping date labels
        if i == n_methods - 1: # only bottom row shows dates
            ax.xaxis.set_major_locator(mdates.MonthLocator(interval=3))
            ax.xaxis.set_major_formatter(mdates.DateFormatter('%Y-%m'))
            plt.setp(ax.get_xticklabels(), rotation=45, ha='right')
        else:
            ax.set_xticklabels([])

plt.suptitle(f"Rolling\u222a VaR\u222a vs\u222a Realized\u222a Losses\u222a({int(confidence*100)}%\u222a Confidence)", fontsize=16, fontweight="bold")
plt.tight_layout(rect=[0, 0, 1, 0.96])
plt.show()

plt.suptitle(f"Rolling\u222a VaR\u222a vs\u222a Realized\u222a Losses\u222a({int(confidence*100)}%\u222a Confidence)", fontsize=16, fontweight="bold")
plt.tight_layout(rect=[0, 0, 1, 0.97])
plt.show()

methods_to_plot = ["historical", "fhs", "gaussian", "student", "parametric"]
methods_to_plot = ["historical", "fhs", "gaussian", "student", "parametric"]

plot_var_subplots(rolling_results=rolling_results, log_returns=log_returns, tickers=TICKERS,
                  methods_to_plot=methods_to_plot, confidence=0.95)

```

```

# %% [markdown]
# From this plot, we observe that the non-parametric methods produce relatively flat VaR
# curves, since they do not explicitly react to changes in volatility as it is washed out
# by the estimation window. In contrast, the parametric approaches such as FHS and the
# conditional parametric model adjust the VaR threshold dynamically and capture local
# volatility spikes.

# %%
def kupiec_pof_test(VaR_series, L_series, alpha):
    """
    Kupiec (1995) Proportion of Failures (POF) test.
    H0: The observed failure rate equals alpha.
    """

    df = pd.concat([VaR_series, L_series], axis=1)
    df.columns = ['VaR', 'Loss']
    df = df.dropna()

    # Identify breaches
    df['breached'] = df['Loss'] > df['VaR']

    # Compute statistics
    x = df['breached'].sum()
    n = len(df)
    alpha_hat = x / n

    # Edge cases
    if alpha_hat == 0 or alpha_hat == 1:
        return np.nan

    # Kupiec POF statistic
    POF = 2 * np.log(
        ((1 - alpha_hat) / (1 - alpha))**(n - x) *
        (alpha_hat / alpha)**x
    )

    return POF

# %%
def christoffersen(VaR_series, L_series):
    """
    Christoffersen (1998) conditional coverage test
    """

    VaR = pd.Series(VaR_series, name="VaR").astype(float)
    Loss = pd.Series(L_series, name="Loss").astype(float)
    df = pd.concat([VaR, Loss], axis=1).dropna()

    # Breach indicator
    df['breached'] = df['Loss'] > df['VaR']

    # Previous breach via shift
    df['breached_previous'] = df['breached'].shift(1)

    # Keep rows where previous state is defined
    df = df.dropna(subset=['breached_previous'])

    # Transition counts using boolean masks
    N_0_0 = ((~df['breached']) & (~df['breached_previous'])).sum()
    N_0_1 = ((df['breached']) & (~df['breached_previous'])).sum()
    N_1_0 = ((~df['breached']) & (df['breached_previous'])).sum()

```

```

N_1_1 = ((df['breached']) & (df['breached_previous'])).sum()

N = N_0_0 + N_0_1 + N_1_0 + N_1_1
if N == 0:
    return np.nan # not enough transitions

denom0 = N_0_0 + N_0_1 # times with prev=0
denom1 = N_1_0 + N_1_1 # times with prev=1
if denom0 == 0 or denom1 == 0:
    return np.nan

pi_0 = N_0_1 / denom0
pi_1 = N_1_1 / denom1
p = (N_0_1 + N_1_1) / N

# In case log(0)
eps = 1e-12
pi_0 = np.clip(pi_0, eps, 1 - eps)
pi_1 = np.clip(pi_1, eps, 1 - eps)
p = np.clip(p, eps, 1 - eps)

L0_log = (N_0_0 + N_1_0) * np.log(1 - p) + (N_0_1 + N_1_1) * np.log(p)
L1_log = (N_0_0 * np.log(1 - pi_0) + N_0_1 * np.log(pi_0) +
           N_1_0 * np.log(1 - pi_1) + N_1_1 * np.log(pi_1))
LR_ind = -2 * (L0_log - L1_log)

return LR_ind

# %%
def chi_square_test(t_stat, df, method_name, test_name) :
    print(f'Performing {test_name}')
    p_value = 1 - chi2.cdf(t_stat, df = df)

    print(f'The test statistic and p-value for {method_name} are: t_stat={t_stat}, p_value={p_value}')

    if p_value < 0.05 :
        print(f'Reject the null hypothesis at 5% significance level for {method_name}. \n')
    else :
        print(f'Fail to reject the null hypothesis at 5% significance level for {method_name}.\n')
    return p_value

# %%
confidence = 0.95
alpha = 1 - confidence
METHODS = ['historical', 'gaussian', 'student', 'parametric', 'fhs']
result = pd.DataFrame(columns = ['Ticker', 'Method', 'test_statistic_pof', 'p_value_pof', ,
    'test_statistic_cd_coverage', 'p_value_cd_coverage'])
for ticker in rolling_results.keys() :
    print('*'*115)
    for method in METHODS:
        ticker_df = rolling_results[ticker][method]['VaR'][confidence]
        ticker_loss = -log_returns[ticker]

        df = pd.DataFrame({'Loss': ticker_loss, "VaR": ticker_df})
        # Shift loss series to align with VaR forecasts
        df["Loss"] = df["Loss"].shift(-1)
        df = df.dropna()
        POF = kupiec_pof_test(df['VaR'], df['Loss'], alpha)
        INC = christoffersen(df['VaR'], df['Loss'])

```

```

result = pd.concat([
    result,
    pd.DataFrame([
        {'Ticker': ticker,
         'Method': method,
         'test_statistic_pof': POF,
         'p_value_pof': chi_square_test(
             POF, df=1, method_name=f'{ticker}_{method}', test_name='POF_Utest'
         ),
         'test_statistic_cd_coverage': INC + POF,
         'p_value_cd_coverage': chi_square_test(
             INC + POF, df=2, method_name=f'{ticker}_{method}', test_name=
                 Christoffersen_Uconditional_Ucoverage_Utest'
         )
    ])
], ignore_index=True)

print('*'*115)

# %%
result.pivot(index=['Ticker', 'Method'], columns=[], values=['test_statistic_pof', 'p_value_pof', 'test_statistic_cd_coverage', 'p_value_cd_coverage'])

# %% [markdown]
# Across assets, all VaR models pass both the coverage and independence tests for META and JPM, indicating that their VaR and ES forecasts are statistically consistent with observed losses.
# In contrast, every model fails for AAPL, indicating an underestimation of the tail risk.
#
# We identify whether failures come from biased coverage or clustered exceptions by checking which test rejects the null: Kupiec detects incorrect violation frequency, while Christoffersens independence test reveals clustering, combining both indicates full model adequacy.
#
# Since AAPL fails both the unconditional coverage test and the joint conditional coverage test for each methods, this confirms biased coverage, indicating that the models systematically underestimate its tail risk captured.

# %%
def es_backtest_z1(VaR_series, ES_series, Loss_series):
    """
    Implements the MSCI_Z1_Expected_Shortfall_backtest (Acerbi & Szekely, 2014)
    Using the 'X_t' definition: P&L (positive=profit, negative=loss)
    """
    VaR = pd.Series(VaR_series, name="VaR", dtype="float64")
    ES = pd.Series(ES_series, name="ES", dtype="float64")
    Loss = pd.Series(Loss_series, name="Loss", dtype="float64")

    df = pd.concat([VaR, ES, Loss], axis=1).dropna()

    df["I_t"] = (df["Loss"] > df["VaR"]).astype(int)

    N_T = int(df["I_t"].sum())

    if N_T == 0:
        return np.nan # No breaches, Z1 undefined

    Z_1 = ( (df["Loss"] * df["I_t"]) / df["ES"] ).sum() / N_T - 1

    return Z_1

```

```

def monte_carlo_es_backtest_z1(VaR_series, ES_series, Loss_series, sampler, window, M=1000):
    """
    VaR_series, ES_series:pandas.Series
    Out-of-sample VaR/ES forecasts, same index (e.g. from rolling_forecast).
    Loss_series:pandas.Series
    Full loss series (in-sample + out-of-sample), NOT truncated.
    sampler(window_array, M) -> np.ndarray of shape (M,)
    Function that, given a window of losses, returns M simulated losses.
    window:int
    Rolling window length used to estimate VaR/ES.
    M:int
    Number of Monte Carlo paths.

    Returns
    ----
    Z1:np.ndarray of shape (M,)
    Monte Carlo distribution of the ES backtest statistic Z1.
    """
    # Make sure everything is aligned and sorted
    VaR = VaR_series.sort_index()
    ES = ES_series.reindex(VaR.index) # ensure same index/order

    # Full loss series (sorted)
    Loss_series = Loss_series.sort_index()
    x = np.asarray(Loss_series, dtype=float)
    loss_index = Loss_series.index

    idx = VaR.index # out-of-sample dates
    T = len(idx)

    sims = np.empty((M, T), dtype=float)

    # For each date, build the same rolling window used to estimate VaR/ES
    for j, date in enumerate(idx):
        loc = loss_index.get_loc(date) # position of this date in the full loss series

        # Window of length 'window' ending at 'date'
        w = x[loc - window + 1 : loc + 1]
        draws = sampler(w, M)
        sims[:, j] = draws

    # Compute Z1 for each simulated path
    Z1 = np.empty(M, dtype=float)
    for m in range(M):
        sim_series = pd.Series(sims[m], index=idx)
        Z1[m] = es_backtest_z1(VaR, ES, sim_series)

    return Z1

# %%
sampler = {'historical': historical_sampler, 'gaussian': gaussian_sampler, 'student':
           student_sampler, 'parametric': conditional_parametric_sampler, 'lhs': lhs_sampler}

result = pd.DataFrame(columns=['Ticker', 'Method', 'Z1_stat', 'p_value_es_backtest'])
confidence = 0.95

for ticker in rolling_results.keys():

```

```

print('-' * 115)

# Full loss series for this ticker (same as in rolling_forecast)
ticker_loss = -log_returns[ticker].dropna()
for method in METHODS:
    ticker_df_es = rolling_results[ticker][method]['ES'][confidence]
    ticker_df_var = rolling_results[ticker][method]['VaR'][confidence]

    Z_1_simulated = monte_carlo_es_backtest_z1(ticker_df_var, ticker_df_es, ticker_loss,
                                                sampler[method], window=WINDOW, M=1000)

    real_Z1 = es_backtest_z1(ticker_df_var, ticker_df_es, ticker_loss)

    p_value = (Z_1_simulated >= real_Z1).sum() / len(Z_1_simulated)

    result = pd.concat([result,pd.DataFrame([{'Ticker': ticker, 'Method': method, 'Z1_stat': real_Z1, 'p_value_es_backtest': p_value}])], ignore_index=True)

print(f'The test statistic and p-value for {ticker} {method} are: {real_Z1}, {p_value}')
if p_value < 0.05:
    print(f'Reject the null hypothesis at 5% significance level for {ticker} {method}. \n')
else:
    print(f'Fail to reject the null hypothesis at 5% significance level for {ticker} {method}. \n')

# %%
result.pivot(index=['Ticker', 'Method'], columns=[], values=['Z1_stat', 'p_value_es_backtest'])

# %% [markdown]
# For AAPL, only the Gaussian ES model fails the  $Z_1$  test, with a p-value of $0$, indicating that the Gaussian assumption systematically underestimates tail severity.
# All other models exhibit p-values above 5%, indicating that their ES estimates are statistically compatible with the realized tail losses.
#
# For META, the Gaussian model fails, with a very small p-value of $0$, indicating that the Gaussian assumption systematically underestimates tail severity.
# The conditional parametric model also fails with a p-value of  $2 \cdot 10^{-2}$ 
# All other models exhibit p-values above 5% showing no statistical evidence of ES misspecification.
#
#
# For JPM, the Gaussian model has p-value $0$ and the parametric model $0$ and the student model  $3.50 \cdot 10^{-2}$  both fail the  $Z_1$  test. These models underpredict the magnitude of extreme losses. The historical and FHS models pass, indicating that their ES forecasts align with realized tail behavior.
#
#
# Across the three assets, the results show that ES model performance is strongly asset dependent, reflecting differences in tail behaviour across return distributions. A common pattern is that the Gaussian model systematically fails for all assets, confirming that its thin-tailed distributional assumption leads to persistent underestimation of extreme losses. However, the fact that different models fail for different assets highlights that the appropriateness of an ES model depends critically on asset-specific tail dynamics. Notably, the filtered historical simulation performs well across all cases, suggesting that its semi-parametric structure offers robust flexibility in practice.

```

```

# %% [markdown]
# ## 4) Copula fitting (first window)

# %% [markdown]
# - a. Visualize dependence in returns and copula space using pseudo-observations.
# - b. Fit Gaussian and t copulas; report parameters.
# - c. Simulate from fitted copulas and map to empirical marginals; compare with original
#      returns.

# %%
# TODO

# Use copulae package for copula fitting and sampling
# cop = copulae.elliptical.GaussianCopula(dim=len(TICKERS))
# cop.fit(data)
# samples = cop.random(n)

# %% [markdown]
# ### a. Pseudo-observations and dependence visualization
#
# Pseudo-observations transform the marginal distributions to uniform  $[0, 1]$  using
# empirical ranks:
#
# $$
# U_{t,i} = \frac{\text{rank}(R_{t,i})}{W+1}, \quad t = 1, \dots, W
# $$
#
# This allows us to visualize the **pure dependence structure** (copula) separately from
# marginal behavior.

# %%
# Extract the first estimation window (W=252 observations) across assets
returns_window = log_returns.iloc[:WINDOW].copy()

# Compute pseudo-observations (empirical quantiles)
def compute_pseudo_obs(data):
    """
    Transform data to pseudo-observations in [0, 1] using empirical ranks.

    Parameters:
    data (pd.DataFrame): Input data with columns corresponding to assets.

    Returns:
    pd.DataFrame: Pseudo-observations where each row is a vector of length n_assets.
    """
    n = len(data)
    U = data.rank() / (n + 1)
    return U

pseudo_obs = compute_pseudo_obs(returns_window)

print("Pseudo-observations (first 5 rows):")
print(pseudo_obs.head())
print(f"\nShape: {pseudo_obs.shape}")
print(f"Range: [{pseudo_obs.min().min():.4f}, {pseudo_obs.max().max():.4f}]")

# %%
returns_window.describe()

# %%
# Visualize dependence: raw returns space vs copula (pseudo-observations) space
pairs = list(combinations(TICKERS, 2))
n_pairs = len(pairs)

fig, axes = plt.subplots(2, n_pairs, figsize=(6*n_pairs, 10))

```

```

for j, (asset_i, asset_j) in enumerate(pairs):
    # Raw returns space
    ax_raw = axes[0, j]
    ax_raw.scatter(returns_window[asset_i], returns_window[asset_j],
                   alpha=0.5, s=20, edgecolors='k', linewidths=0.5)
    ax_raw.set_xlabel(f'{asset_i}_returns')
    ax_raw.set_ylabel(f'{asset_j}_returns')
    ax_raw.set_title(f'Raw Returns: {asset_i} vs {asset_j}')
    ax_raw.grid(True, alpha=0.3)

    # Copula space (pseudo-observations)
    ax_copula = axes[1, j]
    ax_copula.scatter(pseudo_obs[asset_i], pseudo_obs[asset_j],
                      alpha=0.5, s=20, edgecolors='k', linewidths=0.5, color='coral')
    ax_copula.set_xlabel(f'U({asset_i})')
    ax_copula.set_ylabel(f'U({asset_j})')
    ax_copula.set_title(f'Copula Space: {asset_i} vs {asset_j}')
    ax_copula.set_xlim(0, 1)
    ax_copula.set_ylim(0, 1)
    ax_copula.grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig(os.path.join(OUT_DIR, 'raw_returns_vs_pseudo_obs.png'))
plt.show()

# %% [markdown]
# ##### Interpretation of Raw Returns and Copula Space
#
# **Raw returns**
#
# The **AAPL META** pair exhibits a clear upward-sloping cloud, indicating a positive dependence and frequent joint large moves. In contrast, **AAPL JPM** and **META JPM** show more diffuse scatter patterns, reflecting weaker dependence between tech firms and a bank. Across all pairs, extreme returns are more dispersed than expected under normality, suggesting heavy-tailed behavior and occasional simultaneous large shocks.
#
# **Copula space**
#
# In copula space, **AAPL META** displays a marked diagonal structure, confirming a positive dependence. The clustering of observations near the corners (0,0) and (1,1) further reveals tail dependence meaning that extreme moves tend to occur together. By comparison, **AAPL JPM** and **META JPM** resemble an almost uniform cloud over  $([0,1]^2)$ , confirming weaker dependence and only marginal tail co-movement.
#
# **Implications for copula choice**
#
# Given the evident heavy tails and clear tail dependence in **AAPL META**, a Gaussian copula, though adequate for symmetric dependence, cannot capture the probability mass in the joint extremes. The Student-t copula, however, explicitly models tail dependence and is therefore the appropriate copula for this portfolio. Its structure aligns with the empirical behavior of the assets, particularly the strong, tail-driven co-movements between the two tech stocks. Consequently, the Student-t copula provides a more realistic representation of joint downside risk than its Gaussian counterpart.

# %% [markdown]
# ### b. Fit Gaussian and Student-t copulas
#
# **Gaussian Copula:**  

# 
$$C^{\text{Gauss}}(u_1, \dots, u_d; \Sigma) = \Phi_{\Sigma}(\Phi^{-1}(u_1), \dots, \Phi^{-1}(u_d))$$


```

```

# $$#
# where  $\Phi_{\Sigma}$  is the multivariate normal CDF with correlation matrix  $\Sigma$ .
#
# **Student-t Copula:**#
# $$#
#  $C^t(u_1, \dots, u_d; \Sigma, \nu) = T_{\{\Sigma, \nu\}}(T_{\nu^{-1}(u_1)}, \dots, T_{\nu^{-1}(u_d)})$ 
# $$#
# where  $T_{\{\Sigma, \nu\}}$  is the multivariate Student-t CDF with correlation  $\Sigma$  and degrees of freedom  $\nu$ .
#
# %%
# Fit Gaussian Copula
gauss_cop = copulae.elliptical.GaussianCopula(dim=len(TICKERS))
gauss_cop.fit(pseudo_obs.values, to_pobs=False)

gauss_corr = pd.DataFrame(gauss_cop.sigma, index=TICKERS, columns=TICKERS)
gauss_loglik = gauss_cop.log_likelihood(pseudo_obs.values, to_pobs=False)
gauss_num_params = len(gauss_cop.params)

print("GAUSSIAN COPULA\n")
print(gauss_corr.round(4))
print(f"\nLog-Likelihood: {gauss_loglik:.4f}")
print(f"Number of Parameters: {gauss_num_params}")

# %%
d = len(TICKERS)

# Fit Student-t Copula
t_cop = copulae.StudentCopula(dim=d)
t_cop.fit(pseudo_obs.values, to_pobs=False)

print("STUDENT-T COPULA\n")

nu = t_cop.params[0]

print("Correlation Matrix sigma:")
corr_matrix = t_cop.sigma
print(pd.DataFrame(corr_matrix, index=TICKERS, columns=TICKERS).round(4))

t_loglik = t_cop.log_likelihood(pseudo_obs.values, to_pobs=False)
t_num_params = d * (d - 1) // 2 + 1

print(f"\nLog-Likelihood: {t_loglik:.4f}")
print(f"Number of Parameters: {t_num_params}")
print(f"Degrees of Freedom nu: {nu:.4f}")

# %% [markdown]
# ##### Explanation of the fitting procedure
#
# We first transform the returns into **pseudo-observations**
# $$#
# U_{t,i} = \frac{\text{rank}(R_{t,i})}{W+1}
# $$#
# so each marginal is approximately uniform on  $[0, 1]$ .
#
# This allows us to fit the copulas **only on the dependence structure**.
#
# Both copulas are then fitted using **maximum likelihood**: the pseudo-observations are transformed using the appropriate inverse CDF, and the log-likelihood is maximized over the copula parameters.

```

```

#
# **Gaussian copula.**
# Maximizes the likelihood over the correlation matrix  $\Sigma$  only, after applying the
# inverse normal CDF.
#
# **Student-t copula.**
# Same procedure as the Gaussian copula, but with one additional parameter: the **degrees of
# freedom**  $\nu$ , which controls tail dependence.
#
# Thus, both models share the same fitting framework. The Student-t copula simply includes
# an extra parameter to capture heavier joint tails.

# %% [markdown]
# ### c. Simulate from copulas and compare with original returns
#
# We simulate from the fitted copulas and map back to the original return scale using
# empirical marginals.

# %% [markdown]
# 1. **Generate uniform samples (dependence only)**
#
# From the fitted copula  $C_\theta$ , we simulate
# $$U^{\{t\}} = (U^{\{t\}}_1, \dots, U^{\{t\}}_d) \sim C_\theta, \quad t = 1, \dots, T,$$
# where  $U^{\{t\}}_i \in (0, 1)$ .
# These uniforms encode the **dependence structure** of the Gaussian or Student-t copula.
#
# 2. **Apply the inverse empirical CDFs (recover marginals)**
#
# For each asset  $i$ , let  $\hat{F}_i^{-1}$  denote the empirical quantile function built
# from the observed returns.
# Each uniform component is transformed via
# $$X^{\{t\}}_i = \hat{F}_i^{-1}(U^{\{t\}}_i),$$
# which ensures the simulated data have the **same marginal distribution** as the
# original returns.
#
# 3. **Construct simulated return series**
#
# Repeating this for all assets and all samples produces two synthetic datasets,
# $$
# X^{\{t\}}_{-t} \quad \text{and} \quad X^{\{t\}}_{-t}
# $$
# each with the original empirical marginals but with dependence dictated by the fitted
# Gaussian or Student-t copula.

# %%
T = len(log_returns)
print(f"Simulating {T} samples from each fitted copula...")

# Sample  $U \sim \text{Copula}$ 
U_gauss = gauss_cop.random(T)
U_t = t_cop.random(T)

# %%
def inverse_empirical(u, data):
    """
    Strictly monotonic inverse-ECDF using interpolation

```

```

uuuuMaps uniform U to real observed distribution of data.
uuuu"""
data_sorted = np.sort(data)
ranks = (np.arange(1, len(data_sorted) + 1)) / (len(data_sorted) + 1)
return np.interp(u, ranks, data_sorted)

# %%
sim_gauss = pd.DataFrame({
    ticker: inverse_empirical(U_gauss[:, i], log_returns[ticker].values)
    for i, ticker in enumerate(TICKERS)
})
sim_t = pd.DataFrame({
    ticker: inverse_empirical(U_t[:, i], log_returns[ticker].values)
    for i, ticker in enumerate(TICKERS)
})

print("\nSimulated_returns_(Gaussian):")
display(sim_gauss.head())

print("\nSimulated_returns_(Student-t):")
display(sim_t.head())

# %%
fig, axes = plt.subplots(
    2, len(TICKERS), figsize=(5 * len(TICKERS), 8), squeeze=False
)

for j, tk in enumerate(TICKERS):

    # Shared bins across original + both simulations
    pooled = np.concatenate([log_returns[tk], sim_gauss[tk], sim_t[tk]])
    bins = np.histogram_bin_edges(pooled, bins='auto')

    # ---- Row 1: Original vs Gaussian ----
    ax = axes[0, j]
    ax.hist(log_returns[tk], bins=bins, density=True, alpha=0.5,
            label="Original", edgecolor="k", color="tab:blue")
    ax.hist(sim_gauss[tk], bins=bins, density=True, alpha=0.5,
            label="Gaussian_copula", edgecolor="k", color="tab:orange")
    ax.set_title(f"{tk} _Original_vs_Gaussian")
    ax.set_xlabel("Return")
    ax.set_ylabel("Density")
    ax.grid(True, alpha=0.3)

    # ---- Row 2: Original vs Student-t ----
    ax = axes[1, j]
    ax.hist(log_returns[tk], bins=bins, density=True, alpha=0.5,
            label="Original", edgecolor="k")
    ax.hist(sim_t[tk], bins=bins, density=True, alpha=0.5,
            label="Student-t_copula", edgecolor="k", color="tab:olive")
    ax.set_title(f"{tk} _Original_vs_Student-t")
    ax.set_xlabel("Return")
    ax.set_ylabel("Density")
    ax.grid(True, alpha=0.3)

# Define legend patches manually
original_patch = mpatches.Patch(facecolor="tab:blue", alpha=0.5, label="Original", edgecolor="k")
gaussian_patch = mpatches.Patch(facecolor="tab:orange", alpha=0.5, label="Gaussian_copula",
                                edgecolor="k")
student_t_patch = mpatches.Patch(facecolor="tab:olive", alpha=0.5, label="Student-t_copula",

```

```

    edgecolor="k")

fig.legend(
    handles=[original_patch, gaussian_patch, student_t_patch],
    loc="upper_center", ncols=3, frameon=False
)

plt.tight_layout(rect=[0, 0, 1, 0.92])
plt.savefig(os.path.join(OUT_DIR, 'copula_marginal_distributions_split.png'))
plt.show()

# %% [markdown]
# The histograms show that both simulated datasets closely match the original return
distributions, as expected from using empirical quantile mapping.

# %%
# Dependence structure
pairs = list(combinations(TICKERS, 2))
fig, axes = plt.subplots(3, len(pairs), figsize=(6 * len(pairs), 14), squeeze=False)

for col, (a, b) in enumerate(pairs):
    # 1) original
    ax = axes[0, col]
    ax.scatter(log_returns[a], log_returns[b], alpha=0.5, s=15)
    ax.set_title(f"Original:{a} vs {b}")
    ax.grid(True, alpha=0.3)

    # 2) Gaussian
    ax = axes[1, col]
    ax.scatter(sim_gauss[a], sim_gauss[b], alpha=0.5, s=15)
    ax.set_title(f"Gaussian_Sim:{a} vs {b}")
    ax.grid(True, alpha=0.3)

    # 3) Student-t
    ax = axes[2, col]
    ax.scatter(sim_t[a], sim_t[b], alpha=0.5, s=15)
    ax.set_title(f"Student-t_Sim:{a} vs {b}")
    ax.grid(True, alpha=0.3)

# annotate rows
axes[0,0].set_ylabel("Original")
axes[1,0].set_ylabel("Gaussian")
axes[2,0].set_ylabel("Student-t")

plt.tight_layout()
plt.savefig(os.path.join(OUT_DIR, 'copula_dependence_structure.png'))
plt.show()

# %% [markdown]
# The scatter plots indicate that both copulas reproduce the overall dependence structure.
# The Student-t copula generates slightly more joint extremes, but the differences remain
modest given the estimated degrees of freedom.

# %%
# Align simulated data index with original returns
sim_gauss.index = log_returns.index[:len(sim_gauss)]
sim_t.index      = log_returns.index[:len(sim_t)]

# Consistent colors with your histogram
COL_ORIG = "tab:blue"
COL_GAUS = "tab:orange"

```

```

COL_T    = "tab:olive"

fig, axes = plt.subplots(
    2, len(TICKERS), figsize=(5 * len(TICKERS), 8),
    squeeze=False, sharex='col'
)

for j, tk in enumerate(TICKERS):

    # ----- Row 1: Original vs Gaussian -----
    ax = axes[0, j]
    ax.plot(log_returns.index, log_returns[tk],
            lw=1, label="Original", color=COL_ORIG)
    ax.plot(sim_gauss.index, sim_gauss[tk],
            lw=1, label="Gaussian\u2225copula", color=COL_GAUS)
    ax.set_title(f"\{tk}\u2225Original\u2225vs\u2225Gaussian\u2225(Time\u2225Series)")
    ax.set_ylabel("Return")
    ax.grid(True, alpha=0.3)
    # no x-label on top row

    # ----- Row 2: Original vs Student-t -----
    ax = axes[1, j]
    ax.plot(log_returns.index, log_returns[tk],
            lw=1, label="Original", color=COL_ORIG)
    ax.plot(sim_t.index, sim_t[tk],
            lw=1, label="Student-t\u2225copula", color=COL_T)
    ax.set_title(f"\{tk}\u2225Original\u2225vs\u2225Student-t\u2225(Time\u2225Series)")
    ax.set_xlabel("Time")
    ax.set_ylabel("Return")
    ax.grid(True, alpha=0.3)
    ax.tick_params(axis="x", rotation=45) # nicer date labels

handles1, labels1 = axes[0, 0].get_legend_handles_labels()
handles2, labels2 = axes[1, 0].get_legend_handles_labels()

handles = list(handles1)
labels = list(labels1)
for h, lab in zip(handles2, labels2):
    if lab not in labels:
        handles.append(h)
        labels.append(lab)

fig.legend(handles, labels,
           loc="upper\u2225center", ncols=3, frameon=False)

plt.tight_layout(rect=[0, 0, 1, 0.92])
plt.savefig(os.path.join(OUT_DIR, "copula_time_series_comparison.png"))
plt.show()

# %% [markdown]
# The plots above compare the original return series with returns simulated using Gaussian
# and Student-t copulas for AAPL, META, and JPM.
#
# Key observations:
#
# - **Overall return dynamics:** All series fluctuate around zero with frequent sign changes
#   , consistent with daily stock return behavior.
#
# - **Volatility clustering:** The **original data** exhibit periods where large movements
#   are followed by other large movements (both positive and negative), particularly visible
#   in META and AAPL around early 2024. Both copula simulations **partially replicate**

```

this phenomenon, but the effect is **more pronounced in the Student-t copula**, where clusters of large returns appear more frequently, especially for META.

```

#
# - **Gaussian copula behavior:** The Gaussian generated series tracks the general variability of the original returns but displays **smoother fluctuations**. Large spikes are **less frequent and less intense**, indicating that the Gaussian copula **underestimates tail events** and does not fully capture periods of heightened volatility .
.

#
# - **Student-t copula behavior:** The Student-t simulated paths show **higher amplitude jumps and more erratic bursts**, aligning more closely with the original heavy-tailed structure. This copula better reproduces **extreme shocks and volatility clustering**, particularly visible in META and JPM, where large returns tend to appear in temporal clusters.

#
# While both copulas mimic the central fluctuations of the return series, **the Student-t copula provides a more realistic replication of extreme events and volatility clustering **, making it more suitable for modeling financial return dynamics and risk.

# %% [markdown]
# ## 5) Portfolio VaR/ES with copulas (rolling)

# %% [markdown]
# Equal-weighted portfolio of AAPL, MSFT, JPM. Compare univariate models (as in Exercise 3) vs copula-based VaR/ES with rolling windows.

#
# At each time, fit copulas on last W weeks, simulate N scenarios, estimate VaR/ES from simulated portfolio returns, then backtest.

# %% [markdown]
# ### a. Univariate models on portfolio returns
#
# Apply the same backtesting procedure as in question 3, but now on equal-weighted portfolio returns.

#
# Compute equal-weighted portfolio returns
# Portfolio return = (1/3) * R_AAPL + (1/3) * R_META + (1/3) * R_JPM
weights = np.array([1/3, 1/3, 1/3])

portfolio_returns = (log_returns[TICKERS] * weights).sum(axis=1)

print("Equal-weighted_portfolio_returns:")
print(f"Shape:{portfolio_returns.shape}")
print(f"Mean:{portfolio_returns.mean():.6f}")
print(f"Std_Dev:{portfolio_returns.std():.6f}")
print(f"Min:{portfolio_returns.min():.6f}")
print(f"Max:{portfolio_returns.max():.6f}")
print("\nFirst5_portfolio_returns:")
print(portfolio_returns.head())

#
# Apply rolling-window forecasts to portfolio returns (same as question 3)
print("Computing_rolling_VaR/ES_for_portfolio_using_univariate_models...")
print(f"Window_size:{WINDOW}days")
print(f"Confidence_levels:{CONFIDENCE}")

METHODS = {
    "historical": historical_simulation,
    "gaussian": gaussian_normal_distribution,
    "student": student_var_es,
```

```

"parametric": conditional_parametric_var_es,
"lhs": filtered_historical_simulation
}

L_portfolio = -portfolio_returns
portfolio_rolling_results = rolling_forecast(L_portfolio, WINDOW, METHODS, CONFIDENCE)

print(f"\nRolling forecasts complete")
print(f"Results shape: {portfolio_rolling_results.shape}")
print(f"Number of forecasts: {len(portfolio_rolling_results.dropna())}")

# %%
# Backtest portfolio VaR using Kupiec POF test
print("-"*60)
print("PORTFOLIO_VaR_BACKTESTING_CHRISTOFFERSEN_CONDITIONAL_COVERAGE_TEST_KUPIEC_POF_TEST")
print("-"*60)

result = pd.DataFrame(columns = ['Confidence_Level', 'Method', 'POF_p_value', 'Christoffersen_p_value'])
for confidence in CONFIDENCE:
    alpha = 1 - confidence
    print(f"\n{'-'*60}")
    print(f"Confidence_Level: {int(confidence*100)}%")
    print(f"{'-'*60}\n")

    for method in METHODS.keys():
        VaR_portfolio = portfolio_rolling_results[method]['VaR'][confidence]
        print(f'\nPof test for Portfolio_{method} method: ')
        POF = kupiec_pof_test(VaR_portfolio, L_portfolio, alpha)
        LR_ind = christoffersen(VaR_portfolio, L_portfolio)
        result = pd.concat([
            result,
            pd.DataFrame([{
                'Confidence_Level': f'{int(confidence*100)}%',
                'Method': method,
                'POF_p_value': chi_square_test(POF, df=1, method_name=f'Portfolio_{method}', test_name='POF test'),
                'Christoffersen_p_value': chi_square_test(LR_ind + POF, df=2, method_name=f'Portfolio_{method}', test_name='Christoffersen conditional coverage test')
            }])
        ], ignore_index=True)

# %%
print("-"*60)
print("PORTFOLIO_VaR_ES_BACKTESTING")
print("-"*60)

# Result table: add ES columns
result = pd.DataFrame(columns=[
    'Confidence_Level',
    'Method',
    'POF_p_value',
    'Christoffersen_p_value',
    'ES_Z1_stat',
    'ES_p_value'
])

portfolio_loss = L_portfolio.sort_index()

```

```

for confidence in CONFIDENCE:
    alpha = 1 - confidence
    print(f"\n{'-'*60}")
    print(f"Confidence_Level:{int(confidence*100)}%")
    print(f"{'-'*60}\n")

for method in METHODS.keys():
    # ----- VaR backtests -----
    VaR_portfolio = portfolio_rolling_results[method] ['VaR'] [confidence]

    print(f'\nPof & Christoffersen tests for Portfolio-{method} method:')
    POF = kupiec_pof_test(VaR_portfolio, L_portfolio, alpha)
    LR_ind = christoffersen(VaR_portfolio, L_portfolio)

    p_pof = chi_square_test(
        POF, df=1,
        method_name=f'Portfolio-{method}',
        test_name='POF test'
    )
    p_christoffersen = chi_square_test(
        LR_ind + POF, df=2,
        method_name=f'Portfolio-{method}',
        test_name='Christoffersen conditional coverage test'
    )

    # ----- ES backtest (AcerbiSzekely Z1) -----
    ES_portfolio = portfolio_rolling_results[method] ['ES'] [confidence]

    Z1_simulated = monte_carlo_es_backtest_z1(
        VaR_series=VaR_portfolio,
        ES_series=ES_portfolio,
        Loss_series=portfolio_loss,
        sampler=sampler[method],
        window=WINDOW,
        M=1000
    )

    real_Z1 = es_backtest_z1(
        VaR_series=VaR_portfolio,
        ES_series=ES_portfolio,
        Loss_series=portfolio_loss
    )

    p_es = (Z1_simulated >= real_Z1).sum() / len(Z1_simulated)

    result = pd.concat([
        result,
        pd.DataFrame([{
            'Confidence_Level': f'{int(confidence*100)}',
            'Method': method,
            'POF_p_value': p_pof,
            'Christoffersen_p_value': p_christoffersen,
            'ES_Z1_stat': real_Z1,
            'ES_p_value': p_es
        }])
    ], ignore_index=True)

summary = result.set_index(['Confidence_Level', 'Method'])[
    ['POF_p_value', 'Christoffersen_p_value', 'ES_Z1_stat', 'ES_p_value']
]

```

```

# %%
summary

# %% [markdown]
# ### b. Copula-based portfolio VaR/ES (rolling)
#
# At each time  $t$ , fit the copulas as in 4.b) on the last  $W$  days of portfolio component
# returns. Simulate  $N = 1000$  returns from each fitted copula, then estimate VaR/ES from
# the simulated portfolio returns.

# %%
def copula_portfolio_var_es(returns_df, tickers, weights, W, N_sim=1000, confidence=
    CONFIDENCE):
    """
    Rolling copula-based VaR/ES estimation for portfolio.

    For each time  $t$ , fit Gaussian and Student-t copulas on the last  $W$  days of returns,
    simulate  $N_{sim}$  scenarios, construct portfolio returns, and compute VaR/ES.

    Parameters:
        returns_df: pd.DataFrame
            Log returns for all assets
        tickers: list
            List of ticker symbols
        weights: np.array
            Portfolio weights (should sum to 1)
        W: int
            Rolling window size
        N_sim: int
            Number of Monte Carlo simulations
        alphas: list
            Confidence levels

    Returns:
        results: dict
            Dictionary with 'gaussian' and 'student' keys, each containing VaR and ES
            forecasts
    """
    T = len(returns_df)

    # Initialize result storage
    results = {
        'gaussian': {
            'VaR': {a: pd.Series(index=returns_df.index[W:], dtype=float) for a in confidence},
            'ES': {a: pd.Series(index=returns_df.index[W:], dtype=float) for a in confidence}
        },
        'student': {
            'VaR': {a: pd.Series(index=returns_df.index[W:], dtype=float) for a in confidence},
            'ES': {a: pd.Series(index=returns_df.index[W:], dtype=float) for a in confidence}
        }
    }

    # Rolling window estimation
    for t in range(W, T):
        # Extract window data
        window_returns = returns_df.iloc[t-W:t][tickers].copy()
        pseudo_obs = compute_pseudo_obs(window_returns)

        # Fit Gaussian copula

```

```

gauss_cop = copulae.elliptical.GaussianCopula(dim=len(tickers))
gauss_cop.fit(pseudo_obs.values, to_pobs=False)

# Fit Student-t copula
t_cop = copulae.StudentCopula(dim=len(tickers))
t_cop.fit(pseudo_obs.values, to_pobs=False)

# Simulate from copulas
U_gauss = gauss_cop.random(N_sim)
U_t = t_cop.random(N_sim)

# Convert copula samples to returns
sim_returns_gauss = np.zeros((N_sim, len(tickers)))
sim_returns_t = np.zeros((N_sim, len(tickers)))

for i, ticker in enumerate(tickers):
    sim_returns_gauss[:, i] = inverse_empirical(U_gauss[:, i], window_returns[ticker].values)
    sim_returns_t[:, i] = inverse_empirical(U_t[:, i], window_returns[ticker].values)

# Compute portfolio returns and losses
portfolio_returns_gauss = sim_returns_gauss @ weights
portfolio_returns_t = sim_returns_t @ weights

portfolio_losses_gauss = -portfolio_returns_gauss
portfolio_losses_t = -portfolio_returns_t

# Compute VaR and ES for each confidence level
current_date = returns_df.index[t]

for a in confidence:
    # Gaussian copula
    VaR_gauss = np.percentile(portfolio_losses_gauss, a * 100)
    ES_gauss = portfolio_losses_gauss[portfolio_losses_gauss >= VaR_gauss].mean()
    results['gaussian']['VaR'][a].loc[current_date] = VaR_gauss
    results['gaussian']['ES'][a].loc[current_date] = ES_gauss

    # Student-t copula
    VaR_t = np.percentile(portfolio_losses_t, a * 100)
    ES_t = portfolio_losses_t[portfolio_losses_t >= VaR_t].mean()
    results['student']['VaR'][a].loc[current_date] = VaR_t
    results['student']['ES'][a].loc[current_date] = ES_t

return results

# %%
def copula_sampler(L_window, M, type='gaussian'):
    """
    L_window: DataFrame or 2D array of shape (T_window, d)
    rows = time, columns = assets (returns)
    M: number of simulated scenarios
    type: 'gaussian' or 'student'
    """
    L_window = pd.DataFrame(L_window)
    d = L_window.shape[1]

    # Pseudo-observations
    U = compute_pseudo_obs(L_window)

    if type == 'gaussian':
        cop = copulae.elliptical.GaussianCopula(dim=d)

```

```

else:
    cop = copulae.StudentCopula(dim=d)

cop.fit(U.values, to_pobs=False)

# Simulate U
U_sim = cop.random(M)

# Map each marginal back
sim_returns = np.zeros((M, d))
for i in range(d):
    sim_returns[:, i] = inverse_empirical(U_sim[:, i], L_window.iloc[:, i].values)

# Compute portfolio returns and losses
portfolio_returns = sim_returns @ weights
portfolio_losses = -portfolio_returns

return portfolio_losses

# %%
def monte_carlo_es_backtest_z1_copula(VaR_series, ES_series, portfolio_loss, asset_returns,
    window, copula_type, M=1000):
    """
    VaR_series, ES_series: rolling forecasts (aligned with portfolio_loss)
    portfolio_loss: Series of portfolio losses (same index as VaR/ES)
    asset_returns: DataFrame of asset returns (same index as portfolio_loss)
    window: rolling window length used to estimate VaR/ES
    copula_type: 'gaussian' or 'student'
    M: number of Monte Carlo paths
    """
    x = np.asarray(portfolio_loss, dtype=float)
    n = len(x)

    T = n - window + 1
    idx = portfolio_loss.index[window-1:]

    sims = np.empty((M, T), dtype=float)

    for t in range(T):
        # Window of asset returns used for estimation at time t
        L_window = asset_returns.iloc[t:t+window, :]
        draws = copula_sampler(L_window, M, type=copula_type)
        sims[:, t] = draws

    VaR = VaR_series.loc[idx]
    ES = ES_series.loc[idx]

    Z1 = np.empty(M, dtype=float)
    for m in range(M):
        sim_series = pd.Series(sims[m], index=idx)
        Z1[m] = es_backtest_z1(VaR, ES, sim_series)

    return Z1

# %%
def monte_carlo_es_backtest_z1_copula(
    VaR_series,
    ES_series,
    portfolio_loss,
    asset_returns,
    window,

```

```

copula_type,
M=1000
):
"""
    VaR_series, ES_series:pandas.Series
    Out-of-sample VaR/ES forecasts, same index (e.g. from rolling_forecast).
    portfolio_loss:pandas.Series
    Full portfolio loss series (in-sample + out-of-sample), NOT truncated.
    asset_returns:pandas.DataFrame
    Asset returns used to build the copula, same (or superset) index as portfolio_loss.
    window:int
    Rolling window length used to estimate VaR/ES and the copula.
    copula_type:str
    'gaussian' or 'student' (or any type supported by copula_sampler).
    M:int
    Number of Monte Carlo paths.

Returns
-----
Z1: ndarray of shape (M,)
Monte Carlo distribution of the ES backtest statistic Z1.
"""

# 1) Align and sort VaR/ES
VaR = VaR_series.sort_index()
ES = ES_series.reindex(VaR.index) # ensure same index/order as VaR

# 2) Sort portfolio losses and asset returns
portfolio_loss = portfolio_loss.sort_index()
asset_returns = asset_returns.sort_index()

loss_index = portfolio_loss.index
ret_index = asset_returns.index

# Out-of-sample dates (where VaR/ES are defined)
idx = VaR.index
T = len(idx)

sims = np.empty((M, T), dtype=float)

# 3) For each OOS date, reconstruct the window used to estimate VaR/ES & copula
for j, date in enumerate(idx):
    # Position of this date in the full series
    loc_ret = ret_index.get_loc(date)

    # Window of length 'window' ending at 'date'
    # (assumes VaR/ES start only after at least 'window' observations)
    L_window = asset_returns.iloc[loc_ret - window + 1 : loc_ret + 1, :]

    # Draw portfolio losses from the copula-based sampler
    draws = copula_sampler(L_window, M, type=copula_type)
    draws = np.asarray(draws, dtype=float).reshape(M,) # safety reshape
    sims[:, j] = draws

# 4) Compute Z1 for each simulated path
Z1 = np.empty(M, dtype=float)
for m in range(M):
    sim_series = pd.Series(sims[m], index=idx)
    Z1[m] = es_backtest_z1(VaR, ES, sim_series)

return Z1

```

```

# %%
# Compute copula-based rolling VaR/ES for portfolio
print("Computing copula-based rolling VaR/ES for portfolio...")
print(f"Window size: {WINDOW} days")
print(f"Monte Carlo simulations: N=1000")
print(f"Confidence levels: {CONFIDENCE}")
print(f"Portfolio weights: {weights}")

copula_results = copula_portfolio_var_es(
    returns_df=log_returns,
    tickers=TICKERS,
    weights=weights,
    W=WINDOW,
    N_sim=1000,
    confidence=CONFIDENCE
)

print(f"\nCopula-based rolling forecasts complete")
print(f"Number of Gaussian VaR forecasts (alpha=0.05): {copula_results['gaussian']['VaR'] [0.95].notna().sum()}")
print(f"Number of Student-t VaR forecasts (alpha=0.05): {copula_results['student']['VaR'] [0.95].notna().sum()}\n")

# %% [markdown]
# ### c. Backtesting and comparison: univariate vs copula-based approaches
#
# Compare portfolio VaR/ES backtests across univariate methods (part a) and copula-based approaches (part b). Assess whether explicit dependence modeling improves accuracy, especially in the tails.

# %%
# Backtest copula-based VaR using Kupiec POF test
print("-"*60)
print("COPULA-BASED PORTFOLIO VaR BACKTESTING - KUPIEC POF TEST")
print("-"*60)

result = pd.DataFrame(columns = ['Confidence_Level', 'Method', 'POF_p_value'])
for confidence in CONFIDENCE:
    alpha = 1 - confidence
    print(f"\n{'-'*60}")
    print(f"Confidence Level: {int(confidence*100)}%")
    print(f"{'-'*60}\n")

    # Gaussian copula
    VaR_gauss = copula_results['gaussian']['VaR'][confidence]
    POF_gauss = kupiec_pof_test(VaR_gauss, L_portfolio, alpha)
    chi_square_test(POF_gauss, df=1, method_name='Portfolio - Gaussian Copula', test_name='POF test')

    # Student-t copula
    VaR_t = copula_results['student']['VaR'][confidence]
    POF_t = kupiec_pof_test(VaR_t, L_portfolio, alpha)
    chi_square_test(POF_t, df=1, method_name='Portfolio - Student-t Copula', test_name='POF test')

# %%
# Backtest copula-based VaR using Christoffersen independence test
print("-"*60)
print("COPULA-BASED PORTFOLIO VaR BACKTESTING - CHRISTOFFERSEN CONDITIONAL COVERAGE TEST")
print("-"*60)

```

```

for confidence in CONFIDENCE:
    alpha = 1 - confidence
    print(f"\n{'-'*60}")
    print(f"Confidence_Level:{int(confidence*100)}%")
    print(f"{'-'*60}\n")

    # Gaussian copula
    VaR_gauss = copula_results['gaussian']['VaR'][confidence]
    POF_gauss = kupiec_pof_test(VaR_gauss, L_portfolio, alpha)
    LR_ind_gauss = christoffersen(VaR_gauss, L_portfolio)
    chi_square_test(LR_ind_gauss + POF_gauss, df=2, method_name='Portfolio-Gaussian_Copula',
                    test_name='Christoffersen_conditional_coverage_test')

    # Student-t copula
    VaR_t = copula_results['student']['VaR'][confidence]
    POF_t = kupiec_pof_test(VaR_t, L_portfolio, alpha)
    LR_ind_t = christoffersen(VaR_t, L_portfolio)
    chi_square_test(LR_ind_t + POF_t, df=2, method_name='Portfolio-Student-t_Copula',
                    test_name='Christoffersen_conditional_coverage_test')

# %%
print("-"*60)
print("COPULA-BASED PORTFOLIO ES BACKTESTING (Acerbi-Szekely_Z1_Test)")
print("-"*60)

asset_returns = log_returns[TICKERS]      # full asset returns
portfolio_loss = L_portfolio # align with rolling VaR/ES

for confidence in CONFIDENCE:
    alpha = 1 - confidence
    print(f"\n{'-'*60}")
    print(f"Confidence_Level:{int(confidence*100)}%")
    print(f"{'-'*60}\n")

    # Align VaR/ES first
    VaR_gauss = copula_results['gaussian']['VaR'][confidence]
    ES_gauss = copula_results['gaussian']['ES'][confidence]

    # Gaussian copula ES backtest
    Z_1_sim_gauss = monte_carlo_es_backtest_z1_copula(
        VaR_gauss, ES_gauss, portfolio_loss, asset_returns, window=WINDOW,
        copula_type='gaussian', M=1000
    )
    real_Z1_gauss = es_backtest_z1(VaR_gauss, ES_gauss, portfolio_loss)
    p_gauss = (np.sum(Z_1_sim_gauss >= real_Z1_gauss) + 1) / (len(Z_1_sim_gauss) + 1)

    print(f"The test statistic and p_value for Portfolio-Gaussian_Copula are: "
          f"Z1_stat:{real_Z1_gauss}, p_value:{p_gauss}")
    if p_gauss < 0.05:
        print("Reject the null hypothesis at 5% significance level for Portfolio-Gaussian_Copula.\n")
    else:
        print("Fail to reject the null hypothesis at 5% significance level for Portfolio-Gaussian_Copula.\n")

    # Student-t copula ES backtest
    VaR_student = copula_results['student']['VaR'][confidence]
    ES_student = copula_results['student']['ES'][confidence]

    Z_1_sim_student = monte_carlo_es_backtest_z1_copula(

```

```

    VaR_student, ES_student, portfolio_loss, asset_returns, window=WINDOW,
    copula_type='student', M=1000
)
real_Z1_student = es_backtest_z1(VaR_student, ES_student, portfolio_loss)
p_student = (np.sum(Z_1_sim_student >= real_Z1_student) + 1) / (len(Z_1_sim_student) +
1)

print(f"The test statistic and p-value for Portfolio - Student Copula are: "
      f"Z1_stat:{real_Z1_student}, p_value:{p_student}")
if p_student < 0.05:
    print("Reject the null hypothesis at 5% significance level for Portfolio - Student "
          "Copula.\n")
else:
    print("Fail to reject the null hypothesis at 5% significance level for Portfolio - "
          "Student Copula.\n")

# %%
def compute_backtest_statistics(VaR_forecasts,
                                ES_forecasts,
                                actual_losses,
                                alpha,
                                sampler_fn,
                                window,
                                M=1000):
    """
    Compute VaR and ES backtest statistics for a given method.
    """

# ----- VaR-based statistics -----
valid_idx = VaR_forecasts.notna() & actual_losses.notna()
VaR_clean = VaR_forecasts[valid_idx]
L_clean = actual_losses[valid_idx]

violations = (L_clean > VaR_clean).astype(int)
n_violations = violations.sum()
n_total = len(violations)
violation_rate = n_violations / n_total if n_total > 0 else np.nan
expected_violations = alpha * n_total

# Kupiec POF
POF = kupiec_pof_test(VaR_clean, L_clean, alpha)

# Christoffersen independence component
LR_ind = christoffersen(VaR_clean, L_clean)

# ----- NEW: p-values -----
POF_p_value = chi_square_test(
    POF, df=1,
    method_name='Portfolio',
    test_name='POF test'
)

Christoffersen_p_value = chi_square_test(
    LR_ind + POF, df=2,
    method_name='Portfolio',
    test_name='Christoffersen test'
)

# ----- ES-based statistics -----
Z1_simulated = monte_carlo_es_backtest_z1(
    VaR_forecasts, ES_forecasts, actual_losses,

```

```

        sampler_fn, window=window, M=M
    )
Z1_real = es_backtest_z1(VaR_forecasts, ES_forecasts, actual_losses)
Z1_p_value = (Z1_simulated >= Z1_real).sum() / len(Z1_simulated)

return {
    "n_violations": n_violations,
    "n_total": n_total,
    "violation_rate": violation_rate,
    "expected_rate": alpha,
    "expected_violations": expected_violations,
    "POF_statistic": POF,
    "POF_p_value": POF_p_value, # NEW
    "Christoffersen_statistic": LR_ind + POF,
    "Christoffersen_p_value": Christoffersen_p_value, # NEW
    "Z1_stat": Z1_real,
    "Z1_p_value": Z1_p_value,
}
}

# %%
print("-"*60)
print("COMPREHENSIVE COMPARISON: UNIVARIATE vs COPULA-BASED APPROACHES")
print("-"*60)

comparison_results = []

# Full asset returns and portfolio loss (for copula ES MC)
asset_returns = log_returns[TICKERS]
portfolio_loss = L_portfolio.sort_index()

# Choose a sampler for the VaR null for copulas (only used inside
compute_backtest_statistics for ES,
# but we will override ES, so this can be anything reasonable, e.g. historical)
sampler_for_VaR = sampler['historical']

for confidence in CONFIDENCE:
    alpha = 1 - confidence

    # ----- Univariate methods -----
    for method in METHODS:
        VaR_forecast = portfolio_rolling_results[method] ['VaR'][confidence]
        ES_forecast = portfolio_rolling_results[method] ['ES'][confidence]

        stats = compute_backtest_statistics(
            VaR_forecast,
            ES_forecast,
            portfolio_loss,
            alpha,
            sampler_fn=sampler[method],
            window=WINDOW,
            M=1000
        )

        comparison_results.append({
            "Confidence": f"{int(confidence*100)}%",
            "Approach": "Univariate",
            "Method": method,
            "Violations": stats["n_violations"],
            "Total": stats["n_total"],
            "Violation_Rate": f"{stats['violation_rate']:.4f}",
            "Z1_p_value": Z1_p_value,
            "Christoffersen_p_value": Christoffersen_p_value,
            "POF_p_value": POF_p_value
        })
    
```

```

    "Expected_Rate": f"{{stats['expected_rate']:.4f}}",
    "POF_Stat": f"{{stats['POF_statistic']:.4f}}",
    "POF_p-value": f"{{stats['POF_p_value']:.4f}}",
    "Christoffersen_Stat": f"{{stats['Christoffersen_statistic']:.4f}}",
    "Christoffersen_p-value": f"{{stats['Christoffersen_p_value']:.4f}}",
    "Z1_Stat": f"{{stats['Z1_stat']:.4f}}",
    "Z1_p-value": f"{{stats['Z1_p_value']:.4f}}",
}

# ----- Copula-based methods -----
for copula_type in ["gaussian", "student"]:
    VaR_forecast = copula_results[copula_type]['VaR'][confidence]
    ES_forecast = copula_results[copula_type]['ES'][confidence]

    # 1) VaR-based statistics from the generic function
    stats = compute_backtest_statistics(
        VaR_forecast,
        ES_forecast,
        portfolio_loss,
        alpha,
        sampler_fn=sampler_for_VaR, # only used for ES in generic MC, but we'll override
        ES
        window=WINDOW,
        M=1000
    )

    # 2) Override ES statistics using copula-based Monte Carlo
    Z1_sim = monte_carlo_es_backtest_z1_copula(
        VaR_forecast,
        ES_forecast,
        portfolio_loss,
        asset_returns,
        window=WINDOW,
        copula_type=copula_type,
        M=1000
    )
    Z1_real = es_backtest_z1(VaR_forecast, ES_forecast, portfolio_loss)
    Z1_p_value = (Z1_sim >= Z1_real).sum() / len(Z1_sim)

    stats["Z1_stat"] = Z1_real
    stats["Z1_p_value"] = Z1_p_value

    copula_name = "Gaussian_Copula" if copula_type == "gaussian" else "Student-t_Copula"

    comparison_results.append({
        "Confidence": f"{{int(confidence*100)}}",
        "Approach": "Copula",
        "Method": copula_name,
        "Violations": stats["n_violations"],
        "Total": stats["n_total"],
        "Violation_Rate": f"{{stats['violation_rate']:.4f}}",
        "Expected_Rate": f"{{stats['expected_rate']:.4f}}",
        "POF_Stat": f"{{stats['POF_statistic']:.4f}}",
        "POF_p-value": f"{{stats['POF_p_value']:.4f}}",
        "Christoffersen_Stat": f"{{stats['Christoffersen_statistic']:.4f}}",
        "Christoffersen_p-value": f"{{stats['Christoffersen_p_value']:.4f}}",
        "Z1_Stat": f"{{stats['Z1_stat']:.4f}}",
        "Z1_p-value": f"{{stats['Z1_p_value']:.4f}}",
    })
}

```

```

# Create DataFrame for better visualization
comparison_df = pd.DataFrame(comparison_results)

print("\nSUMMARY_TABLE:")
display(comparison_df)

# Save to CSV
comparison_df.to_csv(OUT_DIR / 'portfolio_var_es_comparison.csv', index=False)
print(f"\nComparison_table saved to {OUT_DIR / 'portfolio_var_es_comparison.csv'}")

# %%
# Visualize violation rates comparison
fig, axes = plt.subplots(1, 2, figsize=(16, 6))

for i, confidence in enumerate(CONFIDENCE):
    ax = axes[i]
    alpha = 1 - confidence

    # Filter data for this confidence level
    data = comparison_df[comparison_df['Confidence'] == f"{int(confidence*100)}%"].copy()

    # Convert violation rate to float for plotting
    data['Violation_Rate_Numeric'] = data['Violation_Rate'].astype(float)
    data['Method_Full'] = data['Approach'] + '-' + data['Method']

    # Create bar plot
    x_pos = np.arange(len(data))
    bars = ax.bar(x_pos, data['Violation_Rate_Numeric'], alpha=0.7, edgecolor='black')

    # Color bars by approach
    colors = ['steelblue' if approach == 'Univariate' else 'coral'
              for approach in data['Approach']]
    for bar, color in zip(bars, colors):
        bar.set_color(color)

    # Add expected rate line
    ax.axhline(y=alpha, color='red', linestyle='--', linewidth=2, label=f'Expected_Rate({alpha:.2f})')

    # Formatting
    ax.set_xticks(x_pos)
    ax.set_xticklabels(data['Method'], rotation=45, ha='right')
    ax.set_ylabel('Violation_Rate', fontsize=12)
    ax.set_title(f'VaR_Violation_Rates_{int(confidence*100)}%_Confidence', fontsize=14,
                fontweight='bold')
    ax.grid(True, alpha=0.3, axis='y')
    ax.legend()

    # Add value labels on bars
    for j, (idx, row) in enumerate(data.iterrows()):
        height = row['Violation_Rate_Numeric']
        ax.text(j, height + 0.002, f'{height:.3f}', ha='center', va='bottom', fontsize=9)

# Add legend for approaches
legend_elements = [
    Patch(facecolor='steelblue', edgecolor='black', label='Univariate_Methods'),
    Patch(facecolor='coral', edgecolor='black', label='Copula-Based_Methods')
]
fig.legend(handles=legend_elements, loc='upper_center', ncol=2, frameon=True, bbox_to_anchor=(0.5, 0.98))

```

```

plt.tight_layout(rect=[0, 0, 1, 0.96])
plt.savefig(OUT_DIR / 'portfolio_varViolation_rates.png')
plt.show()

# %% [markdown]
# At the 95% confidence level, both Gaussian and Student-*t* copulas exhibit violation rates
# that exceed the theoretical benchmark, performing similarly to or slightly worse than
# several univariate methods. At the 99% level, copula models show a modest improvement
# relative to some univariate approaches; however, their violation rates remain above the
# expected threshold, indicating persistent underestimation of tail risk.
#
# In contrast, univariate models that account for time-varying volatility, in particular the
# Filtered Historical Simulation (FHS) and parametric GARCH, achieve violation rates
# closest to the theoretical expectations across both confidence levels. This suggests
# that, for our particular portfolio and the period under study, modeling temporal
# dependence and volatility clustering is more impactful for accurate risk measurement
# than modeling cross-sectional dependence alone.

# %%
def plot_portfolio_var_methods_grid(L_portfolio,
                                      portfolio_rolling_results,
                                      copula_results=None,
                                      methods_univariate=None,
                                      methods_copula=None,
                                      confidences=(0.95, 0.99)):

    """
    Plot portfolio losses vs VaR for each method in a grid.
    One row per method, one column per confidence level.

    Parameters
    -----
    L_portfolio: pd.Series
        Portfolio losses time series (e.g., L_portfolio = -portfolio_returns).
    portfolio_rolling_results: dict-like
        Output of rolling forecast for the portfolio, indexed as:
        portfolio_rolling_results[method] ['VaR'] [confidence].
    copula_results: dict, optional
        Output of copula_portfolio_var_es, with keys 'gaussian' and 'student',
        accessed as copula_results[copula_type] ['VaR'] [confidence].
    methods_univariate: list of str, optional
        List of univariate method names to include, e.g.
        ["historical", "fhs", "gaussian", "student", "parametric"].
    methods_copula: list of str, optional
        List of copula method keys, subset of ["gaussian", "student"].
    confidences: iterable of float
        Confidence levels to plot (e.g., (0.95, 0.99)).
    """

    confidences = list(confidences)
    n_conf = len(confidences)

    if methods_univariate is None:
        methods_univariate = list(portfolio_rolling_results.keys())

    if (copula_results is not None) and (methods_copula is None):
        methods_copula = ["gaussian", "student"]
    elif copula_results is None:
        methods_copula = []

    # Build list of (source_type, key, label)
    # source_type: 'univariate' or 'copula'

```

```

method_list = []

# Univariate methods
for method in methods_univariate:
    method_list.append(("univariate", method, f"Univariate_{method}"))

# Copula methods
for cop in methods_copula:
    label = "Gaussian_Copula" if cop == "gaussian" else "Student-t_Copula"
    method_list.append(("copula", cop, label))

n_methods = len(method_list)
if n_methods == 0:
    print("No methods provided to plot.")
    return

fig, axes = plt.subplots(
    n_methods, n_conf,
    figsize=(6 * n_conf, 3 * n_methods),
    sharex=True
)

# Ensure axes is 2D
if n_methods == 1 and n_conf == 1:
    axes = np.array([[axes]])
elif n_methods == 1:
    axes = axes.reshape(1, -1)
elif n_conf == 1:
    axes = axes.reshape(-1, 1)

for i, (src_type, key, base_label) in enumerate(method_list):
    for j, conf in enumerate(confidences):
        ax = axes[i, j]

        # Get VaR series for this method & confidence
        if src_type == "univariate":
            VaR_series = portfolio_rolling_results[key] ["VaR"] [conf].dropna()
        else: # copula
            VaR_series = copula_results[key] ["VaR"] [conf].dropna()

        # Align portfolio loss with VaR series
        common_idx = L_portfolio.index.intersection(VaR_series.index)
        L_aligned = L_portfolio.loc[common_idx]
        V_aligned = VaR_series.loc[common_idx]

        # Plot full loss series (context)
        ax.plot(
            L_portfolio.index, L_portfolio.values,
            color="black", alpha=0.35, linewidth=0.9,
            label="Portfolio_Loss"
        )

        # Plot VaR
        ax.plot(
            V_aligned.index, V_aligned.values,
            linewidth=1.6,
            label=f"{base_label}_VaR_{int(conf*100)}%"
        )

        # Violations
        violations = L_aligned > V_aligned

```

```

        ax.scatter(
            L_aligned[violation].index,
            L_aligned[violation].values,
            color="red", marker="x", s=50,
            label="Violations"
        )

        # Labels / titles
        if j == 0:
            ax.set_ylabel("Loss / VaR", fontsize=11)
            ax.set_title(
                f"{{base_label}}-{int(conf*100)}%",
                fontsize=13, fontweight="bold"
            )

        # Grid & legend
        ax.grid(alpha=0.3)
        ax.legend(fontsize=8, loc="upper_left")

        # Date formatting: only bottom row shows tick labels
        if i == n_methods - 1:
            ax.xaxis.set_major_locator(mdates.MonthLocator(interval=3))
            ax.xaxis.set_major_formatter(mdates.DateFormatter('%Y-%m'))
            plt.setp(ax.get_xticklabels(), rotation=45, ha='right')
        else:
            ax.set_xticklabels([])

    fig.suptitle(
        "Portfolio_Rolling_VaR_vs_Realized_Losses\nMethods_x_Confidence_Levels",
        fontsize=16, fontweight="bold",
        y=0.995 # move title closer to the top
    )
    plt.tight_layout(pad=1.0) # reduce padding between title and axes
    fig.subplots_adjust(top=0.95) # key line: moves plots upward
    plt.savefig(OUT_DIR / 'portfolio_var_portfolio_methods.png')
    plt.show()

# %%
methods_univariate = ["historical", "fhs", "gaussian", "student", "parametric"]
methods_copula = ["gaussian", "student"] # keys in copula_results

plot_portfolio_var_methods_grid(
    L_portfolio=L_portfolio,
    portfolio_rolling_results=portfolio_rolling_results,
    copula_results=copula_results,
    methods_univariate=methods_univariate,
    methods_copula=methods_copula,
    confidences=(0.95, 0.99)
)

# %% [markdown]
# The figure presents the realized portfolio losses alongside the Value-at-Risk (VaR) estimates produced by different methods, at both **95%** and **99%** confidence levels. Red crosses indicate violations instances where the realized loss exceeds the forecasted VaR. These plots help assess the temporal behavior, responsiveness, and adequacy of each risk model.

#
# A well-calibrated VaR model should:
#
# - Adjust dynamically to changing market volatility,
# - Anticipate periods of elevated risk,

```

```

# - Produce violations at a frequency consistent with the chosen confidence level (5% or 1%)
,
# - Avoid large clusters of violations during stress periods.
#
# A visual inspection reveals that copula-based approaches do not systematically outperform
# univariate models such as FHS or GARCH-type specifications. Although copulas provide a
# more flexible framework to model cross-sectional dependence, their VaR curves often
# remain too smooth and fail to adjust rapidly during periods of heightened volatility. As
# a result, copula-based VaR estimates are frequently exceeded during market stress,
# leading to a substantial number of violations.
#
# In contrast, univariate methods that explicitly incorporate time-varying volatility,
# particularly FHS, tend to align more closely with the observed loss dynamics. These
# models react more promptly to volatility bursts and adjust their risk estimates
# accordingly, which explains their lower violation frequencies. This behavior is
# especially visible during the sharp increase in losses towards the end of the sample:
# while FHS-based VaR rapidly escalates to reflect the changing risk environment, copula-
# based VaR remains comparatively muted and is repeatedly breached.

# %%
def plot_portfolio_es_methods_grid(L_portfolio,
                                    portfolio_rolling_results,
                                    copula_results=None,
                                    methods_univariate=None,
                                    methods_copula=None,
                                    confidences=(0.95, 0.99)):

    """
    Plot portfolio losses vs ES for each method in a grid:
    one row per method, one column per confidence level.

    Parameters
    -----
    L_portfolio: pd.Series
        Portfolio losses time series (e.g., L_portfolio=portfolio_returns).
    portfolio_rolling_results: dict-like
        Output of rolling forecast for the portfolio, indexed as:
        portfolio_rolling_results[method]['ES'][confidence].
    copula_results: dict, optional
        Output of copula_portfolio_var_es, with keys 'gaussian' and 'student',
        accessed as copula_results[copula_type]['ES'][confidence].
    methods_univariate: list of str, optional
        List of univariate method names to include, e.g.
        ["historical", "fhs", "gaussian", "student", "parametric"].
    methods_copula: list of str, optional
        List of copula method keys, subset of ["gaussian", "student"].
    confidences: iterable of float
        Confidence levels to plot (e.g., (0.95, 0.99)).
    """

    confidences = list(confidences)
    n_conf = len(confidences)

    if methods_univariate is None:
        methods_univariate = list(portfolio_rolling_results.keys())

    if (copula_results is not None) and (methods_copula is None):
        methods_copula = ["gaussian", "student"]
    elif copula_results is None:
        methods_copula = []

    # Build list of (source_type, key, label)

```

```

method_list = []

# Univariate methods
for method in methods_univariate:
    method_list.append(("univariate", method, f"Univariate_{method}"))

# Copula methods
for cop in methods_copula:
    label = "Gaussian_Copula" if cop == "gaussian" else "Student-t_Copula"
    method_list.append(("copula", cop, label))

n_methods = len(method_list)
if n_methods == 0:
    print("No methods provided to plot.")
    return

fig, axes = plt.subplots(
    n_methods, n_conf,
    figsize=(6 * n_conf, 3 * n_methods),
    sharex=True
)

# Ensure axes is 2D
if n_methods == 1 and n_conf == 1:
    axes = np.array([[axes]])
elif n_methods == 1:
    axes = axes.reshape(1, -1)
elif n_conf == 1:
    axes = axes.reshape(-1, 1)

for i, (src_type, key, base_label) in enumerate(method_list):
    for j, conf in enumerate(confidences):
        ax = axes[i, j]

        # Get ES series for this method & confidence
        if src_type == "univariate":
            ES_series = portfolio_rolling_results[key]["ES"][conf].dropna()
        else: # copula
            ES_series = copula_results[key]["ES"][conf].dropna()

        # Align portfolio loss with ES series
        common_idx = L_portfolio.index.intersection(ES_series.index)
        L_aligned = L_portfolio.loc[common_idx]
        ES_aligned = ES_series.loc[common_idx]

        # Plot full loss series (context)
        ax.plot(
            L_portfolio.index, L_portfolio.values,
            color="black", alpha=0.35, linewidth=0.9,
            label="Portfolio_Loss"
        )

        # Plot ES
        ax.plot(
            ES_aligned.index, ES_aligned.values,
            linewidth=1.6, linestyle="--",
            label=f"{base_label}_ES_{int(conf*100)}%"
        )

        # Violations: loss > ES
        violations = L_aligned > ES_aligned

```

```

        ax.scatter(
            L_aligned[violations].index,
            L_aligned[violations].values,
            color="red", marker="x", s=50,
            label="Loss > ES"
        )

        # Labels / titles
        if j == 0:
            ax.set_ylabel("Loss / ES", fontsize=11)
            ax.set_title(
                f"{base_label}-{int(conf*100)}%",
                fontsize=13, fontweight="bold"
            )

        # Grid & legend
        ax.grid(alpha=0.3)
        ax.legend(fontsize=8, loc="upper_left")

        # Date formatting: only bottom row shows tick labels
        if i == n_methods - 1:
            ax.xaxis.set_major_locator(mdates.MonthLocator(interval=3))
            ax.xaxis.set_major_formatter(mdates.DateFormatter('%Y-%m'))
            plt.setp(ax.get_xticklabels(), rotation=45, ha='right')
        else:
            ax.set_xticklabels([])

    fig.suptitle(
        "Portfolio_Rolling_ES_vs_Realized_Losses\nMethods_x_Confidence_Levels",
        fontsize=16, fontweight="bold",
        y=0.995 # move title closer to the top
    )
    plt.tight_layout(pad=1.0)
    fig.subplots_adjust(top=0.95)
    plt.savefig(OUT_DIR / 'portfolio_es_portfolio_methods.png')
    plt.show()

# %%
methods_univariate = ["historical", "fhs", "gaussian", "student", "parametric"]
methods_copula = ["gaussian", "student"] # keys in copula_results

plot_portfolio_es_methods_grid(
    L_portfolio=L_portfolio,
    portfolio_rolling_results=portfolio_rolling_results,
    copula_results=copula_results,
    methods_univariate=methods_univariate,
    methods_copula=methods_copula,
    confidences=(0.95, 0.99)
)

# %% [markdown]
# Expected Shortfall (ES) measures the *average loss* in the tail beyond the VaR threshold.
# A well-calibrated ES model should therefore:
#
# - **Consistently exceed realized losses** during normal periods, since ES is a *worst-case conditional loss*,
# - **Increase sharply during market stress**, capturing escalating downside risk,
# - **Avoid repeated exceedances** losses should only rarely be higher than ES,
# - Reflect changes in **both volatility and dependence** across assets, since ES is highly sensitive to tail co-movements.
#

```

```

# In short, ES should form a **protective upper boundary** for losses in extreme conditions.
# If losses frequently cross above ES, the model is underestimating tail risk.
#
# These plots above compare realized portfolio losses with Expected Shortfall (ES) forecasts
# from different univariate and copula-based models at **95%** and **99%** confidence
# levels. Red crosses indicate instances where losses exceed the ES estimate.
#
# Similar to the VaR analysis, the copula-based approaches do not demonstrate a clear
# performance advantage over univariate methods. Although ES is a coherent risk measure
# and theoretically more sensitive to tail events, the copula-based ES estimates remain
# relatively stable throughout the sample and fail to escalate during periods of
# heightened market stress. Consequently, several realized losses exceed the copula-based
# ES thresholds, indicating that these models tend to underestimate tail risk.
#
# By contrast, univariate models that incorporate time-varying volatility, particularly the
# Filtered Historical Simulation (FHS) and, to a lesser extent, parametric GARCH-type,
# provide more reactive ES estimates. These methods adjust rapidly to volatility bursts,
# producing higher ES values when market conditions deteriorate. This dynamic behavior
# results in fewer ES breaches, especially in the latter part of the sample, where the
# portfolio experiences the most extreme losses. The responsiveness of FHS to evolving
# volatility patterns highlights the importance of modeling temporal dependence and
# volatility clustering when forecasting downside risk.

# %% [markdown]
# ## Conclusion
#
# The objective of this analysis was to assess whether dependence modeling improves
# portfolio risk estimation. While copula-based approaches explicitly capture the cross-
# sectional dependence structure between assets, the empirical results do not provide
# clear evidence of superior performance compared to univariate models.
#
# An explanation for these findings lies in the characteristics of the underlying portfolio.
# The three assets considered do not exhibit pronounced dependence, which limits the
# potential gains from employing copula-based models. In such a context, explicitly
# modeling cross-sectional dependence offers little improvement. Moreover, copulas do not
# account for volatility clustering or sudden volatility bursts if we give them simple
# Gaussian or Student-$t$ marginals, unlike FHS and GARCH. Since extreme losses during the
# sample period appear to be primarily driven by volatility dynamics rather than
# dependence structures, correctly modeling time-varying volatility proves more important
# than capturing marginal dependence. Consequently, the relative underperformance of
# copula-based methods is consistent with the underlying data-generating process and the
# nature of the portfolio risk.
#
# Therefore, we cannot conclude that dependence modeling via copulas systematically improves
# the results. Copula-based methods do not outperform well-specified univariate models
# such as FHS, and their benefits become visible only at extreme confidence levels without
# consistently translating into better backtesting outcomes. While copulas generally
# perform better than some simpler univariate approaches, this incremental improvement
# remains insufficient to claim superiority over methods that explicitly model volatility
# dynamics. As a result, the evidence does not support the claim that copulas provide
# superior VaR or ES estimates in this setting.

```